

QMCPACK

User's Guide and Developer's Manual

v1.0
February 17, 2016

Contents

1	Introduction	7
1.1	Quickstart and a first QMCPACK calculation	7
1.2	Authors and History	9
1.3	Support and Contacting the Developers	10
1.4	Performance	10
1.5	Open source license	11
1.6	Contributing to QMCPACK	12
1.7	QMCPACK Roadmap	12
1.7.1	Code	13
1.7.2	Documentation	13
2	Obtaining, installing and validating QMCPACK	14
2.1	Installation steps	14
2.2	Obtaining the latest release version	15
2.3	Obtaining the latest development version	15
2.4	Prerequisites	15
2.5	Building with CMake	16
2.5.1	Quick build instructions (try first)	17
2.5.2	Environment variables	17
2.5.3	Configuration options	18
2.5.4	Configure and build using cmake and make	19
2.5.5	Example configure and build	19
2.5.6	Build scripts	20
2.6	Installation instructions for common workstations and supercomputers	20
2.6.1	Installing on Ubuntu Linux or other apt-get based distributions	21
2.6.2	Installing on CentOS Linux or other yum based distributions	21
2.6.3	Installing on Mac OS X using Macports	22
2.6.4	Installing on Mac OS X using Homebrew (brew)	23
2.6.5	Installing on ANL ALCF Mira/Cetus IBM Blue Gene/Q	24

2.6.6	Installing on ORNL OLCF Titan Cray XK7 (NVIDIA GPU accelerated)	25
2.6.7	Installing on ORNL OLCF Titan Cray XK7 (CPU version)	25
2.6.8	Installing on ORNL OLCF Eos Cray XC30	26
2.6.9	Installing on NERSC Edison Cray XC30	26
2.6.10	Installing on NERSC Cori (Phase 1) Cray XC40	27
2.7	Testing and validation of QMCPACK	28
2.7.1	Automatic tests of QMCPACK	29
2.8	Building ppconvert, a pseudopotential format converter	30
2.9	Installing and patching Quantum Espresso	31
2.10	How to build the fastest executable version of QMCPACK	31
2.11	Troubleshooting the installation	32
3	Running QMCPACK	33
3.1	Command line options	33
3.2	Input files	33
3.3	Output files	33
3.4	Running in parallel	33
3.4.1	MPI	33
3.4.2	Use of OpenMP threads	34
3.4.3	Running on GPU machines	35
4	Units used in QMCPACK	38
5	Input file overview	39
6	Specifying the system to be simulated	42
6.1	Specifying the simulation cell	42
6.2	Specifying the particle set	42
6.2.1	Input specification	42
6.2.2	Detailed attribute description	43
6.2.3	Example use cases	45
7	Trial wavefunction specification	48
7.1	Single-particle orbitals	48
7.1.1	Spline basis sets	49
7.1.2	Gaussian basis sets	50
7.1.3	Plane-wave basis sets	51
7.1.4	Homogeneous electron gas	52
7.2	Jastrow Factors	53
7.3	One-body Jastrow functions	54

7.3.1	Input Specification	54
7.3.2	Spline form	54
7.4	Multideterminant wavefunctions	57
7.5	Backflow wavefunctions	57
7.5.1	Input Specifications	57
7.5.2	Example Use Case	58
7.5.3	Additional Information	59
8	Hamiltonian and Observables	60
8.1	The Hamiltonian	60
8.2	Pair potentials	62
8.2.1	Coulomb potentials	63
8.2.2	Pseudopotentials	65
8.2.3	Modified periodic Coulomb interaction/correction	67
8.3	General estimators	67
8.3.1	Chiesa-Ceperley-Martin-Holzmam kinetic energy correction	69
8.3.2	Density estimator	69
8.3.3	Spin density estimator	71
8.3.4	Pair correlation function, $g(r)$	72
8.3.5	Static structure factor, $S(k)$	74
8.3.6	Energy density estimator	75
8.3.7	One body density matrix	80
8.4	Forward Walking Estimators	85
8.5	“Force” estimators	86
9	Quantum Monte Carlo Methods	88
9.1	Variational Monte Carlo	89
9.2	Wavefunction Optimization	91
9.3	Diffusion Monte Carlo	91
9.4	Reptation Monte Carlo	95
10	Analysing QMCPACK data	98
10.1	Using the qmca tool	98
10.2	Densities and spin-densities	98
10.3	Energy densities	98
11	Examples	99
11.1	Using Nexus	99
11.1.1	H ₂ O Molecule with Quantum ESPRESSO Orbitals	99
11.1.2	LiH Crystal with Quantum ESPRESSO Orbitals	103

12 Lab 1: Monte Carlo Statistical Analysis	108
12.1 Topics covered in this Lab	108
12.2 Lab directories and files	109
12.3 Atomic units	111
12.4 Reviewing statistics	111
12.5 Inspecting Monte Carlo data	113
12.6 Averaging quantities in the MC data	115
12.7 Evaluating MC simulation quality	116
12.7.1 Tracing MC quantities	117
12.7.2 Blocking away autocorrelation	117
12.7.3 Balancing autocorrelation and acceptance ratio	119
12.7.4 Considering variance	119
12.8 Reducing statistical error bars	120
12.8.1 Increasing MC sampling	120
12.8.2 Improving the basis set	121
12.8.3 Adding a Jastrow factor	121
12.9 Scaling to larger numbers of electrons	122
12.9.1 Calculating the efficiency	122
12.9.2 Scaling up	122
13 Lab 2: QMC Basics	124
13.1 Topics covered in this Lab	124
13.2 Lab directories and files	125
13.3 Obtaining and converting a pseudopotential for oxygen	126
13.4 Optimization walkthrough with QMCPACK: neutral O atom	127
13.5 DMC timestep extrapolation I: neutral O atom	132
13.6 DMC timestep extrapolation II: O atom ionization potential	136
13.7 DMC workflow automation with Nexus	137
13.8 Automated binding curve of the oxygen dimer	139
13.9 (Optional) Running your system with QMCPACK	147
14 Lab 4: Using PWSCF and QMCPACK to perform total energy calculations of condensed systems	152
14.1 Preliminaries	152
14.2 Total energy of BCC beryllium	154
14.3 Handling a 2D system: graphene	158
14.4 Conclusion	159
14.5 Acknowledgment	159
15 Contributing to the Manual	160

Chapter 1

Introduction

QMCPACK is an open-source, high-performance electronic structure code that implements numerous Quantum Monte Carlo algorithms. Its main applications are electronic structure calculations of molecular, periodic 2D and periodic 3D solid-state systems. Variational Monte Carlo (VMC), diffusion Monte Carlo (DMC) and a number of other advanced QMC algorithms are implemented. By directly solving the Schrodinger equation, QMC methods offer greater accuracy than methods such as density functional theory, but at a trade-off of much greater computational expense. Distinct from many other correlated many-body methods, QMC methods are readily applicable to both bulk (periodic) and isolated molecular systems.

QMCPACK is written in C++ and designed with the modularity afforded by object-oriented programming. It makes extensive use of template metaprogramming to achieve high computational efficiency. Due to the modular architecture, the addition of new wavefunctions, algorithms, and observables is relatively straightforward. For parallelization QMCPACK utilizes a fully hybrid (OpenMP,CUDA)/MPI approach to optimize memory usage and to take advantage of the growing number of cores per SMP node or graphical processing units (GPUs) and accelerators. High parallel and computational efficiencies are achievable on the largest supercomputers. Finally, QMCPACK utilizes standard file formats for input and output in XML and HDF5 to facilitate data exchange.

This manual currently serves as an introduction to the essential features of QMCPACK and a guide to installing and running it. Over time this manual will be expanded to including a fuller introduction to QMC methods in general and to include more of the specialized features in QMCPACK.

1.1 Quickstart and a first QMCPACK calculation

If you are keen to get started this section describes how to quickly build and run QMCPACK on standard UNIX or Linux-like system. The autoconfiguring build system usually works without much fuss on these systems. If C++, MPI, BLAS/LAPACK, FFTW, HDF5, and CMake are

already installed, QMCPACK can be built and run within five minutes. For supercomputers, cross-compilation systems, and other computer clusters the build system may require hints on the locations of libraries and which versions to use, typical of any code, see Chapter 2. Section includes complete examples for common workstations and supercomputers that you can reuse.

To build QMCPACK:

1. Download the latest QMCPACK distribution from <http://www.qmcpack.org>
2. Untar the archive, e.g., `tar xvf qmcpack.v1.3.tar.gz`
3. Check the instructions in the README
4. Run CMake in a suitable build directory to configure QMCPACK for your system: `cd qmcpack/build; cmake ..`
5. If CMake is unable to find all needed libraries, see Chapter 2 for instructions and specific build instructions for common systems.
6. Build QMCPACK: `make` or `make -j 16`, the latter for a faster parallel build on a system using, e.g., 16 processes.
7. The QMCPACK executable is `bin/qmcpack`

QMCPACK is distributed with examples illustrating different capabilities. Most of the examples are designed to run quickly with modest resources. We'll run a short diffusion Monte Carlo calculation of a water molecule:

1. Go to the appropriate example directory: `cd ../examples/molecules`
2. (Optional) Put the QMCPACK binary on your path:
`export PATH=$PATH:location-of-qmcpack/build/bin`
3. Run QMCPACK: `.././build/bin/qmcpack simple-H2O.xml` or `qmcpack simple-H2O.xml` if you followed the step above.
4. The run will output to the screen and generate a number of files:

```
$ls H2O*
H2O.HF.wfs.xml      H2O.s001.scalar.dat  H2O.s002.cont.xml
H2O.s002.qmc.xml    H2O.s002.stat.h5     H2O.s001.qmc.xml
H2O.s001.stat.h5    H2O.s002.dmc.dat     H2O.s002.scalar.dat
```

5. Partially summarized results are in the standard text files with the suffixes `scalar.dat` and `dmc.dat`. They are viewable with any standard editor.

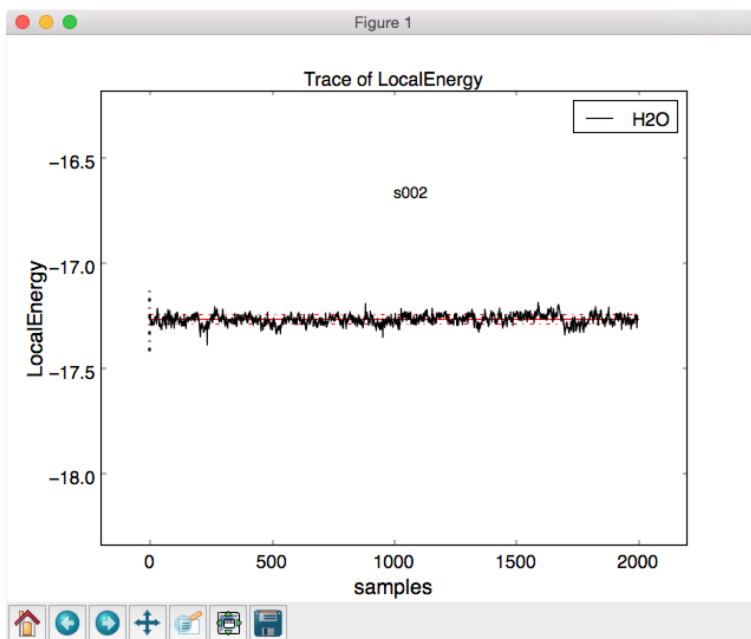


Figure 1.1: Trace of walker energies produced by qmca tool for simple water molecule example.

If you have python and matplotlib installed, you can use the `qmca` analysis utility to produce statistics and plots of the data. See Chapter 10 for information on analysing QMCPACK data.

```
export PATH=$PATH:location-of-qmcpack/nexus/executables
export PYTHONPATH=$PYTHONPATH:location-of-qmcpack/nexus/library
qmca H2O.s002.scalar.dat          # For statistical analysis of the DMC data
qmca -t -q e H2O.s002.scalar.dat # Graphical plot of DMC energy
```

The last command will produce a graph as per Fig. 1.1. This shows the average energy of the DMC walkers at each timestep. In a real simulation we would have to check equilibration, convergence with walker population, timestep etc.

Congratulations, you have completed a DMC calculation with QMCPACK!

1.2 Authors and History

QMCPACK was initially written by Jeongnim Kim while in the group of Prof. David Ceperley at the University of Illinois at Urbana-Champaign, with later contributions at Oak Ridge National Laboratory. Over the years, many others have contributed, particularly students and researchers in the groups of Prof. David Ceperley and Prof. Richard M. Martin, as well as staff at Lawrence

Livermore National Laboratory, Sandia National Laboratories, Argonne National Laboratory, and Oak Ridge National Laboratory.

The primary and original author of the code is Jeongnim Kim. Additional developers, contributors, and advisors include: Anouar Benali, Mark A. Berrill, David M. Ceperley, Simone Chiesa, Raymond C. III Clay, Bryan Clark, Kris T. Delaney, Kenneth P. Esler, Paul R. C. Kent, Jaron T. Krogel, Ying Wai Li, Ye Luo, Jeremy McMinis, Miguel A. Morales, William D. Parker, Nichols A. Romero, Luke Shulenburger, Norman M. Tubman, and Jordan E. Vincent.

If you should be added to this list please let us know.

Development of QMCPACK has been supported financially by several grants, including:

- “Network for ab initio many-body methods: development, education and training” supported through the Predictive Theory and Modeling for Materials and Chemical Science program by the U.S. Department of Energy Office of Science, Basic Energy Sciences.
- “QMC Endstation”, supported by Accelerating Delivery of Petascale Computing Environment at the DOE Leadership Computing Facility at ORNL.
- PetaApps, supported by the U. S. National Science Foundation.
- Materials Computational Center, supported by the U.S. National Science Foundation.

1.3 Support and Contacting the Developers

Questions about installing, applying or extending QMCPACK can be posted on the QMCPACK Google group <https://groups.google.com/forum/#!forum/qmcpack>. You may also email any of the developers, but we recommend checking the group first. Particular attention is given to any problem reports.

1.4 Performance

QMCPACK implements modern Monte Carlo algorithms, is highly parallel, and is also written using very efficient code for high per-CPU or on node performance. In particular the code is highly vectorizable, giving high performance on modern CPUs and GPUs. We believe QMCPACK delivers performance either comparable to or better than other QMC codes when similar calculations are run, particularly for the most common QMC methods and for large systems. If you find a calculation where this is not the case, or you simply find performance slower than expected, please post on the Google group or contact one of the developers. These reports are valuable. If your calculation is sufficiently mainstream we will optimize QMCPACK to improve the performance.

1.5 Open source license

QMCPACK is distributed under the University of Illinois/NCSA Open Source License.

University of Illinois/NCSA Open Source License

Copyright (c) 2003, University of Illinois Board of Trustees.
All rights reserved.

Developed by:

Jeongnim Kim
Condensed Matter Physics,
National Center for Supercomputing Applications, University of Illinois
Materials computation Center, University of Illinois
<http://www.mcc.uiuc.edu/qmc/>

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the ``Software''), to deal with the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

- * Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimers.
- * Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimers in the documentation and/or other materials provided with the distribution.
- * Neither the names of the NCSA, the MCC, the University of Illinois, nor the names of its contributors may be used to endorse or promote products derived from this Software without specific prior written permission.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE CONTRIBUTORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE,

ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS WITH THE SOFTWARE.

Copyright is generally believed to remain with the authors of the individual sections of code. See the various notations in the source code as well as the code history.

1.6 Contributing to QMCPACK

QMCPACK is fully open source and we welcome contributions. Please post on the QMCPACK Google group or contact the developers. If you are planning a development, early discussions are encouraged. We will be able to tell you if anyone else is working on a similar feature or if any related work has been done in the past. Credit for your contribution can be obtained, e.g., through citation of a paper, or becoming one of the authors on the next version of the standard QMCPACK reference citation.

Please note the following guidelines for a contributions:

- Additions should be fully synchronized with the latest release version and ideally the latest development SVN. Merging of code developed on older versions is error prone.
- Code should be cleanly formatted, commented, portable, and accessible to other programmers. i.e. If you need to use any clever tricks, add a comment to note this, why the trick is needed, how it works etc. Although we like high performance, ease of maintenance and accessibility are also considerations.
- Comment your code. You are not only writing it for the compiler but also for other humans! (We know this is a repeat of the previous point, but it is important enough to repeat.)
- Write a brief description of the method, algorithms and inputs and outputs suitable for inclusion in this manual.
- Develop some short tests that exercise the functionality that can be used for validation and for examples. We can help with this and their integration into the test system.

1.7 QMCPACK Roadmap

A general outline of the QMCPACK roadmap is given below. Suggestions for improvements are welcome, particularly those that would facilitate new scientific applications. For example, if an interface to a particular quantum chemical or density functional code would help, this would be given strong consideration.

1.7.1 Code

We will continue improving the accessibility and usability of QMCPACK, by combinations of more convenient input parameters, improved workflow, integration with more quantum chemical and density functional codes, and a wider range of examples.

In terms of methodological development, we expect to significantly increase the range of QMC algorithms in QMCPACK in the near future.

Computationally, we are porting QMCPACK to the next generation of supercomputer systems. The internal changes required to run on these systems efficiently are expected to benefit *all* platforms due to improved vectorization, cache utilization and memory performance.

1.7.2 Documentation

This manual currently describes the core features of QMCPACK that are required for routine research calculations. i.e. the VMC and DMC methods, how to obtain and optimize trial wavefunctions, and simple observables. Over time this manual will be expanded to include a broader introduction to QMC methods and to describe more features of the code.

Due to its history as a research code, QMCPACK contains a variety of additional QMC methods, trial wavefunction forms, potentials (etc.) that, although not critical, may be very useful for specialized calculations or particular material or chemical systems. These “secret features” (every code has these) are not actually secret but simply lack descriptions, example inputs, and tests. You are encouraged to browse and read the source code to find them. New descriptions will be added over time, but can also be prioritized and added on request, e.g. if a specialized Jastrow factor would help or an historical Jastrow form is needed for benchmarking.

Chapter 2

Obtaining, installing and validating QMCPACK

This chapter describes how to obtain, build and validate QMCPACK. This process is designed to be as simple as possible and should be no harder than building a modern plane-wave density functional theory code such as Quantum Espresso, QBox, or VASP. Parallel builds enable a complete compilation in under 2 minutes on a fast multicore system. If you are unfamiliar with building codes we suggest working with your system administrator to install QMCPACK.

2.1 Installation steps

To install QMCPACK, follow the steps listed below. Full details of each step are given in the referenced sections.

1. Download the source code, Sections 2.2 or 2.3.
2. Verify that you have the required compilers, libraries and tools installed, Section 2.4.
3. Run the cmake configure step and build with make, Section 2.5 and 2.5.1. Some examples for common systems are given in Section 2.6.
4. Run the tests to verify QMCPACK, Section 2.7.
5. Build the ppconvert utility in QMCPACK, Section 2.8.
6. Download and patch Quantum Espresso. This patch adds the pw2qmcpack utility, Section 2.9.

Hints for high performance are in Section 2.10. Troubleshooting suggestions are in Section 2.11.

Note that there are two different QMCPACK executables that can be produced: the general one, which is the default, and the “complex” version which support periodic calculations at arbitrary twist angles and k-points. This second version is enabled via a cmake configuration parameter, see Section 2.5.3. The general version only supports wavefunctions that can be made real. If you run a calculation that needs the complex version, QMCPACK will stop and inform you.

2.2 Obtaining the latest release version

Major releases of QMCPACK are distributed from <http://www.qmcpack.org>. These releases undergo the most testing. Unless there are specific reasons we encourage all production calculations to use the latest release versions.

Releases are usually compressed tar files indicating the version number, date, and often the source code revision control number corresponding to the release.

- Download the latest QMCPACK distribution from <http://www.qmcpack.org>.
- Untar the archive, e.g., `tar xvf qmcpack_v1.3.tar.gz`

2.3 Obtaining the latest development version

The most recent development version of QMCPACK can be obtained anonymously via

`svn checkout https://svn.qmcpack.org/svn/trunk`

Once checked-out, updates can be made via the standard `svn update`.

The subversion repository contains the day-to-day development source with the latest updates, bugfixes etc. This may be useful for updates to the build system to support new machines, for support of the latest versions of Quantum Espresso, or for updates to the documentation. Note that the development version may not be fully consistent with the online documentation. We attempt to keep the development version fully working. However, please be sure to run the tests and compare with previous release versions before using for any serious calculations. We try to keep bugs out, but occasionally they crawl in! Reports of any breakages are appreciated.

2.4 Prerequisites

The following are required to build QMCPACK. For workstations, these are available via the standard package manager. On shared supercomputers this software is usually installed by default and is often access via a modules environment - check your system documentation.

Use of the latest versions of all compilers and libraries is strongly encouraged, but not absolutely essential. Generally newer versions are faster - see Section 2.10 for performance suggestions.

- C/C++ compilers such as GCC, Intel, IBM XLC. CLANG-based compilers are not yet supported by the build system, but the source code is ready.
- MPI library such as OpenMPI <http://open-mpi.org>
- BLAS/LAPACK, numerical and linear algebra libraries. Use platform-optimized libraries where available, such as Intel MKL. ATLAS or other optimized open-source libraries may also be used <http://math-atlas.sourceforge.net>
- CMake, build utility, <http://www.cmake.org>
- Libxml2, XML parser, <http://xmlsoft.org>
- HDF5, portable I/O library, <http://www.hdfgroup.org/HDF5/>
- BOOST, peer-reviewed portable C++ source libraries, <http://www.boost.org>
- FFTW, FFT library, <http://www.fftw.org/>

To build the GPU accelerated version of QMCPACK an installation of NVIDIA CUDA development tools is required. Ensure that this is compatible with the C and C++ compiler versions you plan to use. Supported versions are included in the NVIDIA release notes.

Many of the utilities provided with QMCPACK use python (v2). The numpy and matplotlib libraries are required for full functionality.

Note that the standalone einspline library used by previous versions of QMCPACK is no longer required. A more optimized version is included inside. The standalone version should *not* be on any standard search paths because conflicts between the old and new include files can result.

2.5 Building with CMake

The build system for QMCPACK is based on CMake. It will autoconfigure based on the detected compilers and libraries. The most recent version of CMake has the best detection for the greatest variety of systems - at the time of writing this means CMake 3.4.3. The much older CMake 2.8 is known to work, but might not work optimally on your system.

Previously QMCPACK made extensive use of toolchains, but the build system has since been updated to eliminate the use of toolchain files for most cases. The build system is verified to work with GNU, Intel, and IBM XLC compilers. Specific compile options can be specified either through specific environmental or CMake variables. When the libraries are installed in standard locations, e.g., /usr, /usr/local, there is no need to set environmental or cmake variables for the packages.

2.5.1 Quick build instructions (try first)

If you are feeling lucky and are on a standard UNIX-like system such as a Linux workstation, the following might quickly give a working QMCPACK:

The safest quick build option is to specify the C and C++ compilers through their MPI wrappers. Here we use Intel MPI and Intel compilers. Move to the build directory, run `cmake` and `make`

```
cd build
cmake -DCMAKE_C_COMPILER=mpiicc -DCMAKE_CXX_COMPILER=mpicpc ..
make -j 8
```

You can increase the “8” to the number of cores on your system for faster builds. Substitute `mpicc` and `mpicxx` or other wrapped compiler names to suit your system. e.g. With OpenMPI use

```
cd build
cmake -DCMAKE_C_COMPILER=mpicc -DCMAKE_CXX_COMPILER=mpicxx ..
make -j 8
```

If you are feeling particularly lucky, you can skip the compiler specification:

```
cd build
cmake ..
make -j 8
```

The complexities of modern computer hardware and software systems are such that you should check that the autoconfiguration system has made good choices and picked optimized libraries and compiler settings before doing significant production. i.e. Check the details below. We give examples for a number of common systems in Section 2.6.

2.5.2 Environment variables

A number of environmental variables affect the build. In particular they can control the default paths for libraries, the default compilers, etc. The list of environmental variables is given below:

CXX	C++ compiler
CC	C Compiler
MKL_HOME	Path for MKL
LIBXML2_HOME	Path for libxml2
HDF5_ROOT	Path for HDF5
BOOST_ROOT	Path for Boost
FFTW_HOME	Path for FFTW

2.5.3 Configuration options

In addition to reading the environmental variables, CMake provides a number of optional variables that can be set to control the build and configure steps. When passed to CMake, these variables will take precedent over the environmental and default variables. To set them add `-D FLAG=VALUE` to the configure line between the `cmake` command and the path to the source directory.

- Key QMCPACK build options

QMC_CUDA	Enable CUDA and GPU acceleration (1:yes, 0:no)
QMC_COMPLEX	Build the complex (general twist/k-point) version (1:yes, 0:no)

- General build options

CMAKE_BUILD_TYPE	A variable which controls the type of build (defaults to Release). Possible values are: None (Do not set debug/optimize flags, use CMAKE_C_FLAGS or CMAKE_CXX_FLAGS) Debug (create a debug build) Release (create a release/optimized build) RelWithDebInfo (create a release/optimized build with debug info) MinSizeRel (create an executable optimized for size)
CMAKE_C_COMPILER	Set the C compiler
CMAKE_CXX_COMPILER	Set the C++ compiler
CMAKE_C_FLAGS	Set the C flags. Note: to prevent default debug/release flags from being used, set the CMAKE_BUILD_TYPE=None Also supported: CMAKE_C_FLAGS_DEBUG, CMAKE_C_FLAGS_RELEASE, and CMAKE_C_FLAGS_RELWITHDEBINFO
CMAKE_CXX_FLAGS	Set the C++ flags. Note: to prevent default debug/release flags from being used, set the CMAKE_BUILD_TYPE=None Also supported: CMAKE_CXX_FLAGS_DEBUG, CMAKE_CXX_FLAGS_RELEASE, and CMAKE_CXX_FLAGS_RELWITHDEBINFO

- Additional QMCPACK build options

QMC_DATA	Specify data directory for QMCPACK (currently unused, but likely to be used for performance tests)
QMC_INCLUDE	Add extra include paths
QMC_EXTRA_LIBS	Add extra link libraries
QMC_BUILD_STATIC	Add <code>-static</code> flags to build

- libxml related

```
Libxml2_INCLUDE_DIRS  Specify include directories for libxml2
Libxml2_LIBRARY_DIRS  Specify library directories for libxml2
```

- FFTW related

```
FFTW_INCLUDE_DIRS     Specify include directories for FFTW
FFTW_LIBRARY_DIRS     Specify library directories for FFTW
```

2.5.4 Configure and build using cmake and make

To configure and build QMPACK, move to build directory, run cmake and make

```
cd build
cmake ..
make -j 8
```

As you will have gathered, cmake encourages “out of source” builds, where all the files for a specific build configuration reside in their own directory separate from the source files. This allows multiple builds to be created from the same source files which is very useful where the filesystem is shared between different systems. You can also build versions with different settings (e.g. QMC.COMPLEX) and different compiler settings. The build directory does not have to be called build - use something descriptive such as build_machinename or build_complex. The “..” in the cmake line refers to the directory containing CMakeLists.txt. Update the “..” for other build directory locations.

2.5.5 Example configure and build

- Set the environments (the examples below assume bash, Intel compilers and MKL library)

```
export CXX=icpc
export CC=icc
export MKL_HOME=/usr/local/intel/mkl/10.0.3.020
export LIBXML2_HOME=/usr/local
export HDF5_ROOT=/usr/local
export BOOST_ROOT=/usr/local/boost
export FFTW_HOME=/usr/local/fftw
```

- Move to build directory, run cmake and make

```
cd build
cmake -D CMAKE_BUILD_TYPE=Release ..
make -j 8
```

2.5.6 Build scripts

It is recommended to create a helper script that contains the configure line for CMake. This is particularly useful when avoiding environmental variables, packages are installed in custom locations, or if the configure line is long or complex. In this case it is also recommended to add "rm -rf CMake*" before the configure line to remove existing CMake configure files to ensure a fresh configure each time that the script is called. Deleting all the files in the build directory is also acceptable. If you do so we recommend to add some sanity checks in case the script is run from the wrong directory, e.g., checking for the existence of some QMCPACK files.

Some build script examples for different systems are given in the config directory. For example, on Cray systems these scripts might load the appropriate modules to set the appropriate programming environment, specific library versions etc.

An example script build.sh is given below. It is overly complex for the sake of example:

```
export CXX=mpic++
export CC=mpicc
export ACML_HOME=/opt/acml-5.3.1/gfortran64
export HDF5_ROOT=/opt/hdf5
export BOOST_ROOT=/opt/boost

rm -rf CMake*

cmake \
  -D CMAKE_BUILD_TYPE=Debug \
  -D Libxml2_INCLUDE_DIRS=/usr/include/libxml2 \
  -D Libxml2_LIBRARY_DIRS=/usr/lib/x86_64-linux-gnu \
  -D FFTW_INCLUDE_DIRS=/usr/include \
  -D FFTW_LIBRARY_DIRS=/usr/lib/x86_64-linux-gnu \
  -D QMC_EXTRA_LIBS="-ldl ${ACML_HOME}/lib/libacml.a -lgfortran" \
  -D QMC_DATA=/projects/QMCPACK/qmc-data \
  ..
```

2.6 Installation instructions for common workstations and supercomputers

This section describes how to build QMCPACK on various common systems including multiple Linux distributions, Apple OS X, and various supercomputers. The examples should serve as good starting points for building QMCPACK on similar machines. For example, the software environment on modern Crays is very consistent. Note that updates to operating systems and

system software may require small modifications to these recipes. See Section 2.10 for key points to check to obtain highest performance and Section 2.11 for troubleshooting hints.

2.6.1 Installing on Ubuntu Linux or other apt-get based distributions

The following is designed to obtain a working QMCPACK build on e.g. a student laptop, starting from a basic Linux installation with none of the developer tools installed. Fortunately, all the required packages are available in the default repositories making for a quick installation. Note that for convenience we use a generic BLAS. For production a platform optimized BLAS should be used.

```
apt-get subversion cmake g++ openmpi-bin libopenmpi-dev libboost-dev
apt-get libatlas-base-dev liblapack-dev libhdf5-dev libxml2-dev fftw3-dev
export CXX=mpiCC
cd build
cmake ..
make -j 8
ls -l bin/qmcpack
```

For qmca and other tools to function, we install some python libraries:

```
sudo apt-get install python-numpy python-matplotlib
```

2.6.2 Installing on CentOS Linux or other yum based distributions

The following is designed to obtain a working QMCPACK build on e.g. a student laptop, starting from a basic Linux installation with none of the developer tools installed. CentOS 7 (Red Hat compatible) is using gcc 4.8.2. The installation is only complicated by the need to install another repository to obtain HDF5 packages which are not available by default. Note that for convenience we use a generic BLAS. For production a platform optimized BLAS should be used.

```
sudo yum install make cmake gcc gcc-c++ subversion openmpi openmpi-devel fftw fftw-devel \
                boost boost-devel libxml2 libxml2-devel
sudo yum install blas-devel lapack-devel atlas-devel
module load mpi
```

To setup repoforge as a source for the HDF5 package, go to <http://repoforge.org/use> . Install the appropriate up to date release package for your OS. By default the CentOS Firefox will offer to run the installer. The CentOS 6.5 settings were usable for HDF5 on CentOS 7 in July 2014, but use CentOS 7 versions when they become available.

```
sudo yum install hdf5 hdf5-devel
```

To build QMCPACK

```
module load mpi/openmpi-x86_64
which mpirun
# Sanity check; should print something like /usr/lib64/openmpi/bin/mpirun
export CXX=mpiCC
cd build
cmake ..
make -j 8
ls -l bin/qmcpack
```

2.6.3 Installing on Mac OS X using Macports

These instructions assume a fresh installation of macports and for consistency with current Linux distributions, use the gcc 4.8.2 compiler. More recent versions are fine, but it is vital to ensure matching compilers/options for all packages and to force use of what is installed in /opt/local. As with the Linux examples above, this build is very good if not optimal, and is easily good enough to learn QMCPACK or experiment on a travel laptop.

Note that we utilize the Apple provided Accelerate framework for optimized BLAS.

Follow the Macports install instructions <https://www.macports.org/>

- Install Xcode and the Xcode Command Line Tools
- Agree to Xcode license in Terminal: `sudo xcodebuild -license`
- Install MacPorts for your version of OS X

Install the required tools:

```
sudo port install gcc48
sudo port select gcc mp-gcc48 # Set default

sudo port install openmpi-devel-gcc48
sudo port select set mpi openmpi-devel-gcc48-fortran # Set default

# Sanity check
mpiCXX -v
#should return gcc version 4.8.2 (MacPorts gcc48 4.8.2_2) or similar.

sudo port install fftw-3 +gcc48
sudo port install cmake # already cmake 3 as of 2014/7/29
```

```

sudo port install boost +gcc48
sudo port install libxml2
sudo port install hdf5-18 +gcc48

sudo port select set python python27
sudo port install py27-matplotlib # For qmca

```

QMCPACK build:

```

export CXX=mpiCXX
export CC=/opt/local/bin/gcc
export LIBXML2_HOME=/opt/local/
export HDF5_HOME=/opt/local
export BOOST_HOME=/opt/local
export FFTW_HOME=/opt/local
cd build
cmake ..
make -j 6 # Adjust for available core count
ls -l bin/qmcpack

```

2.6.4 Installing on Mac OS X using Homebrew (brew)

Homebrew is a package manager for OS X that provides a convenient route to install all the QMCPACK dependencies. The following recipe will install the latest available versions of each package. This was successfully tested under OS X 10.11 “El Capitan” in February 2016. Note that it is necessary to build the MPI software from source to use the brew-provided gcc instead of Apple CLANG.

1. Install Homebrew from <http://brew.sh/>

```

/usr/bin/ruby -e "$(curl -fsSL
  https://raw.githubusercontent.com/Homebrew/install/master/install)"

```

2. Install the prerequisites

```

brew install gcc # Builds full gcc 5 from scratch, will take 30 minutes
export HOMEBREW_CXX=g++-5
export HOMEBREW_CC=gcc-5
brew install mpich2 --build-from-source
# Build from source required to use homebrew compiled compilers as
# opposed to Apple CLANG. Check "mpicc -v" indicates Homebrew gcc 5.x.x

```

```
brew install cmake
brew install fftw
brew install boost
brew install homebrew/science/hdf5
#Note: Libxml2 is not required via brew since OS X already includes it.
```

3. Configure and build QMCPACK

```
cmake -DCMAKE_C_COMPILER=/usr/local/bin/mpicc \
      -DCMAKE_CXX_COMPILER=/usr/local/bin/mpicxx ..
make -j 12
```

4. Run the short tests. When mpich is used for the first time, OS X will request approval of the network connection.

```
ctest -R short
```

2.6.5 Installing on ANL ALCF Mira/Cetus IBM Blue Gene/Q

Mira/Cetus is a Blue Gene/Q supercomputer at Argonne National Laboratory's Argonne Leadership Computing Facility (ANL ALCF). Mira has 49152 compute nodes and each node has a 16-core PowerPC A2 processor with 16 GB DDR3 memory. Due to the fact that the login nodes and the compute nodes have different processors with distinct instruction sets, cross-compiling is required on this platform. See details about using Blue Gene/Q at <http://www.alcf.anl.gov/user-guides/compiling-linking>. On Mira, compilers are loaded via softenv and users need to add +mpiwrapper-xl and +cmake in \$HOME/.soft. In order to build QMCPACK, a toolchain file is provided for setting up CMake and the cmake command should be executed twice.

```
cd build
cmake -DCMAKE_TOOLCHAIN_FILE=../config/BGQToolChain.cmake ..
cmake -DCMAKE_TOOLCHAIN_FILE=../config/BGQToolChain.cmake ..
make -j 16
ls -l bin/qmcpack
```

In addition, adding a very useful cmake option `-DCMAKE_VERBOSE_MAKEFILE=TRUE` allows printing all the build commands during the make step. Alternatively you can use `make VERBOSE=1`.

2.6.6 Installing on ORNL OLCF Titan Cray XK7 (NVIDIA GPU accelerated)

Titan is a GPU accelerated supercomputer at Oak Ridge National Laboratory's Oak Ridge Leadership Computing Facility (ORNL OLCF). Each compute node has a 16 core AMD 2.2GHz Opteron 6274 (Interlagos) and an NVIDIA Kepler accelerator. The standard Cray software environment is available, with libraries accessed via modules. The only extra settings required to build the GPU version are the cudatoolkit module and specifying -DQMC.CUDA=1 on the cmake configure line.

Note that on Crays the compiler wrappers "CC" and "cc" are used. The build system checks for these and does not (should not) use the compilers directly.

```
module swap PrgEnv-pgi PrgEnv-gnu # Use gnu compilers
module load cudatoolkit           # CUDA for GPU build
module load cray-hdf5
module load cmake
module load fftw
export FFTW_HOME=$FFTW_DIR/..
module load boost
mkdir build_titan_gpu
cd build_titan_gpu
cmake -DQMC_CUDA=1 ..             # Must enable CUDA capabilities
make -j 8
ls -l bin/qmcpack
```

2.6.7 Installing on ORNL OLCF Titan Cray XK7 (CPU version)

As noted in Section 2.6.6 for the GPU, building on Crays requires only loading the appropriate library modules.

```
module swap PrgEnv-pgi PrgEnv-gnu # Use gnu compilers
module unload cudatoolkit         # No CUDA for CPU build
module load cray-hdf5
module load cmake
module load fftw
export FFTW_HOME=$FFTW_DIR/..
module load boost
mkdir build_titan_cpu
cd build_titan_cpu
cmake ..
make -j 8
ls -l bin/qmcpack
```

2.6.8 Installing on ORNL OLCF Eos Cray XC30

Eos is Cray XC30 with 16 core Intel Xeon E5-2670 processors connected by the Aries interconnect. The build process is identical to Titan, except that we use the default Intel programming environment. This is usually preferred to GNU.

```
module load cray-hdf5
module load cmake
module load fftw
export FFTW_HOME=$FFTW_DIR/..
module load boost
mkdir build_eos
cd build_eos
cmake ..
make -j 8
ls -l bin/qmcpack
```

2.6.9 Installing on NERSC Edison Cray XC30

Edison is a Cray XC30 with dual 12-core Intel "Ivy Bridge" nodes installed at NERSC. The build settings are identical to eos.

```
module load cray-hdf5
module load cmake
module load fftw
export FFTW_HOME=$FFTW_DIR/..
module load boost
mkdir build_edison
cd build_edison
cmake ..
make -j 8
ls -l bin/qmcpack
```

When the above was tested on 1 February 2016, the following module and software versions were present:

```
qmcpack@edison04:trunk> module list
```

Currently Loaded Modulefiles:

- | | |
|-------------------------------------|--------------------------------------|
| 1) modules/3.2.10.3 | 16) alps/5.2.3-2.0502.9295.14.14.ari |
| 2) nsg/1.2.0 | 17) rca/1.0.0-2.0502.57212.2.56.ari |
| 3) eswrap/1.1.0-1.020200.1130.0 | 18) atp/1.8.3 |
| 4) switch/1.0-1.0502.57058.1.58.ari | 19) PrgEnv-intel/5.2.56 |

- | | |
|--|----------------------|
| 5) craype-network-aries | 20) craype-ivybridge |
| 6) craype/2.5.0 | 21) cray-shmem/7.3.0 |
| 7) intel/15.0.1.133 | 22) cray-mpich/7.3.0 |
| 8) cray-libsci/13.3.0 | 23) slurm/edison |
| 9) udreg/2.3.2-1.0502.9889.2.20.ari | 24) altd/2.0 |
| 10) ugni/6.0-1.0502.10245.9.9.ari | 25) darshan/2.3.0 |
| 11) pmi/5.0.10-1.0000.11050.0.0.ari | 26) subversion/1.7.9 |
| 12) dmapp/7.0.1-1.0502.10246.8.47.ari | 27) cray-hdf5/1.8.14 |
| 13) gni-headers/4.0-1.0502.10317.9.2.ari | 28) cmake/2.8.11.2 |
| 14) xpmem/0.1-2.0502.57015.1.15.ari | 29) fftw/3.3.4.6 |
| 15) dvs/2.5_0.9.0-1.0502.1958.2.55.ari | 30) boost/1.54 |

2.6.10 Installing on NERSC Cori (Phase 1) Cray XC40

Cori is a Cray XC40 with 16-core Intel "Haswell" nodes installed at NERSC. The build settings are identical to eos.

```
module load cray-hdf5
module load cmake
module load fftw
export FFTW_HOME=$FFTW_DIR/..
module load boost
mkdir build_cori
cd build_cori
cmake ..
make -j 8
ls -l bin/qmcpack
```

When the above was tested on 1 February 2016, the following module and software versions were present:

```
qmcpack@cori05:trunk> module list
Currently Loaded Modulefiles:
  1) nsrg/1.2.0
  2) modules/3.2.10.3
  3) eswrap/1.1.0-1.020200.1231.0
  4) switch/1.0-1.0502.60522.1.61.ari
  5) intel/16.0.0.109
  6) craype-network-aries
  7) craype/2.4.2
  8) cray-libsci/13.2.0
 15) dvs/2.5_0.9.0-1.0502.2188.1.116.ari
 16) alps/5.2.4-2.0502.9774.31.11.ari
 17) rca/1.0.0-2.0502.60530.1.62.ari
 18) atp/1.8.3
 19) PrgEnv-intel/5.2.82
 20) craype-haswell
 21) cray-shmem/7.2.5
 22) cray-mpich/7.2.5
```

9) udreg/2.3.2-1.0502.10518.2.17.ari	23) slurm/cori
10) ugni/6.0-1.0502.10863.8.29.ari	24) cray-hdf5/1.8.14
11) pmi/5.0.9-1.0000.10911.0.0.ari	25) gcc/5.1.0
12) dmapp/7.0.1-1.0502.11080.8.76.ari	26) cmake/3.3.2
13) gni-headers/4.0-1.0502.10859.7.8.ari	27) fftw/3.3.4.5
14) xpmem/0.1-2.0502.64982.5.3.ari	28) boost/1.59

2.7 Testing and validation of QMCPACK

We **strongly encourage** running the included tests each time QMCPACK is built. These compare the results from the executable with known-good mean-field, quantum chemical, and other QMC results.

The tests included with QMCPACK currently test only the VMC code with single determinant wavefunction and simple spline Jastrow wavefunctions, and for gaussian and periodic spline basis sets. Although not yet comprehensive, it is extremely unlikely that, e.g., DMC will be correct if the VMC tests do not pass. We check that the known mean field results are obtained with no Jastrow. When Jastrow functions are included we test against previous QMC data. The tests are statistical with a generous 3σ tolerance, however the system sizes are small, typically < 10 electrons, so the error bars are typically small.

The “short” tests only take a few minutes on a 16 core machine. You can run these tests using the command below in the build directory:

```
ctest -R short    # Run the tests with "short" in their name
```

The output should be similar to the following:

```
Test project build_gcc
  Start 1: short-LiH_dimer_ae-vmc_hf_noj-16-1
1/44 Test #1: short-LiH_dimer_ae-vmc_hf_noj-16-1 ..... Passed 11.20 sec
  Start 2: short-LiH_dimer_ae-vmc_hf_noj-16-1-kinetic
2/44 Test #2: short-LiH_dimer_ae-vmc_hf_noj-16-1-kinetic ..... Passed 0.13 sec
..
42/44 Test #42: short-mono0_1x1x1_pp-vmc_sdj-1-16 ..... Passed 10.02 sec
  Start 43: short-mono0_1x1x1_pp-vmc_sdj-1-16-totenergy
43/44 Test #43: short-mono0_1x1x1_pp-vmc_sdj-1-16-totenergy ..... Passed 0.08 sec
  Start 44: short-mono0_1x1x1_pp-vmc_sdj-1-16-samples
44/44 Test #44: short-mono0_1x1x1_pp-vmc_sdj-1-16-samples ..... Passed 0.08 sec

100% tests passed, 0 tests failed out of 44

Total Test time (real) = 167.14 sec
```

Note that the number of tests that are run varies between the standard, complex, and GPU compilations.

The full set of tests consist of significantly longer versions of the short tests. They require several hours each to run yielding a much more stringent test of the code. To run all the tests simply run `ctest` in the build directory:

```
ctest          # Run all the tests. This will take several hours.
```

You can also run verbose tests which direct the QMCPACK output to the standard output:

```
ctest -V -R short  # Verbose short tests
```

The test system includes specific tests for the complex version of the code.

The data files for the tests are located in the tests directory. The runs occur in `build/src/QMCAApp/test/test_name`. The numerical comparisons and test definitions are in `src/QMCAApp/test/CMakeLists.txt`. If *all* the QMC tests fail it is likely that the appropriate `mpiexec` (or `aprun`, `srun`) is not being called or found. If the QMC runs appear to work but all the other tests fail it is possible that python is not working on your system - we suggest checking some of the test outputs in `build/src/QMCAApp/test/test_name`.

Note that because the tests are very small, consisting of only a few electrons, the performance is not representative of larger calculations. For example, while the calculations might fit in cache, there will be essentially no vectorization due to the small electron counts. **The tests should not be used for any benchmarking or performance analysis.** Dedicated larger runs are required.

2.7.1 Automatic tests of QMCPACK

The QMCPACK developers run automatic tests of QMCPACK on several different computer systems, many on a continuous basis. We currently test the following combinations nights (workstation) or weekly (supercomputers):

- On a Red Hat Linux workstation:
 - GCC 4.8.2 with OpenMPI and CUDA 7.0 (GPU build, run on NVIDIA K40s)
 - GCC 4.8.2 with OpenMPI
 - Intel 2016 with Intel MPI and MKL
 - Intel 2015 with Intel MPI and MKL and CUDA 7.0 (GPU build, run on NVIDIA K40s)
 - Intel 2015 with Intel MPI and MKL
- On Eos, a Cray XC30 Intel machine:
 - The default Intel programming environment and compiler with Cray MPI and Intel MKL

2.9 Installing and patching Quantum Espresso

For trial wavefunctions obtained in a plane-wave basis we mainly support Quantum Espresso. Note that ABINIT and QBox were supported historically and could be reactivated.

Quantum Espresso currently stores wavefunctions in a non-standard internal “save” format. To convert these to a conventional HDF5 format file we have developed a converter, pw2qmcpack. This is an add on to the Quantum Espresso distribution.

To simplify the process of patching Quantum Espresso we have developed a script that will automatically download and patch the source code. The patches are specific to each version. e.g. To download and patch QE v5.3.0:

```
cd external_codes/quantum_espresso
./download_and_patch_qe5.3.0.sh
```

After running the patch, you must configure Quantum Espresso with the HDF5 capability enabled, i.e.

```
cd espresso-5.3.0
./configure --with-hdf5 HDF5_DIR=/opt/local # Specify HDF5 base directory
```

The complete process is described in external_codes/quantum_espresso/README.

2.10 How to build the fastest executable version of QMCPACK

To build the fastest version of QMCPACK we recommend the following:

- Use the latest C++ compilers available for your system. Substantial gains have been made optimizing C++ in recent years.
- Use a vendor optimized BLAS library such as Intel MKL and AMD ACML. Although QMC does not make extensive use of linear algebra, it is used in the VMC wavefunction optimizer and also to apply the orbital coefficients in local basis calculations.
- Use a vector math library such as Intel VML. For periodic calculations, the calculation of the structure factor and Ewald potential benefit from vectorized evaluation of sin and cos. Currently we only autodetect Intel VML, as provided with MKL, but support for MASSV and AMD LibM is included via #defines. See, e.g. src/Numerics/e2iphi.h. For large supercells, this optimization can gain 10% in performance.

Note that greater speedups of QMC calculations can usually be obtained by carefully choosing the required statistics for each investigation. i.e. Do not compute smaller error bars than necessary.

2.11 Troubleshooting the installation

Some tips to help troubleshoot installations of QMCPACK:

- First, build QMCPACK on a workstation that you control, or on any system that has a simple and up-to-date set of development tools. You can compare the results of cmake and QMCPACK on this system with any more difficult systems you encounter.
- Use up to date development software, particularly a recent CMake.
- Verify that the compilers and libraries that you expect are being configured. It is common to have multiple versions installed. The configure system will stop at the first version it finds which might not be the most recent. If this occurs, specify the appropriate directories and files directly (Section 2.5.3). e.g. `cmake -DCMAKE_C_COMPILER=/full/path/to/mpicc -DCMAKE_CXX_COMPILER=/full/path/to/mpicxx ..`
- To monitor the compiler and linker settings, use a verbose build, “make VERBOSE=1”. If an individual source file fails to compile you can experiment by hand using the output of the verbose build to reconstruct the full compilation line.

If you still have problems please post to the QMCPACK Google group with full details, or contact a developer.

Chapter 3

Running QMCPACK

3.1 Command line options

3.2 Input files

3.3 Output files

scalar.dat
dmc.dat
stat.h5
config.h5

3.4 Running in parallel

3.4.1 MPI

QMCPACK is fully parallelized with MPI. When performing an ensemble job, all the MPI ranks are first equally divided into groups which perform individual QMC calculations. Within one calculation, all the walkers are fully distributed across all the MPI ranks in the group. Since MPI requires distributed memory, there must be at least one MPI per node. To maximize the efficiency, more facts should be taken into account. When using MPI+threads on compute nodes with more than one NUMA domain (e.g., AMD Interlagos CPU on Titan or a node with multiple CPU sockets), it is recommended to place as many MPI ranks as the number of NUMA domains if the memory is sufficient. On clusters with more than just one GPU per node (NVIDIA Tesla K80), it requires to use the same number of MPI ranks as the number of GPUs per node in order to let each MPI rank take one GPU.

3.4.2 Use of OpenMP threads

Modern processors integrate multiple identical cores even with hardware threads on a single die to increase the total performance and maintain a reasonable power draw. QMCPACK takes advantage of all that compute capability on a processor by using threads via OpenMP programming model as well as threaded linear algebra libraries. By default, QMCPACK is always built with OpenMP enabled. When launching calculations, users should instruct QMCPACK to create the right number of threads per MPI rank by specifying environmental variable `OMP_NUM_THREADS`. Even in the GPU accelerated version, using threads significantly reduces the time spent on the calculations performed on CPU.

Performance consideration

As walkers are the basic units of workload in QMC algorithms, they are loosely coupled and distributed across all the threads. For this reason, the best strategy to run QMCPACK efficiently is to feed enough walkers to the available threads.

In a VMC calculation, the code automatically raises the actual number of walkers per MPI rank to the number of available threads if the user-specified number of walkers is smaller, see “walkers/mpi=XXX” in the VMC output. In a DMC calculation, the target number of walkers should be chosen to be slightly smaller than a multiple of the total number of available threads across all the MPI ranks belongs to this calculation. Since the number of walkers varies from generation to generation, its dynamical value should be slightly smaller or equal to that multiple most of the time.

Memory consideration

When using threads, some memory objects shared by all the threads. Usually these memory are read-only when the walkers are evolving, for instance the ionic distance table and wavefunction coefficients. If a wavefunction is represented by B-splines, the whole table is shared by all the threads. It usually takes a large chunk of memory when a large primitive cell was used in the simulation. Its actual size is reported as “MEMORY increase XXX MB BsplineSetReader” in the output file. See details about how to reduce it in section 7.1.1.

The other memory objects which are distinct for each walker during random walk need to be associated with individual walkers and can not be shared. This part of memory grows linearly as the number of walkers per MPI rank. Those objects include wavefunction values (Slater determinants) at given electronic configurations and electron related distance tables (electron-electron distance table). Those matrices dominate the N^2 scaling of the memory usage per walker.

3.4.3 Running on GPU machines

The GPU version on the NVIDIA CUDA platform is fully incorporated into the main trunk. Currently some commonly used functionalities for solid-state and molecular systems using B-spline single-particle orbitals is supported. A detailed description of the GPU implementation can be found in Ref. [?].

Current GPU implementation assumes one MPI process per GPU. Vectorization is achieved over walkers, that is, all walkers are propagated in parallel. In each GPU kernel, loops over electrons, atomic cores or orbitals are further vectorized to exploit an additional level of parallelism and to allow coalesced memory access.

Supported GPU features

1. Quantum Monte Carlo methods:
 - (a) Variational Monte Carlo (VMC).
 - (b) Diffusion Monte Carlo (DMC).
 - (c) Limited support for wavefunction optimization.
2. Boundary conditions:
 - (a) Periodic and open boundary conditions are fully supported.
 - (b) Twist-averaged boundary condition is supported for only real-valued wavefunctions.
 - (c) Mixed boundary conditions and complex wavefunctions (e.g. fixed phase) are not yet supported.
3. Wavefunctions:
 - (a) Single Slater determinants with 3D B-spline orbitals. Only real-valued wavefunctions is supported, but tiling complex orbitals to supercells is supported as long as each k-point is a multiple of half a G-vector of the supercell.
 - (b) Mixed basis representation in which orbitals are represented as 1D splines times spherical harmonics in spherical regions (muffin tins) around atoms, and 3D B-splines in the interstitial region.
 - (c) One-body and two-body Jastrows represented as 1D B-splines are supported. Note that only single-precision arithmetic is fully functional at the time of writing.
4. Interaction types:
 - (a) Semilocal (nonlocal and local) pseudopotentials.

- (b) Coulomb interaction (electron-electron, electron-ion).
- (c) Model periodic Coulomb (MPC) interaction.

Compiling the GPU code

To build the executable `qmcpack` with GPU support, follow these steps:

1. Make sure NVIDIA's CUDA compiler, `nvcc`, is in the search path. In most cases, CMake should be able to locate the `nvcc` compiler on the system automatically.
2. (a) Run CMake with the argument `QMC_CUDA` switched on:


```
cd build
cmake -D QMC_CUDA=1 ..
make
```

or

- (b) If a CMake toolchain file is used, switch on `QMC_CUDA` by including this line in the toolchain file:


```
SET (QMC_CUDA 1)
```

 Then compile the code as before:


```
cd build
cmake -D CMAKE_TOOLCHAIN_FILE=[toolchain name] ..
make
```

CMake variables for adjusting CUDA code build features

These values can be changed by passing them as CMake's command line options with the `-D` flag, or using a toolchain file to overwrite the default values.

1. `QMC_CUDA`
 - `=0` (default): no GPU support, build QMCPACK as a CPU code
 - `=1` : build QMCPACK with GPU support
2. `CUDA_PRECISION`
 - `=float` (default): single precision arithmetics and data types will be used for GPU kernels
 - `=double` : double precision arithmetics and data types will be used for GPU kernels (Warning: not fully functional!)

Performance consideration

The relative speedup of the GPU implementation increases with both the number of electrons and the number of walkers running on a GPU. Typically, 128-256 walkers per GPU utilize sufficient number of threads to operate the GPU efficiently and to hide memory-access latency.

To achieve better performance, current implementation utilizes single precision operations on most GPU calculations, except for matrix inversions where double precision is required to retain high accuracy. The single precision GPU code is as accurate as the double precision CPU code up to a certain system size. Cross checking and verification of accuracy are encouraged for systems with more than approximately 1500 electrons.

Memory consideration

In the GPU implementation, each walker has an anonymous buffer on the GPU's global memory to store temporary data associated with the wavefunctions. Therefore, the amount of memory available on a GPU limits the number of walkers and eventually the system size that it can process.

If the GPU memory is exhausted, reduce the number of walkers per GPU. Coarsening the grids of the B-splines representation (by decreasing the value of meshfactor in the input file) can also lower the memory usage, at the expense (risk) of obtaining inaccurate results. Proceed with caution if this option has to be considered.

Chapter 4

Units used in QMCPACK

Internally, QMCPACK uses atomic units throughout. Unless stated, all inputs and outputs are also in atomic units. For convenience the analysis tools offer conversions to eV, Ry, Angstrom, Bohr etc.

Chapter 5

Input file overview

This chapter introduces XML as it is used in QMCPACK's input file. The focus is on the XML file format itself and the general structure of the input file rather than an exhaustive discussion of all keywords and structure elements.

QMCPACK uses XML to represent structured data in its input file. Instead of text blocks like

```
begin project
  id      = vmc
  series = 0
end project

begin vmc
  move      = pbyp
  blocks    = 200
  steps     = 10
  timestep  = 0.4
end vmc
```

QMCPACK input looks like

```
<project id="vmc" series="0">
</project>
```

```

<qmc method="vmc" move="pbyp">
  <parameter name="blocks" > 200 </parameter>
  <parameter name="steps" > 10 </parameter>
  <parameter name="timestep"> 0.4 </parameter>
</qmc>

```

XML elements start with `<element_name>`, end with `</element_name>`, and can be nested within each other to denote substructure (the trial wavefunction is composed of a Slater determinant and a Jastrow factor, which are each further composed of ...). `id` and `series` are attributes of the `<project/>` element. XML attributes are generally used to represent simple values, like names, integers, or real values. Similar functionality is also commonly provided by `<parameter/>` elements like those shown above.

The overall structure of the input file reflects different aspects of the QMC simulation: the simulation cell, particles, trial wavefunction, Hamiltonian, and QMC run parameters. A condensed version of the actual input file is shown below:

```

<?xml version="1.0"?>
<simulation>

  <project id="vmc" series="0">
    ...
  </project>

  <qmcsystem>

    <simulationcell>
      ...
    </simulationcell>

    <particleset name="e">
      ...
    </particleset>

    <particleset name="ion0">
      ...
    </particleset>

```



```

<wavefunction name="psi0" ... >
  ...
  <determinantset>
    <slaterdeterminant>
      ..
    </slaterdeterminant>
  </determinantset>
  <jastrow type="One-Body" ... >
    ...
  </jastrow>
  <jastrow type="Two-Body" ... >
    ...
  </jastrow>
</wavefunction>

<hamiltonian name="h0" ... >
  <pairpot type="coulomb" name="ElecElec" ... />
  <pairpot type="coulomb" name="IonIon" ... />
  <pairpot type="pseudo" name="PseudoPot" ... >
    ...
  </pairpot>
</hamiltonian>

</qmcsystem>

<qmc method="vmc" move="pbyp">
  <parameter name="warmupSteps"> 20 </parameter>
  <parameter name="blocks" > 200 </parameter>
  <parameter name="steps" > 10 </parameter>
  <parameter name="timestep" > 0.4 </parameter>
</qmc>

</simulation>

```

The omitted portions (...) are more fine-grained inputs such as the axes of the simulation cell, the number of up and down electrons, positions of atomic species, external orbital files, starting Jastrow parameters, and external pseudopotential files.

Chapter 6

Specifying the system to be simulated

6.1 Specifying the simulation cell

6.2 Specifying the particle set

The `particleset` blocks specify the particles in the QMC simulations: their types, attributes (mass, charge, valence), and positions.

6.2.1 Input specification

particleset element				
parent elements:	simulation			
child elements:	group, attrib			
attribute :				
name	datatype	values	default	description
name/id	text	<i>any</i>	e	Name of particle set
size ^o	integer	<i>any</i>	0	Number of particles in set
random ^o	text	yes/no	no	Randomize starting positions
randomsrc/ random_source ^o	text	particleset.name	<i>none</i>	Particle set to randomize

group element					
parent elements:	particleset				
child elements:	parameter, attrib				
attribute :					
name	datatype	values	default	description	
name	text	any	e	Name of particle set	
size ^o	integer	any	0	Number of particles in set	
mass ^o	real	any	1	Mass of particles in set	
unit ^o	text	au/amu	au	Units for mass of particles	
parameters					
name	datatype	values	default	description	
charge	real	any	0	Charge of particles in set	
valence	real	any	0	Valence charge of particles in set	
atomicnumber	integer	any	0	Atomic number of particles in set	

attrib element					
parent elements:	particleset,group				
attribute :					
name	datatype	values	default	description	
name	string	any	none	Name of attrib	
datatype	string	intArray, realArray, posArray, stringArray	none	Type of data in attrib	
size ^o	string	any	none	Size of data in attrib	

6.2.2 Detailed attribute description

particleset required attributes

- **name/id**
Unique name for the particle set. Default is “e” for electrons. “i” or “ion0” is typically used for ions.

particleset optional attributes

- **size**
Number of particles in set
- **random**
Randomize starting positions of particles. Each component of each particle’s position is

randomized independently in the range of the simulation cell in that component's direction.

- **randomsrc/random_source**
Specify source particle set around which to randomize the initial positions of this particle set.

name required attributes

- **name/id**
Unique name for the particle set group. Typically, element symbols are used for ions and “u” or “d” for spin-up and spin-down electron groups, respectively.

group optional attributes

- **mass**
Mass of particles in set.
- **unit**
Units for mass of particles in set (au [$m_e = 1$] or amu [$\frac{1}{12}m_{12C} = 1$]).

6.2.3 Example use cases

Listing 6.1: particleset elements for ions and electrons randomizing electron start positions.

```
<particleset name="i" size="2">
  <group name="Li">
    <parameter name="charge">3.000000</parameter>
    <parameter name="valence">3.000000</parameter>
    <parameter name="atomicnumber">3.000000</parameter>
  </group>
  <group name="H">
    <parameter name="charge">1.000000</parameter>
    <parameter name="valence">1.000000</parameter>
    <parameter name="atomicnumber">1.000000</parameter>
  </group>
  <attrib name="position" datatype="posArray" condition="1">
    0.0  0.0  0.0
    0.5  0.5  0.5
  </attrib>
  <attrib name="ionid" datatype="stringArray">
    Li H
  </attrib>
</particleset>
<particleset name="e" random="yes" randomsrc="i">
  <group name="u" size="2">
    <parameter name="charge">-1</parameter>
  </group>
  <group name="d" size="2">
    <parameter name="charge">-1</parameter>
  </group>
</particleset>
```

Listing 6.2: particleset elements for ions and electrons specifying electron start positions

```

<particleset name="e">
  <group name="u" size="4">
    <parameter name="charge">-1</parameter>
    <attrib name="position" datatype="posArray">
      2.9151687332e-01 -6.5123272502e-01 -1.2188463918e-01
      5.8423636048e-01 4.2730406357e-01 -4.5964306231e-03
      3.5228575807e-01 -3.5027014639e-01 5.2644808295e-01
      -5.1686250912e-01 -1.6648002292e+00 6.5837023441e-01
    </attrib>
  </group>
  <group name="d" size="4">
    <parameter name="charge">-1</parameter>
    <attrib name="position" datatype="posArray">
      3.1443445436e-01 6.5068682609e-01 -4.0983449009e-02
      -3.8686061749e-01 -9.3744432997e-02 -6.0456005388e-01
      2.4978241724e-02 -3.2862514649e-02 -7.2266047173e-01
      -4.0352404772e-01 1.1927734805e+00 5.5610824921e-01
    </attrib>
  </group>
</particleset>
<particleset name="ion0" size="3">
  <group name="O">
    <parameter name="charge">6</parameter>
    <parameter name="valence">4</parameter>
    <parameter name="atomicnumber">8</parameter>
  </group>
  <group name="H">
    <parameter name="charge">1</parameter>
    <parameter name="valence">1</parameter>
    <parameter name="atomicnumber">1</parameter>
  </group>
  <attrib name="position" datatype="posArray">
    0.0000000000e+00 0.0000000000e+00 0.0000000000e+00
    0.0000000000e+00 -1.4308249289e+00 1.1078707576e+00
    0.0000000000e+00 1.4308249289e+00 1.1078707576e+00
  </attrib>
  <attrib name="ionid" datatype="stringArray">
    O H H
  </attrib>
</particleset>

```

Listing 6.3: particleset elements for ions specifying positions by ion type

```
<particleset name="ion0">
  <group name="O" size="1">
    <parameter name="charge">6</parameter>
    <parameter name="valence">4</parameter>
    <parameter name="atomicnumber">8</parameter>
    <attrib name="position" datatype="posArray">
      0.0000000000e+00 0.0000000000e+00 0.0000000000e+00
    </attrib>
  </group>
  <group name="H" size="2">
    <parameter name="charge">1</parameter>
    <parameter name="valence">1</parameter>
    <parameter name="atomicnumber">1</parameter>
    <attrib name="position" datatype="posArray">
      0.0000000000e+00 -1.4308249289e+00 1.1078707576e+00
      0.0000000000e+00 1.4308249289e+00 1.1078707576e+00
    </attrib>
  </group>
</particleset>
```

Chapter 7

Trial wavefunction specification

7.1 Single-particle orbitals

7.1.1 Spline basis sets

7.1.2 Gaussian basis sets

7.1.3 Plane-wave basis sets

7.1.4 Homogeneous electron gas

The interacting Fermi Liquid has its own special determinantset for filling up a Fermi surface. The shell number can be specified separately for both spin up and spin down. This determines how many electrons to include of each time, only closed shells are currently implemented. The shells are filled according to the rules of a square box, if other lattice vectors are used, the electrons may not fill up a complete shell.

This following example can also be used for Helium simulations too, by specifying the proper pair interaction in the Hamiltonian section.

Listing 7.1: 2D Fermi Liquid example: particle specification

```
<qmcsystem>
<simulationcell name="global">
  <parameter name="rs" pol="0" condition="74">6.5</parameter>
  <parameter name="bconds">p p p</parameter>
  <parameter name="LR_dim_cutoff">15</parameter>
</simulationcell>
<particleset name="e" random="yes">
  <group name="u" size="37">
    <parameter name="charge">-1</parameter>
    <parameter name="mass">1</parameter>
  </group>
  <group name="d" size="37">
    <parameter name="charge">-1</parameter>
    <parameter name="mass">1</parameter>
  </group>
</particleset>
</qmcsystem>
```

Listing 7.2: 2D Fermi Liquid example (Slater Jastrow wave function)

```
<qmcsystem>
  <wavefunction name="psi0" target="e">
    <determinantset type="electron-gas" shell="7" shell2="7" randomize="true">
    </determinantset>
    <jastrow name="J2" type="Two-Body" function="Bspline" print="no">
      <correlation speciesA="u" speciesB="u" size="8" cusp="0">
        <coefficients id="uu" type="Array" optimize="yes">
        </coefficients>
      </correlation>
      <correlation speciesA="u" speciesB="d" size="8" cusp="0">
        <coefficients id="ud" type="Array" optimize="yes">
        </coefficients>
      </correlation>
    </jastrow>
  </wavefunction>
</qmcsystem>
```

7.2 Jastrow Factors

7.3 One-body Jastrow functions

The text below is rough draft and to be replaced with a complete and accurate description! In particular, the table columns and entries require consideration.

The one-body Jastrow is designed to... and is normally used in conjunction with an additional two-body term. Many different one-body are implemented...

The jastrow function is specified within a **wavefunction** element and must contain one of more **correlation** elements specifying additional parameters as well as the actual coefficients. Section 7.3.2 gives examples of the typical nesting of **jastrow**, **correlation**, and **coefficient** elements.

7.3.1 Input Specification

Jastrow element				
name	datatype	values	defaults	description
name	text	?	?	Unique name for this Jastrow function
type	text	One-body	(required)	Define a one-body function
function	text	Bspline	(required)	BSpline Jastrow
		Pade		Pade form
	
source	text	?	?	
print	text	?	?	

7.3.2 Spline form

The one-body spline Jastrow function is the most commonly used one-body Jastrow for solids. This form was first described and used in [?].

$$J1 = \sum_I^{ion0} \sum_i^e u_{ab}(|r_r - R_I|) \quad (7.1)$$

where u_{ab} is an interpolating spline between zero distance and r_{cut} . In 3D periodic systems the default cutoff distance is the Wigner Seitz cell radius. For other periodicities including isolated molecules the r_{cut} must be specified. The gradient at zero distance is... cusp can be set.... r_i and R_I are most commonly the electron and ion positions, but any particlesets that can provide the needed centers can be used.

Input Specification

Additional information:

Correlation element				
name	datatype	values	defaults	description
elementType	fillme			
size				
rcut				
cusp				
print	text	?	(optional)	
elements				
Coefficients				
Contents				
(None)				

- **rcut**. The cutoff distance for the function in atomic units. For 3D fully periodic systems this parameter is optional and a default of the Wigner Seitz cell radius is used. Otherwise this parameter is required.

Coefficients element				
name	datatype	values	defaults	description
id	text			Unique identifier
type	text	Array		
elements				
(None)				
Contents				
(no name)	real array		zeros	Jastrow coefficients

Example use cases

Specify a spin-independent function with four parameters. Because rcut is not specified, the default cutoff of the Wigner Seitz cell radius is used; this Jastrow must be used with a 3D periodic system such as a bulk solid. The source of the ionic positions is set "i".

```
<jastrow name="J1" type="One-Body" function="Bspline" print="yes" source="i">
  <correlation elementType="C" cusp="0.0" size="4">
    <coefficients id="C" type="Array"> 0 0 0 0 </coefficients>
  </correlation>
</jastrow>
```

Specify a spin-dependent function with seven upspin and seven downspin parameters. The cutoff distance is set to 6 atomic units.

```
<jastrow name="J1" type="One-Body" function="Bspline" source="ion0" spin="yes">
  <correlation speciesA="C" speciesB="u" size="7" rcut="6">
    <coefficients id="eCu" type="Array">
      0.0 0.0 0.0 0.0 0.0 0.0 0.0
    </coefficients>
  </correlation>
  <correlation speciesA="C" speciesB="d" size="7" rcut="6">
    <coefficients id="eCd" type="Array">
      0.0 0.0 0.0 0.0 0.0 0.0 0.0
    </coefficients>
  </correlation>
</jastrow>
```


7.4 Multideterminant wavefunctions

7.5 Backflow wavefunctions

One can perturb the nodal surface of a single-slater/multi-slater wavefunction through use of a backflow-transformation. Specifically, if we have an antisymmetric function $D(\mathbf{x}_{0\uparrow}, \dots, \mathbf{x}_{N\uparrow}, \mathbf{x}_{0\downarrow}, \dots, \mathbf{x}_{N\downarrow})$, and if i_α is the i -th particle of species type α , then the backflow transformation works by making the coordinate transformation $\mathbf{x}_{i_\alpha} \rightarrow \mathbf{x}'_{i_\alpha}$, and evaluating D at these new “quasiparticle” coordinates. QMCPACK currently supports quasiparticle transformations given by:

$$\mathbf{x}'_{i_\alpha} = \mathbf{x}_{i_\alpha} + \sum_{\alpha \leq \beta} \sum_{i_\alpha \neq j_\beta} \eta^{\alpha\beta}(|\mathbf{x}_{i_\alpha} - \mathbf{x}_{j_\beta}|)(\mathbf{x}_{i_\alpha} - \mathbf{x}_{j_\beta}) \quad (7.2)$$

Here, $\eta^{\alpha\beta}(|\mathbf{x}_{i_\alpha} - \mathbf{x}_{j_\beta}|)$ is a radially symmetric back flow transformation between species α and β . In QMCPACK, particle i_α is known as the “target” particle and j_β is known as the “source”. The main types of transformations we’ll talk about are so called one-body terms, which are between an electron and an ion $\eta^{eI}(|\mathbf{x}_{i_e} - \mathbf{x}_{j_I}|)$, and two-body terms. Two body terms are distinguished as those between like and opposite spin electrons: $\eta^{e(\uparrow)e(\uparrow)}(|\mathbf{x}_{i_e(\uparrow)} - \mathbf{x}_{j_e(\uparrow)}|)$ and $\eta^{e(\uparrow)e(\downarrow)}(|\mathbf{x}_{i_e(\uparrow)} - \mathbf{x}_{j_e(\downarrow)}|)$. Henceforth, we will assume that $\eta^{e(\uparrow)e(\uparrow)} = \eta^{e(\downarrow)e(\downarrow)}$.

In the following, I will explain how to describe general terms like Eq. 7.2 in a QMCPACK XML file. For specificity, I will consider a particle set consisting of H and He (in that order). This ordering will be important when we build the XML file, so you can find this out either through your specific declaration of `iparticleset`, by looking at the hdf5 file in the case of plane waves, or by looking at the qmcpack output file in the section labelled “Summary of QMC systems”.

7.5.1 Input Specifications

All backflow declarations occur within a single `<backflow> ... </backflow>` block. Backflow transformations occur in `<transformation>` blocks, and have the following input parameters

Transformation element				
name	datatype	values	defaults	description
name	text		(required)	Unique name for this Jastrow function
type	text	“e-I”	(required)	Define a one-body backflow transformation.
		“e-e”		Define a two-body backflow transformation.
function	text	Bspline	(required)	B-spline type transformation. (No other types supported)
source	text			“e” if two-body, ion particle set if one-body.

Just like one and two-body jastrows, parameterization of the backflow transformations are specified within the `<transformation>` blocks by `<correlation>` blocks. Please refer to 7.3.2 for

more information.

7.5.2 Example Use Case

Having specified the general form, we present a general example of one-body and two-body backflow transformations in a hydrogen-helium mixture. The H and He ions have independent backflow transformations, as do the like and unlike-spin two-body terms. One caveat is in order: ionic backflow transformations must be listed in the order that they appear in the particle set. If in our example, He is listed first, and H is listed second, the following example would be correct. However, switching backflow declaration to H first, then He, will result in an error. Outside of this, declaration of one-body blocks and two-body blocks aren't sensitive to ordering.

```
<backflow>
<!--The One-Body term with independent e-He and e-H terms. IN THAT ORDER -->
<transformation name="eIonB" type="e-I" function="Bspline" source="ion0">
  <correlation cusp="0.0" size="8" type="shortrange" init="no" elementType="He" rcut="3.0">
    <coefficients id="eHeC" type="Array" optimize="yes">
      0 0 0 0 0 0 0 0
    </coefficients>
  </correlation>
  <correlation cusp="0.0" size="8" type="shortrange" init="no" elementType="H" rcut="3.0">
    <coefficients id="eHC" type="Array" optimize="yes">
      0 0 0 0 0 0 0 0
    </coefficients>
  </correlation>
</transformation>

<!--The Two-Body Term with Like and Unlike Spins -->
<transformation name="eeB" type="e-e" function="Bspline" >
  <correlation cusp="0.0" size="7" type="shortrange" init="no" speciesA="u" speciesB="u" rcut="
  1.2">
    <coefficients id="uuB1" type="Array" optimize="yes">
      0 0 0 0 0 0 0
    </coefficients>
  </correlation>
  <correlation cusp="0.0" size="7" type="shortrange" init="no" speciesA="d" speciesB="u" rcut="
  1.2">
    <coefficients id="udB1" type="Array" optimize="yes">
      0 0 0 0 0 0 0
    </coefficients>
  </correlation>
</transformation>
</backflow>
```

Currently, backflow only works with single-slater determinant wavefunctions. When a backflow transformation has been declared, it should be placed within the `<determinantset>` block, but outside of the `<slaterdeterminant>` blocks, like so:

```
<determinantset ... >
  <!--basis set declarations go here, if there are any -->

  <backflow>
    <transformation ...>
      <!--Here's where discussed one and two-body terms are defined -->
    </transformation>
  </backflow>

  <slaterdeterminant>
    <!--Usual determinant definitions -->
  </slaterdeterminant>
</determinantset>
```

7.5.3 Additional Information

- **Optimization:** Optimizable backflow transformation parameters are notoriously nonlinear, and so optimizing backflow wavefunctions can sometimes be difficult. We direct the reader to our provided backflow tutorials for more information.

Chapter 8

Hamiltonian and Observables

QMCPACK is capable of the simultaneous measurement of the Hamiltonian and many other quantum operators. The Hamiltonian attains a special status among the available operators (also referred to as observables) because it ultimately generates all available information regarding the quantum system. This is evident from an algorithmic standpoint as well since the Hamiltonian (embodied in the the projector) generates the imaginary time dynamics of the walkers in DMC and RMC.

This section covers how the Hamiltonian can be specified, component by component, by the user in the XML format native to QMCPACK . It also covers the input structure of statistical estimators corresponding to quantum observables such as the density, the static structure factor, and forces.

8.1 The Hamiltonian

The many-body Hamiltonian in Hartree units is given by

$$\hat{H} = - \sum_i \frac{1}{2m_i} \nabla_i^2 + \sum_i v^{ext}(r_i) + \sum_{i < j} v^{qq}(r_i, r_j) + \sum_{i\ell} v^{qc}(r_i, r_\ell) + \sum_{\ell < m} v^{cc}(r_\ell, r_m). \quad (8.1)$$

Here, the sums indexed by i/j are over quantum particles, while ℓ/m are reserved for classical particles. Often the quantum particles are electrons and the classical particles are ions, though QMCPACK is not limited in this way. The mass of each quantum particle is denoted m_i , $v^{qq}/v^{qc}/v^{cc}$ are pair potentials between quantum-quantum/quantum-classical/classical-classical particles, and v^{ext} denotes a purely external potential.

QMCPACK is designed modularly so that any potential can be supported with minimal additions to the code base. Potentials currently supported include Coulomb interactions in open and periodic boundary conditions, the modified periodic coulomb (MPC) potential, non-local pseudopo-

tentials, helium pair potentials, and various model potentials such as hard sphere, gaussian, and modified Poschl-Teller.

Reference information and examples for the `<hamiltonian/>` XML element is provided below. Detailed descriptions of the input for individual potentials is given in the sections that follow.

hamiltonian element				
parent elements:	simulation, qmcsystem			
child elements:	pairpot extpot estimator constant(deprecated)			
attributes				
name	datatype	values	default	description
name/id ^o	text	<i>anything</i>	h0	Unique id for this Hamiltonian instance
type ^o	text		generic	<i>No current function</i>
role ^o	text	primary/extra	extra	Designate as primary Hamiltonian or not
source ^o	text	particle set .name	i	Identify classical particle set
target ^o	text	particle set .name	e	Identify quantum particle set
default ^o	boolean	yes/no	yes	Include kinetic energy term implicitly

Additional information:

- **target:** Must be set to the name of the quantum particleset. The default value is typically sufficient. In normal usage, no other attributes are provided.

Listing 8.1: All electron Hamiltonian XML element.

```
<hamiltonian target="e">
  <pairpot name="ElecElec" type="coulomb" source="e" target="e"/>
  <pairpot name="ElecIon" type="coulomb" source="i" target="e"/>
  <pairpot name="IonIon" type="coulomb" source="i" target="i"/>
</hamiltonian>
```

Listing 8.2: Pseudopotential Hamiltonian XML element.

```
<hamiltonian target="e">
  <pairpot name="ElecElec" type="coulomb" source="e" target="e"/>
  <pairpot name="PseudoPot" type="pseudo" source="i" wavefunction="psi0" format="xml">
    <pseudo elementType="Li" href="Li.xml"/>
    <pseudo elementType="H" href="H.xml"/>
  </pairpot>
  <pairpot name="IonIon" type="coulomb" source="i" target="i"/>
</hamiltonian>
```

8.2 Pair potentials

Many pair potentials are supported. Though only the most commonly used pair potentials are covered in detail in this section, all currently available potentials are listed briefly below. If a potential you desire is not covered below, or is not present at all, feel free to contact the developers.

pairpot factory element				
parent elements:	hamiltonian			
type selector:	type attribute			
type options:	coulomb	Coulomb/Ewald potential		
	pseudo	Semilocal pseudopotential		
	mpc	Modified Periodic Coulomb interaction/correction		
	cpp	Core polarization potential		
	numerical/*num*	Numerical radial potential		
	skpot	Unknown		
	vhxc	Exchange correlation potential (external)		
	jellium	Atom-centered spherical jellium potential		
	hardsphere	Hard sphere potential		
	gaussian	Gaussian potential		
	modpostel	Modified Poschl-Teller potential		
	huse	Huse quintic potential		
	modInsKE	Model insulator kinetic energy		
	oscillatory	Unknown		
	LJP_smoothed	Helium pair potential		
	HeSAPT_smoothed	Helium pair potential		
	HFDHE2_Moroni1995	Helium pair potential		
	HFDHE2	Helium pair potential		
	eHe	Helium-electron pair potential		
shared attributes:				
name	datatype	values	default	description
type ^r	text	See above	0	Select pairpot type
name ^r	text	anything	any	Unique name for this pairpot
source ^r	text	particleset.name	hamiltonian.target	Identify interacting particles
target ^r	text	particleset.name	hamiltonian.target	Identify interacting particles
units ^o	text		hartree	No current function

Additional information:

- **type:** Used to select the desired pair potential. Must be selected from the list of type options above.

- **name:** A unique name used to identify this pair potential. Block averaged output data will appear under this name in `scalar.dat` and/or `stat.h5` files.
- **source/target:** These specify the particles involved in a pair interaction. If an interaction is between classical (e.g. ions) and quantum (e.g. electrons), **source/target** should be the name of the classical/quantum particleset.
- Only `coulomb`, `pseudo`, `mpc` are described in detail below. The older or less used types (`cpp`, `numerical`, `jellium`, `hardsphere`, `gaussian`, `huse`, `modpostel`, `oscillatory`, `skpot`, `vhxc`, `modInsKE`, `LJP_smoothed`, `HeSAPT_smoothed`, `HFDHE2_Moroni1995`, `eHe`, `HFDHE2`) are not covered.
- Available only if `QMC_BUILD_LEVEL>2` and `QMC_CUDA` is not defined: `hardsphere`, `gaussian`, `huse`, `modpostel`, `oscillatory`, `skpot`.
- Available only if `OHMMS_DIM==3`: `mpc`, `vhxc`, `pseudo`.
- Available only if `OHMMS_DIM==3` and `QMC_BUILD_LEVEL>2` and `QMC_CUDA` is not defined: `cpp`, `LJP_smoothed`, `HeSAPT_smoothed`, `HFDHE2_Moroni1995`, `eHe`, `jellium`, `HFDHE2`, `modInsKE`.

8.2.1 Coulomb potentials

The bare Coulomb potential is used in open boundary conditions:

$$V_c^{open} = \sum_{i<j} \frac{q_i q_j}{|r_i - r_j|} \quad (8.2)$$

When periodic boundary conditions are selected, Ewald summation is used automatically:

$$V_c^{pbc} = \sum_{i<j} \frac{q_i q_j}{|r_i - r_j|} + \frac{1}{2} \sum_{L \neq 0} \sum_{i,j} \frac{q_i q_j}{|r_i - r_j + L|} \quad (8.3)$$

The sum indexed by L is over all non-zero simulation cell lattice vectors. In practice, the Ewald sum is broken into short and long ranged parts in a manner optimized for efficiency (see Ref. [?]) for details.

For information on how to set the boundary conditions, consult Sec. 6.1.

pairpot type=coulomb element				
parent elements:		hamiltonian		
child elements:		None		
attributes				
name	datatype	values	default	description
type ^r	text	coulomb		Must be coulomb
name/id ^r	text	anything	ElecElec	Unique name for interaction
source ^r	text	particleset.name	hamiltonian.target	Identify interacting particles
target ^r	text	particleset.name	hamiltonian.target	Identify interacting particles
pbc ^o	boolean	yes/no	yes	Use Ewald summation
physical ^o	boolean	yes/no	yes	Hamiltonian(yes)/observable(no)
forces	boolean	yes/no	no	Deprecated

Additional information

- **type/source/target** See description for the generic **pairpot** factory element above.
- **name:** Traditional user-specified names for electron-electron, electron-ion, and ion-ion terms are **ElecElec**, **ElecIon**, and **IonIon**, respectively. While any choice can be used, the data analysis tools expect to find columns in ***.scalar.dat** with these names.
- **pbc:** Ewald summation will not be performed if **simulationcell.bconds== n n n**, regardless of the value of **pbc**. Similarly, the **pbc** attribute can only be used to turn off Ewald summation if **simulationcell.bconds!= n n n**. The default value is recommended.
- **physical:** If **physical==yes**, this pair potential is included in the Hamiltonian and will factor into the **LocalEnergy** reported by QMCPACK and also in the DMC branching weight. If **physical==no**, then the pair potential is treated as a passive observable but not as part of the Hamiltonian itself. As such it does not contribute to the outputted **LocalEnergy**. Regardless of the value of **physical** output data will appear in **scalar.dat** in a column headed by **name**.

Listing 8.3: XML element for Coulomb interaction between electrons.

```
<pairpot name="ElecElec" type="coulomb" source="e" target="e"/>
```

Listing 8.4: XML element for Coulomb interaction between electrons and ions (all-electron only).

```
<pairpot name="ElecIon" type="coulomb" source="i" target="e"/>
```

Listing 8.5: XML element for Coulomb interaction between ions.

```
<pairpot name="IonIon" type="coulomb" source="i" target="i"/>
```


8.2.2 Pseudopotentials

QMCPACK supports pseudopotentials in semilocal form, which is local in the radial coordinate and non-local in angular coordinates. When all angular momentum channels above a certain threshold (ℓ_{max}) are well approximated by the same potential ($V_{\bar{\ell}} \equiv V_{loc}$), the pseudopotential separates into a fully local channel and an angularly-nonlocal component:

$$V^{PP} = \sum_{ij} \left(V_{\bar{\ell}}(|r_i - \tilde{r}_j|) + \sum_{\ell \neq \bar{\ell}}^{\ell_{max}} \sum_{m=-\ell}^{\ell} |Y_{\ell m}\rangle [V_{\ell}(|r_i - \tilde{r}_j|) - V_{\bar{\ell}}(|r_i - \tilde{r}_j|)] \langle Y_{\ell m}| \right) \quad (8.4)$$

Here the electron/ion index is i/j and only one type of ion is shown for simplicity.

Evaluation of the localized pseudopotential energy $\Psi_T^{-1} V^{PP} \Psi_T$ requires additional angular integrals. These integrals are evaluated on a randomly shifted angular grid. The size of this grid is determined by ℓ_{max} . See Ref. [?] for further detail.

QMCPACK uses the FSAtom pseudopotential file format associated with the “Free Software Project for Atomic-scale Simulations” initiated in 2002 (see <http://www.tddft.org/fsatom/manifest.php> for general information). The FSAtom format uses XML for structured data. Files in this format do not use a specific identifying file extension; they are simply suffixed with “.xml”. The tabular data format of CASINO is also supported.

pairpot type=pseudo element				
parent elements:		hamiltonian		
child elements:		pseudo		
attributes				
name	datatype	values	default	description
type ^r	text	pseudo		Must be pseudo
name/id ^r	text	anything	PseudoPot	No current function
source ^r	text	particleset.name	i	Ion particleset name
target ^r	text	particleset.name	hamiltonian.target	Electron particleset name
pbcs ^o	boolean	yes/no	yes*	Use Ewald summation
forces	boolean	yes/no	no	Deprecated
wavefunction ^r	text	wavefunction.name	invalid	Identify wavefunction
format ^r	text	xml/table	table	Select file format

Additional information:

- **type/source/target** See description for the generic **pairpot** factory element above.
- **name:** Ignored. Instead default names will be present in ***scalar.dat** output files when pseudopotentials are used. The field **LocalECP** refers to the local part of the pseudopotential.

If non-local channels are present, a `NonLocalECP` field will be added that contains the non-local energy summed over all angular momentum channels.

- **pbcs:** Ewald summation will not be performed if `simulationcell.bconds== n n n`, regardless of the value of `pbcs`. Similarly, the `pbcs` attribute can only be used to turn off Ewald summation if `simulationcell.bconds!= n n n`.
- **format:** If `format==table`, QMCPACK looks for `*.psf` files containing pseudopotential data in a tabular format. The files must be named after the ionic species provided in `particleset` (e.g. `Li.psf` and `H.psf`). If `format==xml`, additional `pseudo` child XML elements must be provided (see below). These elements specify individual file names and formats (both the FSAAtom XML and CASINO tabular data formats are supported).

Listing 8.6: XML element for pseudopotential electron-ion interaction (psf files).

```
<pairpot name="PseudoPot" type="pseudo" source="i" wavefunction="psi0" format="psf"/>
```

Listing 8.7: XML element for pseudopotential electron-ion interaction (xml files).

```
<pairpot name="PseudoPot" type="pseudo" source="i" wavefunction="psi0" format="xml">
  <pseudo elementType="Li" href="Li.xml"/>
  <pseudo elementType="H" href="H.xml"/>
</pairpot>
```

Details of `<pseudo/>` input elements are given below. It is possible to include (or construct) a full pseudopotential directly in the input file without providing an external file via `href`. The full XML format for pseudopotentials is not yet covered.

pseudo element				
parent elements:	pairpot type=pseudo			
child elements:	header local grid			
attributes				
name	datatype	values	default	description
elementType/symbol ^r	text	group.name	none	Identify ionic species
href ^r	text	filepath	none	Pseudopotential file path
format ^r	text	xml/casino	xml	Specify file format
cutoff ^o	real			Non-local cutoff radius
lmax ^o	integer			Largest angular momentum
nrule ^o	integer			Integration grid order

Listing 8.8: XML element for pseudopotential of single ionic species.

```
<pseudo elementType="Li" href="Li.xml"/>
```

8.2.3 Modified periodic Coulomb interaction/correction

The modified periodic Coulomb (MPC) interaction is an alternative to direct Ewald summation. The MPC corrects the exchange correlation hole to more closely match its thermodynamic limit. Because of this, the MPC exhibits smaller finite size errors than the bare Ewald interaction, though a few alternative and competitive finite size correction schemes now exist. The MPC is itself often used just as a finite size correction in postprocessing (set `physical=false` in the input).

pairpot type=mpc element				
parent elements:	hamiltonian			
child elements:	None			
attributes				
name	datatype	values	default	description
type ^r	text	mpc		Must be mpc
name/id ^r	text	anything	MPC	Unique name for interaction
source ^r	text	particleset.name	hamiltonian.target	Identify interacting particles
target ^r	text	particleset.name	hamiltonian.target	Identify interacting particles
physical ^o	boolean	yes/no	no	Hamiltonian(yes)/observable(no)
cutoff	real	> 0	30.0	Kinetic energy cutoff

Remarks

- **physical**: Typically set to **no**, meaning the standard Ewald interaction will be used during sampling and MPC will be measured as an observable for finite-size post correction. If **physical** is **yes**, the MPC interaction will be used during sampling. In this case an electron-electron Coulomb **pairpot** element should not be supplied.
- Developer note: Currently the **name** attribute for the mpc interaction is ignored. The name is always reset to MPC.

Listing 8.9: Modified periodic coulomb for finite size post-correction.

```
<pairpot type="MPC" name="MPC" source="e" target="e" ecut="60.0" physical="no"/>
```

8.3 General estimators

A broad range of estimators for physical observables are available in QMCPACK . The sections below contain input details for the total number density (**density**), number density resolved by particle spin (**spindensity**), spherically averaged pair correlation function (**gofr**), static structure factor (**sk**), energy density (**energydensity**), one body reduced density matrix (**dm1b**), $S(k)$

based kinetic energy correction (**chiesa**), forward walking (**ForwardWalking**), and force (**Force**) estimators. Other estimators are not yet covered.

When an `<estimator/>` element appears in `<hamiltonian/>`, it is evaluated for all applicable chained QMC runs (*e.g.* VMC→DMC→DMC). Estimators are generally not accumulated during wavefunction optimization sections. If an `<estimator/>` element is instead provided in a particular `<qmc/>` element, that estimator is only evaluated for that specific section (*e.g.* during VMC only).

estimator factory element				
parent elements:	hamiltonian, qmc			
type selector:	type attribute			
type options:	density	Density on a grid		
	spindensity	Spin density on a grid		
	gofr	Pair correlation function (quantum species)		
	sk	Static structure factor		
	structurefactor	Species resolved structure factor		
	momentum	Momentum distribution		
	energydensity	Energy density on uniform or Voronoi grid		
	dm1b	One body density matrix in arbitrary basis		
	chiesa	Chiesa-Ceperley-Martin-Holzmam kinetic energy correction		
	Force	Family of “force” estimators (see 8.5)		
	ForwardWalking	Forward walking values for existing estimators		
	orbitalimages	Create image files for orbitals, then exit		
	flux	Checks sampling of kinetic energy		
	localmoment	Atomic spin polarization within cutoff radius		
	numberfluctuations	Spatial number fluctuations		
	HFDHE2	Helium pressure		
	NearestNeighbors	Trace nearest neighbor indices		
	Kinetic	No current function		
	Pressure	No current function		
	ZeroVarObs	No current function		
	DMCCorrection	No current function		
shared attributes:				
name	datatype	values	default	description
type ^r	text	See above	0	Select estimator type
name ^r	text	anything	any	Unique name for this estimator

8.3.1 Chiesa-Ceperley-Martin-Holzmamm kinetic energy correction

This estimator calculates a finite size correction to the kinetic energy following the formalism laid out in Ref. [?]. The total energy can be corrected for finite size effects by using this estimator in conjunction with the MPC correction.

estimator type=chiesa element				
parent elements:	hamiltonian, qmc			
child elements:	None			
attributes				
name	datatype	values	default	description
type ^r	text	chiesa		Must be chiesa
name ^o	text	anything	KEcorr	Always reset to KEcorr
source ^o	text	particleset.name	e	Identify quantum particles
psi ^o	text	wavefunction.name	psi0	Identify wavefunction

Listing 8.10: “Chiesa” kinetic energy finite size post-correction.

```
<estimator name="KEcorr" type="chiesa" source="e" psi="psi0"/>
```

8.3.2 Density estimator

The particle number density operator is given by

$$\hat{n}_r = \sum_i \delta(r - r_i) \quad (8.5)$$

The **density** estimator accumulates the number density on a uniform histogram grid over the simulation cell. The value obtained for a grid cell c with volume Ω_c is then the average number of particles in that cell:

$$n_c = \int dR |\Psi|^2 \int_{\Omega_c} dr \sum_i \delta(r - r_i) \quad (8.6)$$

estimator type=density element				
parent elements:	hamiltonian, qmc			
child elements:	None			
attributes				
name	datatype	values	default	description
type ^r	text	density		Must be density
name ^r	text	anything	any	Unique name for estimator
delta ^o	real array(3)	$0 \leq v_i \leq 1$	0.1 0.1 0.1	Grid cell spacing, unit coords
x_min ^o	real	> 0	0	Grid starting point in x (Bohr)
x_max ^o	real	> 0	lattice[0]	Grid ending point in x (Bohr)
y_min ^o	real	> 0	0	Grid starting point in y (Bohr)
y_max ^o	real	> 0	lattice[1]	Grid ending point in y (Bohr)
z_min ^o	real	> 0	0	Grid starting point in z (Bohr)
z_max ^o	real	> 0	lattice[2]	Grid ending point in z (Bohr)
potential ^o	boolean	yes/no	no	Accumulate local potential, <i>Deprecated</i>
debug ^o	boolean	yes/no	no	<i>No current function</i>

Additional information:

- **name:** The name provided will be used as a label in the `stat.h5` file for the blocked output data. Post-processing tools expect `name="Density"`.
- **delta:** This sets the histogram grid size used to accumulate the density: `delta="0.1 0.1 0.05"` → $10 \times 10 \times 20$ grid, `delta="0.01 0.01 0.01"` → $100 \times 100 \times 100$ grid. The density grid is written to a `stat.h5` file at the end of each Monte Carlo block. If you request many *blocks* in a `<qmc/>` element, or select a large grid, the resulting `stat.h5` file may be many GB in size.
- ***_min/*_max:** Can be used to select a subset of the simulation cell for the density histogram grid. For example if a (cubic) simulation cell is 20 Bohr on a side, setting `*_min=5.0` and `*_max=15.0` will result in a density histogram grid spanning a $10 \times 10 \times 10$ Bohr cube about the center of the box. Use of `x_min`, `x_max`, `y_min`, `y_max`, `z_min`, `z_max` is only appropriate for orthorhombic simulation cells with open boundary conditions.
- When open boundary conditions are used, a `<simulationcell/>` element must be explicitly provided as the first sub-element of `<qmcsystem/>` for the density estimator to work. In this case the molecule should be centered around the middle of the simulation cell ($L/2$) and not the origin (0 since the space within the cell, and hence the density grid, is defined from 0 to L).

Listing 8.11: Density estimator (uniform grid).

```
<estimator name="Density" type="density" delta="0.05 0.05 0.05"/>
```

8.3.3 Spin density estimator

The spin density is similar to the total density described above. In this case, the sum over particles is performed independently for each spin component.

estimator type=spindensity element					
parent elements:		hamiltonian, qmc			
child elements:		None			
attributes					
name	datatype	values	default	description	
type ^r	text	spindensity		Must be spindensity	
name ^r	text	anything	any	Unique name for estimator	
report ^o	boolean	yes/no	no	Write setup details to stdout	
parameters					
name	datatype	values	default	description	
grid ^o	integer array(3)	$v_i > 0$		Grid cell count	
dr ^o	real array(3)	$v_i > 0$		Grid cell spacing (Bohr)	
cell ^o	real array(3,3)	anything		Volume grid exists in	
corner ^o	real array(3)	anything		Volume corner location	
center ^o	real array(3)	anything		Volume center/origin location	
voronoi ^o	text	particleset.name		Under development	
test_moves ^o	integer	≥ 0	0	Test estimator with random moves	

Additional information:

- **name:** The name provided will be used as a label in the `stat.h5` file for the blocked output data. Post-processing tools expect `name="SpinDensity"`.
- **grid:** Sets the dimension of the histogram grid. Input like `<parameter name="grid"> 40 40 40 </parameter>` requests a $40 \times 40 \times 40$ grid. The shape of individual grid cells is commensurate with the supercell shape.
- **dr:** Real space dimensions of grid cell edges (Bohr units). Input like `<parameter name="dr"> 0.5 0.5 0.5 </parameter>` in a supercell with axes of length 10 Bohr each (but of arbitrary shape) will produce a $20 \times 20 \times 20$ grid. The inputted `dr` values are rounded to produce an integer number of grid cells along each supercell axis. Either `grid` or `dr` must be provided, but not both.

- **cell**: When **cell** is provided, a user defined grid volume is used instead of the global supercell. This must be provided if open boundary conditions are used. Additionally, if **cell** is provided, the user must specify where the volume is located in space in addition to its size/shape (**cell**) using either the **corner** or **center** parameters.
- **corner**: The grid volume is defined as $corner + \sum_{d=1}^3 u_d cell_d$ with $0 < u_d < 1$ (“cell” refers to either the supercell or user provided cell).
- **center**: The grid volume is defined as $center + \sum_{d=1}^3 u_d cell_d$ with $-1/2 < u_d < 1/2$ (“cell” refers to either the supercell or user provided cell). **corner/center** can be used to shift the grid even if **cell** is not specified. Simultaneous use of **corner** and **center** will cause QMCPACK to abort.

Listing 8.12: Spin density estimator (uniform grid).

```
<estimator type="spindensity" name="SpinDensity" report="yes">
  <parameter name="grid"> 40 40 40 </parameter>
</estimator>
```

Listing 8.13: Spin density estimator (uniform grid centered about origin).

```
<estimator type="spindensity" name="SpinDensity" report="yes">
  <parameter name="grid">
    20 20 20
  </parameter>
  <parameter name="center">
    0.0 0.0 0.0
  </parameter>
  <parameter name="cell">
    10.0 0.0 0.0
    0.0 10.0 0.0
    0.0 0.0 10.0
  </parameter>
</estimator>
```

8.3.4 Pair correlation function, $g(r)$

The functional form of the species resolved radial pair correlation function operator is

$$g_{ss'}(r) = \frac{V}{4\pi r^2 N_s N_{s'}} \sum_{i_s=1}^{N_s} \sum_{j_{s'}=1}^{N_{s'}} \delta(r - |r_{i_s} - r_{j_{s'}}|). \quad (8.7)$$

Here N_s is the number of particles of species s and V is the supercell volume. If $s = s'$, then the sum is restricted so that $i_s \neq j_s$.

In QMCPACK, an estimate of $g_{ss'}(r)$ is obtained as a radial histogram with a set of N_b uniform bins of width δr . This can be expressed analytically as

$$\tilde{g}_{ss'}(r) = \frac{V}{4\pi r^2 N_s N_{s'}} \sum_{i=1}^{N_s} \sum_{j=1}^{N_{s'}} \frac{1}{\delta r} \int_{r-\delta r/2}^{r+\delta r/2} dr' \delta(r' - |r_{si} - r_{s'j}|), \quad (8.8)$$

where the radial coordinate r is restricted to reside at the bin centers, $\delta r/2, 3\delta r/2, 5\delta r/2, \dots$

estimator type=gofr element				
parent elements:	hamiltonian, qmc			
child elements:	None			
attributes				
name	datatype	values	default	description
type ^r	text	gofr		Must be gofr
name ^o	text	anything	any	No current function
num_bin ^r	integer	> 1	20	# of histogram bins
rmax ^o	real	> 0	10	Histogram extent (Bohr)
dr ^o	real	> 0	0.5	No current function
debug ^o	boolean	yes/no	no	No current function
target ^o	text	particleset.name	hamiltonian.target	Quantum particles
source/sources ^o	text array	particleset.name	hamiltonian.target	Classical particles

Additional information:

- **num_bin:** The number of bins in each species pair radial histogram.
- **rmax:** Maximum pair distance included in the histogram. The uniform bin width is $\delta r = \text{rmax}/\text{num_bin}$. If periodic boundary conditions are used for any dimension of the simulation cell, then the default value of **rmax** is the simulation cell radius instead of 10 Bohr. For open boundary conditions the volume (V) used is 1.0 Bohr³.
- **source/sources:** If unspecified, only pair correlations between each species of quantum particle will be measured. For each classical particleset specified by **source/sources**, additional pair correlations between each quantum and classical species will be measured. Typically there is only one classical particleset (*e.g.* **source**="ion0"), but there can be several in principle (*e.g.* **sources**="ion0 ion1 ion2").
- **target:** The default value is the preferred usage (*i.e.* **target** does not need to be provided).
- Data is outputted to the **stat.h5** for each QMC sub-run. Individual histograms are named according to the quantum particleset and index of the pair. For example, if the quantum

particleset is named "e" and there are two species (up and down electrons, say), then there will be three sets of histogram data in each `stat.h5` file named `gofr_e_0_0`, `gofr_e_0_1`, and `gofr_e_1_1` for up-up, up-down, and down-down correlations, respectively.

Listing 8.14: Pair correlation function estimator element.

```
<estimator type="gofr" name="gofr" num_bin="200" rmax="3.0" />
```

Listing 8.15: Pair correlation function estimator element with additional electron-ion correlations.

```
<estimator type="gofr" name="gofr" num_bin="200" rmax="3.0" source="ion0" />
```

8.3.5 Static structure factor, $S(k)$

Let $\rho_{\mathbf{k}}^e = \sum_j e^{i\mathbf{k} \cdot \mathbf{r}_j^e}$ be the Fourier space electron density, with \mathbf{r}_j^e being the coordinate of the j -th electron. \mathbf{k} is a wavevector commensurate with the simulation cell. QMCPACK allows the user to accumulate the static electron structure factor $S(\mathbf{k})$ at all commensurate \mathbf{k} such that $|\mathbf{k}| \leq (LR_DIM_CUTOFF)r_c$. N^e is the number of electrons, `LR_DIM_CUTOFF` is the optimized breakup parameter, and r_c is the Wigner-Seitz radius. It is defined as follows:

$$S(\mathbf{k}) = \frac{1}{N^e} \langle \rho_{-\mathbf{k}}^e \rho_{\mathbf{k}}^e \rangle \quad (8.9)$$

estimator type=sk element				
parent elements:	hamiltonian, qmc			
child elements:	None			
attributes				
name	datatype	values	default	description
type ^r	text	sk		Must be sk
name ^r	text	<i>anything</i>	any	Unique name for estimator
hdf5 ^o	boolean	yes/no	no	Output to <code>stat.h5</code> (yes) or <code>scalar.dat</code> (no)

Additional information:

- **name:** Unique name for estimator instance. A data structure of the same name will appear in `stat.h5` output files.
- **hdf5:** If `hdf5==yes` output data for $S(k)$ is directed to the `stat.h5` file (recommended usage). If `hdf5==no`, the data is instead routed to the `scalar.dat` file resulting in many columns of data with headings prefixed by **name** and postfixed by the k-point index (*e.g.* `sk_0 sk_1 ...sk_1037 ...`).

- This estimator only works in periodic boundary conditions. Its presence in the input file is ignored otherwise.
- This is not a species resolved structure factor. Additionally, for \mathbf{k} vectors commensurate with the unit cell, $S(\mathbf{k})$ will include contributions from the static electronic density, thus meaning it won't accurately measure the electron-electron density response.

Listing 8.16: Static structure factor estimator element.

```
<estimator type="sk" name="sk" hdf5="yes"/>
```

8.3.6 Energy density estimator

An energy density operator, $\hat{\mathcal{E}}_r$, satisfies

$$\int dr \hat{\mathcal{E}}_r = \hat{H}, \quad (8.10)$$

where the integral is over all space and \hat{H} is the Hamiltonian. In QMCPACK, the energy density is split into kinetic and potential components

$$\hat{\mathcal{E}}_r = \hat{\mathcal{T}}_r + \hat{\mathcal{V}}_r \quad (8.11)$$

with each component given by

$$\begin{aligned} \hat{\mathcal{T}}_r &= \frac{1}{2} \sum_i \delta(r - r_i) \hat{p}_i^2 \\ \hat{\mathcal{V}}_r &= \sum_{i < j} \frac{\delta(r - r_i) + \delta(r - r_j)}{2} \hat{v}^{ee}(r_i, r_j) + \sum_{i\ell} \frac{\delta(r - r_i) + \delta(r - \tilde{r}_\ell)}{2} \hat{v}^{eI}(r_i, \tilde{r}_\ell) \\ &\quad + \sum_{\ell < m} \frac{\delta(r - \tilde{r}_\ell) + \delta(r - \tilde{r}_m)}{2} \hat{v}^{II}(\tilde{r}_\ell, \tilde{r}_m). \end{aligned} \quad (8.12)$$

Here r_i and \tilde{r}_ℓ represent electron and ion positions, respectively, \hat{p}_i is a single electron momentum operator, and $\hat{v}^{ee}(r_i, r_j)$, $\hat{v}^{eI}(r_i, \tilde{r}_\ell)$, $\hat{v}^{II}(\tilde{r}_\ell, \tilde{r}_m)$ are the electron-electron, electron-ion, and ion-ion pair potential operators (including non-local pseudopotentials, if present). This form of the energy density is size consistent, *i.e.* the partially integrated energy density operators of well separated atoms gives the isolated Hamiltonians of the respective atoms. For periodic systems with twist averaged boundary conditions, the energy density is formally correct only for either a set of supercell k-points that correspond to real valued wavefunctions, or a k-point set that has inversion symmetry around a k-point having a real valued wavefunction. For more information about the energy density, see Ref. [?].

In QMCPACK, the energy density can be accumulated on piecewise uniform three dimensional grids in generalized cartesian, cylindrical, or spherical coordinates. The energy density integrated within Voronoi volumes centered on ion positions is also available. The total particle number density is also accumulated on the same grids by the energy density estimator for convenience so that related quantities, such as the regional energy per particle, can be computed easily.

estimator type=EnergyDensity element				
parent elements:	hamiltonian, qmc			
child elements:	reference_points, spacegrid			
attributes				
name	datatype	values	default	description
type ^r	text	EnergyDensity		Must be EnergyDensity
name ^r	text	anything		Unique name for estimator
dynamic ^r	text	particleset.name		Identify electrons
static ^o	text	particleset.name		Identify ions

Additional information:

- **name:** Must be unique. A dataset with blocked statistical data for the energy density will appear in the `stat.h5` files labeled as `name`.

Listing 8.17: Energy density estimator accumulated on a 20x10x10 grid over the simulation cell.

```
<estimator type="EnergyDensity" name="EDcell" dynamic="e" static="ion0">
  <spacegrid coord="cartesian">
    <origin p1="zero"/>
    <axis p1="a1" scale=".5" label="x" grid="-1 (.05) 1"/>
    <axis p1="a2" scale=".5" label="y" grid="-1 (.1) 1"/>
    <axis p1="a3" scale=".5" label="z" grid="-1 (.1) 1"/>
  </spacegrid>
</estimator>
```

Listing 8.18: Energy density estimator accumulated within spheres of radius 6.9 Bohr centered on the first and second atoms in the `ion0` particleset.

```
<estimator type="EnergyDensity" name="EDatom" dynamic="e" static="ion0">
  <reference_points coord="cartesian">
    r1 1 0 0
    r2 0 1 0
    r3 0 0 1
  </reference_points>
  <spacegrid coord="spherical">
    <origin p1="ion01"/>
    <axis p1="r1" scale="6.9" label="r" grid="0 1"/>
  </spacegrid>
</estimator>
```

```

    <axis p1="r2" scale="6.9" label="phi" grid="0 1"/>
    <axis p1="r3" scale="6.9" label="theta" grid="0 1"/>
  </spacegrid>
  <spacegrid coord="spherical">
    <origin p1="ion02"/>
    <axis p1="r1" scale="6.9" label="r" grid="0 1"/>
    <axis p1="r2" scale="6.9" label="phi" grid="0 1"/>
    <axis p1="r3" scale="6.9" label="theta" grid="0 1"/>
  </spacegrid>
</estimator>

```

Listing 8.19: Energy density estimator accumulated within Voronoi polyhedra centered on the ions.

```

<estimator type="EnergyDensity" name="EDvoronoi" dynamic="e" static="ion0">
  <spacegrid coord="voronoi"/>
</estimator>

```

The `<reference_points/>` element provides a set of points for later use in specifying the origin and coordinate axes needed to construct a spatial histogramming grid. Several reference points on the surface of the simulation cell (see Table 8.1) as well as the positions of the ions (see the `energydensity.static` attribute) are made available by default. The reference points can be used, for example, to construct a cylindrical grid along a bond with the origin on the bond center.

reference_points element				
parent elements:	estimator type=EnergyDensity			
child elements:	None			
attributes				
name	datatype	values	default	description
coord ^r	text	cartesian/cell		Specify coordinate system
body text				
The body text is a line formatted list of points with labels				

Additional information

- **coord:** If `coord=cartesian`, labeled points are in cartesian (x,y,z) format in units of Bohr. If `coord=cell`, then labeled points are in units of the simulation cell axes.
- **body text:** The list of points provided in the body text are line formatted, with four entries per line (*label coor1 coor2 coor3*). A set of points referenced to the simulation cell are available by default (see table 8.1). If `energydensity.static` is provided, the location of each individual ion is also available (*e.g.* if `energydensity.static=ion0`, then the location of the first atom is available with label `ion01`, the second with `ion02`, etc.). All points can be used by label when constructing spatial histogramming grids (see the `spacegrid` element below) used to collect energy densities.

label	point	description
zero	0 0 0	Cell center
a1	a_1	Cell axis 1
a2	a_2	Cell axis 2
a3	a_3	Cell axis 3
f1p	$a_1/2$	Cell face 1+
f1m	$-a_1/2$	Cell face 1-
f2p	$a_2/2$	Cell face 2+
f2m	$-a_2/2$	Cell face 2-
f3p	$a_3/2$	Cell face 3+
f3m	$-a_3/2$	Cell face 3-
cppp	$(a_1 + a_2 + a_3)/2$	Cell corner +,+,+
cppm	$(a_1 + a_2 - a_3)/2$	Cell corner +,+,-
cpmp	$(a_1 - a_2 + a_3)/2$	Cell corner +,-,+
cmpp	$(-a_1 + a_2 + a_3)/2$	Cell corner -,+,+
cpmm	$(a_1 - a_2 - a_3)/2$	Cell corner +,-,-
cmpm	$(-a_1 + a_2 - a_3)/2$	Cell corner -,+,-
cmmp	$(-a_1 - a_2 + a_3)/2$	Cell corner -,-,+
cmmm	$(-a_1 - a_2 - a_3)/2$	Cell corner -,-,-

Table 8.1: Reference points available by default. The vectors a_1 , a_2 , and a_3 refer to the simulation cell axes. The representation of the cell is centered around **zero**.

The `<spacegrid/>` element is used to specify a spatial histogramming grid for the energy density. Grids are constructed based on a set of, potentially non-orthogonal, user provided coordinate axes. The axes are based on information available from `reference.points`. Voronoi grids are based only on nearest neighbor distances between electrons and ions. Any number of space grids can be provided to a single energy density estimator.

spacegrid element				
parent elements:	estimator type=EnergyDensity			
child elements:	origin, axis			
attributes				
name	datatype	values	default	description
coord ^r	text	cartesian cylindrical spherical voronoi		Specify coordinate system

The `<origin/>` element gives the location of the origin for a non-Voronoi grid.

origin element				
parent elements:	spacegrid			
child elements:	None			
attributes				
name	datatype	values	default	description
p1 ^r	text	reference_point.label		Select end point
p2 ^o	text	reference_point.label		Select end point
fraction ^o	real		0	Interpolation fraction

Additional information:

- **p1/p2/fraction:** The location of the origin is set to $p1 + fraction * (p2 - p1)$. If only **p1** is provided, the origin is at **p1**.

The `<axis/>` element represents a coordinate axis used to construct the, possibly curved, coordinate system for the histogramming grid. Three `<axis/>` elements must be provided to a non-Voronoi `<spacegrid/>` element.

axis element				
parent elements:	spacegrid			
child elements:	None			
attributes				
name	datatype	values	default	description
label ^r	text	See below		Axis/dimension label
grid ^r	text		"0 1"	Grid ranges/intervals
p1 ^r	text	reference_point.label		Select end point
p2 ^o	text	reference_point.label		Select end point
scale ^o	real			Interpolation fraction

Additional information:

- **label:** The allowed set of axis labels depends on the coordinate system (*i.e.* `spacegrid.coord`). Labels are `x/y/z` for `coord=cartesian`, `r/phi/z` for `coord=cylindrical`, `r/phi/theta` for `coord=spherical`.
- **p1/p2/scale:** The axis vector is set to `p1+scale*(p2-p1)`. If only `p1` is provided, the axis vector is `p1`.
- **grid:** Specifies the histogram grid along the direction specified by `label`. The allowed grid points fall in the range `[-1,1]` for `label=x/y/z` or `[0,1]` for `r/phi/theta`. A grid of 10 evenly spaced points between 0 and 1 can be requested equivalently by `grid="0 (0.1) 1"` or `grid="0 (10) 1"`. Piecewise uniform grids covering portions of the range are supported, *e.g.* `grid="-0.7 (10) 0.0 (20) 0.5"`.
- Note that `grid` specifies the histogram grid along the (curved) coordinate given by `label`. The axis specified by `p1/p2/scale` does not correspond one-to-one with `label` unless `label=x/y/z`, but the full set of axes provided define the (sheared) space on top of which the curved (*e.g.* spherical) coordinate system is built.

8.3.7 One body density matrix

The N-body density matrix in DMC is $\hat{\rho}_N = |\Psi_T\rangle\langle\Psi_{FN}|$ (for VMC, substitute Ψ_T for Ψ_{FN}). The one body reduced density matrix (1RDM) is obtained by tracing out all particle coordinates but one:

$$\hat{n}_1 = \sum_n Tr_{R_n} |\Psi_T\rangle\langle\Psi_{FN}| \quad (8.13)$$

In the formula above, the sum is over all electron indices and $Tr_{R_n}(\ast) \equiv \int dR_n \langle R_n | \ast | R_n \rangle$ with $R_n = [r_1, \dots, r_{n-1}, r_{n+1}, \dots, r_N]$. When the sum is restricted over spin up or down electrons, one

obtains a density matrix for each spin species. The 1RDM computed by QMCPACK is partitioned in this way.

In real space, the matrix elements of the 1RDM are

$$n_1(r, r') = \langle r | \hat{n}_1 | r' \rangle = \sum_n \int dR_n \Psi_T(r, R_n) \Psi_{FN}^*(r', R_n) \quad (8.14)$$

A more efficient and compact representation of the 1RDM is obtained by expanding in the single particle orbitals obtained from a Hartree-Fock or DFT calculation, $\{\phi_i\}$:

$$\begin{aligned} n_1(i, j) &= \langle \phi_i | \hat{n}_1 | \phi_j \rangle \\ &= \int dR \Psi_{FN}^*(R) \Psi_T(R) \sum_n \int dr'_n \frac{\Psi_T(r'_n, R_n)}{\Psi_T(r_n, R_n)} \phi_i(r'_n)^* \phi_j(r_n) \end{aligned} \quad (8.15)$$

The integration over r' in Eq. 8.15 is inefficient when one is also interested in obtaining matrices involving energetic quantities, such as the energy density matrix of Ref. [?] or the related (and more well known) Generalized Fock matrix. For this reason, an approximation is introduced as follows:

$$n_1(i, j) \approx \int dR \Psi_{FN}^*(R) \Psi_T(R) \sum_n \int dr'_n \frac{\Psi_T(r'_n, R_n)^*}{\Psi_T(r_n, R_n)^*} \phi_i(r_n)^* \phi_j(r'_n) \quad (8.16)$$

For VMC, FN-DMC, FP-DMC, and RN-DMC the formula above represents an exact sampling of the 1RDM corresponding to $\hat{\rho}_N^\dagger$ (see appendix A of Ref. [?] for more detail).

estimator type=dm1b element				
parent elements:	hamiltonian, qmc			
child elements:	none			
attributes				
name	datatype	values	default	description
type ^r	text	dm1b		Must be dm1b
name ^r	text	anything		Unique name for estimator
parameters				
name	datatype	values	default	description
basis ^r	text array	sposet.name(s)		Orbital basis
integrator ^o	text	uniform_grid uniform density	uniform_grid	Integration method
evaluator ^o	text	loop/matrix	loop	Evaluation method
scale ^o	real	0 < scale < 1	1.0	Scale integration cell
center ^o	real array(3)	any point		Center of cell
points ^o	integer	> 0	10	Grid points in each dim
samples ^o	integer	> 0	10	MC samples
warmup ^o	integer	> 0	30	MC warmup
timestep ^o	real	> 0	0.5	MC time step
use_drift ^o	boolean	yes/no	no	Use drift in VMC
check_overlap ^o	boolean	yes/no	no	Print overlap matrix
check_derivatives ^o	boolean	yes/no	no	Check density derivatives
acceptance_ratio ^o	boolean	yes/no	no	Print accept ratio
rstats ^o	boolean	yes/no	no	Print spatial stats
normalized ^o	boolean	yes/no	no	basis comes norm'ed
energy_matrix ^o	boolean	yes/no	no	Energy density matrix

Additional information

- **name:** Density matrix results appear in `stat.h5` files labeled according to **name**.
- **basis:** List of `sposet.name`'s. The total set of orbitals contained in all `sposet`'s comprises the basis (subspace) the one body density matrix is projected onto. This set of orbitals generally includes many virtual orbitals that are not occupied in a single reference Slater determinant.
- **integrator:** This selects the method used to perform the additional single particle integration. Options are `uniform_grid` (uniform grid of points over the cell), `uniform` (uniform random sampling over the cell), and `density` (Metropolis sampling of approximate density: $\sum_{b \in \text{basis}} |\phi_b|^2$, not well tested, please check results carefully!). Depending on the integrator selected, different subsets of the other input parameters are active.
- **evaluator:** Select for-loop or matrix multiply implementations. Matrix is preferred for speed. Both implementations should give the same results, but please check as this has not been exhaustively tested.
- **scale:** Resize the simulation cell by scale for use as an integration volume (active for `integrator=uniform/uniform_grid`).
- **center:** Translate the integration volume to center at this point (active for `integrator=uniform/uniform_grid`). If **center** is not provided, the scaled simulation cell is used as is.
- **points:** The number of grid points in each dimension for `integrator=uniform_grid`. For example, `points=10` results in a uniform 10x10x10 grid over the cell.
- **samples:** Sets the number of Monte Carlo samples collected each step (active for `integrator=uniform/density`).
- **warmup:** Number of warmup Metropolis steps at the start of the run, prior to data collection (active for `integrator=density`).
- **timestep:** Drift-diffusion timestep used in Metropolis sampling (active for `integrator=density`).
- **use_drift:** Enable drift in Metropolis sampling (active for `integrator=density`).
- **check_overlap:** Print the overlap matrix (computed via simple Riemann sums) to the log and then abort. Note that subsequent analysis based on the 1RDM is simplest if the input orbitals are orthogonal.
- **check_derivatives:** Print analytic and numerical derivatives of the approximate (sampled) density for several sample points, then abort.
- **acceptance_ratio:** Print the acceptance ratio of the density sampling to the log each step.

- **rstats**: Print statistical information about the spatial motion of the sampled points to the log each step.
- **normalized**: Declare whether the inputted orbitals are normalized or not. If **normalized=no**, direct Riemann integration over a 200x200x200 grid will be used to compute the normalizations prior to use.
- **energy_matrix**: Also accumulate the one body reduced energy density matrix and write it to **stat.h5**. This matrix is not covered in any detail here; the interested reader is referred to Ref. [?].

Listing 8.20: One body density matrix with uniform grid integration.

```
<estimator type="dm1b" name="DensityMatrices">
  <parameter name="basis"      >  spo_u spo_uv  </parameter>
  <parameter name="evaluator"  >  matrix      </parameter>
  <parameter name="integrator"  >  uniform_grid </parameter>
  <parameter name="points"     >  4             </parameter>
  <parameter name="scale"      >  1.0          </parameter>
  <parameter name="center"     >  0 0 0       </parameter>
</estimator>
```

Listing 8.21: One body density matrix with uniform sampling.

```
<estimator type="dm1b" name="DensityMatrices">
  <parameter name="basis"      >  spo_u spo_uv  </parameter>
  <parameter name="evaluator"  >  matrix      </parameter>
  <parameter name="integrator"  >  uniform      </parameter>
  <parameter name="samples"    >  64          </parameter>
  <parameter name="scale"      >  1.0          </parameter>
  <parameter name="center"     >  0 0 0       </parameter>
</estimator>
```

Listing 8.22: One body density matrix with density sampling.

```
<estimator type="dm1b" name="DensityMatrices">
  <parameter name="basis"      >  spo_u spo_uv  </parameter>
  <parameter name="evaluator"  >  matrix      </parameter>
  <parameter name="integrator"  >  density      </parameter>
  <parameter name="samples"    >  64          </parameter>
  <parameter name="timestep"   >  0.5         </parameter>
  <parameter name="use_drift"   >  no           </parameter>
</estimator>
```

Listing 8.23: Example sposet initialization for density matrix use. Occupied and virtual orbital sets are created separately, then joined (**basis="spo_u spo_uv"**).

```

<sposet_builder type="bspline" href="../../dft/pwscf_output/pwscf.pwscf.h5" tilematrix="1 0 0 0 1 0
0 0 1" twistnum="0" meshfactor="1.0" gpu="no" precision="single">
  <sposet type="bspline" name="spo_u" group="0" size="4"/>
  <sposet type="bspline" name="spo_d" group="0" size="2"/>
  <sposet type="bspline" name="spo_uv" group="0" index_min="4" index_max="10"/>
</sposet_builder>

```

Listing 8.24: Example sposet initialization for density matrix use. Density matrix orbital basis created separately (`basis="dm_basis"`).

```

<sposet_builder type="bspline" href="../../dft/pwscf_output/pwscf.pwscf.h5" tilematrix="1 0 0 0 1 0
0 0 1" twistnum="0" meshfactor="1.0" gpu="no" precision="single">
  <sposet type="bspline" name="spo_u" group="0" size="4"/>
  <sposet type="bspline" name="spo_d" group="0" size="2"/>
  <sposet type="bspline" name="dm_basis" size="50" spindataset="0"/>
</sposet_builder>

```

8.4 Forward Walking Estimators

Forward walking is a method by which one can sample the pure fixed-node distribution $\langle \Phi_0 | \Phi_0 \rangle$. Specifically, one multiplies each walker’s DMC mixed estimate for the observable \mathcal{O} , $\frac{\mathcal{O}(\mathbf{R})\Psi_T(\mathbf{R})}{\Psi_T(\mathbf{R})}$, by the weighting factor $\frac{\Phi_0(\mathbf{R})}{\Psi_T(\mathbf{R})}$. As it turns out, this weighting factor for any walker \mathbf{R} is proportional to the total number of descendants the walker will have after a sufficiently long projection time β .

To forward walk on an observable, one declares a generic forward walking estimator within a `<hamiltonian>` block, and then specifies the observables to forward walk on and forward walking parameters. Here is a summary.

estimator type=ForwardWalking element				
parent elements: <code>hamiltonian</code> , <code>qmc</code>				
child elements: <code>Observable</code>				
attributes				
name	datatype	values	default	description
<code>type</code> ^r	text	ForwardWalking		Must be “ForwardWalking”
<code>name</code> ^r	text	<i>anything</i>	any	Unique name for estimator

Additional information:

- **Cost:** Due to having to store histories of observables up to `max` time-steps, one should multiply the memory cost of storing the non-forward walked observables variables by `max`. Not an issue for things like the potential energy, but can be prohibitive for observables like density, forces, etc.

Observable element				
parent elements:	estimator, hamiltonian, qmc			
child elements:	None			
attributes				
name	datatype	values	default	description
name ^r	text	anything	any	Registered name of existing estimator on which to forward
max ^r	integer	> 0		The maximum projection time in steps (max= β/τ).
frequency ^r	text	≥ 1		Dump data only for every frequency-th to scalar.dat file

- **Naming Convention:** Forward walked observables are automatically named `FWE_name.i`, where `i` is the forward walked expectation value at time step `i`, and `name` is whatever name appears in the `<Observable>` block. This is also how it will appear in the `scalar.dat` file.

In the following example case, QMCPACK forward walks on the potential energy for 300 time steps, and dumps the forward walked value at every time step.

Listing 8.25: Forward walking estimator element.

```
<estimator name="fw" type="ForwardWalking">
  <Observable name="LocalPotential" max="300" frequency="1"/>
  <!--- Additional Observable blocks go here -->
</estimator>
```

8.5 “Force” estimators

QMCPACK supports force estimation by use of the Chiesa-Ceperly-Zhang (CCZ) estimator. Currently, open and periodic boundary conditions are supported, but for all-electron calculations only.

Without loss of generality, the CCZ estimator for the z-component of the force on an ion centered at the origin is given by the following expression:

$$F_z = -Z \sum_{i=1}^{N_e} \frac{z_i}{r_i^3} [\theta(r_i - \mathcal{R}) + \theta(\mathcal{R} - r_i) \sum_{\ell=1}^M c_\ell r_i^\ell] \quad (8.17)$$

Z is the ionic charge, M is the degree of the smoothing polynomial, \mathcal{R} is a real-space cutoff of the sphere within which the bare-force estimator is smoothed, and c_ℓ are predetermined coefficients. These coefficients are chosen to minimize the weighted mean square error between the bare force estimate and the s-wave filtered estimator. Specifically,

$$\chi^2 = \int_0^{\mathcal{R}} dr r^m [f_z(r) - \tilde{f}_z(r)]^2 \quad (8.18)$$

Here, m is the weighting exponent, $f_z(r)$ is the unfiltered radial force density for the z force component, and $\tilde{f}_z(r)$ smoothed polynomial function for the same force density. The reader is invited to refer to the original paper for a more thorough explanation of the methodology, but with the notation in hand, QMCPACK takes the following parameters.

estimator type=Force element				
parent elements:	hamiltonian, qmc			
child elements:	parameter			
attributes				
name	datatype	values	default	description
mode ^o	text	See above	bare	Select estimator type
type ^r	text	Force		Must be “Force”
name ^o	text	anything	ForceBase	Unique name for this estimator
pb ^c	boolean	yes/no	yes	Using periodic BC’s or not
addionion ^o	boolean	yes/no	no	Add the ion-ion force contribution to output force estim
parameters				
name	datatype	values	default	description
rcut ^o	real	> 0	1.0	Real space cutoff \mathcal{R} in bohr.
nbasis ^o	integer	> 0	2	Degree of smoothing polynomial M
weightexp ^o	integer	> 0	2	χ^2 weighting exponent m .

Additional information:

- **Naming Convention:** The unique identifier **name** is appended with **name_X.Y** in the **scalar.dat** file, where X is the ion ID number, and Y is the component ID (an integer with x=0, y=1, z=2). All force components for all ions are computed and dumped to the **scalar.dat** file.
- **Miscellaneous:** Usually, the default choice of **weightexp** is sufficient. Different combinations of **rcut** and **nbasis** should be tested though to minimize variance and bias. There is of course a tradeoff, with larger **nbasis** and smaller **rcut** leading to smaller biases and larger variances.

The following is an example use case.

```
<estimator name="myforce" type="Force" mode="cep" addionion="yes">
  <parameter name="rcut">0.1</parameter>
  <parameter name="nbasis">4</parameter>
  <parameter name="weightexp">2</parameter>
</estimator>
```

Chapter 9

Quantum Monte Carlo Methods

qmc factory element				
parent elements:	simulation, loop			
type selector:	method attribute			
type options:	vmc	Variational Monte Carlo		
	linear	Wavefunction optimization with linear method		
	dmc	Diffusion Monte Carlo		
	rmc	Reptation Monte Carlo		
shared attributes:				
name	datatype	values	default	description
method	text	listed above	invalid	QMC driver
move	text	pbyp, alle	pbyp	method used to move electrons
gpu	text	yes, no	dep.	use the GPU
trace	text		no	???
checkpoints	integer	-1, 0, n	-1	checkpoint frequency
target	text			???
completed	text			???
append	text	yes, no	yes	???

Additional information:

- **move.** There are two ways implemented to move electrons. The more used method is the particle-by-particle move. In this method, only one electron is moved for acception or rejection. The other method is the all-electron move, namely all the electrons are moved once for testing acception or rejection.
- **gpu.** When the executable is compiled with CUDA, the target computing device can be

chosen by this switch. With a regular CPU only compilation, this option is not effective.

- **checkpoints.** If Checkpoint=“-1” no checkpoint will be done (default setting). If Checkpoint=“0” dump after the completion of a qmc section. If Checkpoint=“n” where n is an integer ≥ 0 , walkers will be dumped into a *.config.h5 file every n block. The config.h5 file will contain the state of a population to continue a run including the random number sequences; the list of what is included in the .config.h5 is: number of walkers, status of the run, branch mode, energy dataset, ratio to accepted moves, ratio to proposed moves, variance dataset, vParam{tau, taueff, E_trial, E_ref, Branch_Max, BranchCutOff, BranchFilter, Sigma, Accepted_Energy, Accepted_Samples}, IParam{warmupSteps, Energy_Update_Interval, Counter, targetwalkers, Maxwalkers, MinWalkers, Branching Interval, Walker coordinates, Random number size, Random number sequence, version of the code.

9.1 Variational Monte Carlo

vmc method					
parameters					
name	datatype	values	default	description	
walkers	integer	> 0	dep.	number of walkers per node	
blocks	integer	≥ 0	1	number of blocks	
steps	integer	≥ 0	1	number of steps per block	
warmupsteps	integer	≥ 0	0	number of steps for warming up	
substeps	integer	≥ 0	1	number of substeps per step	
usedrift	text	yes, no	no	use the algorithm with drift	
timestep	real	> 0	0.1	time step for each electron move	
samples	integer	≥ 0	0	total number of samples	
stepsbetweensamples	integer	> 0	1	period of the sample accumulation	
samplesperthread	integer	≥ 0	0	number of samples per thread	
storeconfigs	integer	all values	0	store configurations	

Additional information:

- **walkers.** The initial default number of walkers is 1 but in the CPU branch this number will be overwritten as the number of OpenMP threads if the user requested number is smaller than the number of threads.
- **blocks.** This parameter is universal for all the method. At the end of each block, all the statistics accumulated in the block is dumped in to files, e.g. scalar.dat.

- **warmupsteps.** Warm-up steps are steps used only for equilibration. All the samples generated by warm-up steps are discarded. In practice, there's no need to use many warm-up steps because we can always discard more statistics when we perform the post-process.
- **substeps.** In a substep, each of the electrons is moved only once by either particle-by-particle or all-electron move. Because the local energy is evaluated not at each substep but at each step, increasing the number of substeps doesn't accumulate more samples. But in order to reduce the correlation between consecutive samples, increasing substeps is a very good option for its cheaper computational cost.
- **usedrift.** The VMC is implemented in two algorithms with or without drift. In the no-drift algorithm, the move of each electron is proposed with a Gaussian distribution. The standard deviation is chosen as the timestep input. In the drift algorithm, electrons are moved by langevin dynamics.
- **timestep.** The meaning of timestep depends on whether the drift is used or not. In general, larger timestep reduces the time correlation but might also reduces the accept ratio. Users are required to check the accept ratio of the calculation and make sure it's larger than 0.9 or between 0.2 and 0.8 with or without the drift.
- **stepsbetweensamples.** Due to the fact that samples generated by consecutive steps might be still correlated. Having stepsbetweensamples larger than 1 reduces that correlation. In practice, using larger substeps is cheaper than using stepsbetweensamples to decorrelate samples.
- **samples.** This is the total amount of samples generated in the current VMC session. This parameter is not important for VMC only calculation but necessary if optimization or DMC follows.

$$\text{samples} = \frac{\text{blocks} \cdot \text{steps} \cdot \text{walkers}}{\text{stepsbetweensamples}} \cdot \text{number of MPI tasks}$$

- **samplesperthread.** This is an alternative way to set the target amount of samples. More useful in the VMC session preparing the population for the following DMC calculation.

$$\text{samplesperthread} = \frac{\text{blocks} \cdot \text{steps}}{\text{stepsbetweensamples}}$$

- **storeconfigs.** If storeconfigs is set to a non-zero value, then electron configurations during the VMC run will be saved to the files.

The following is an example of VMC section.

```
<qmc method="vmc" move="pbyp" gpu="yes">
  <estimator name="LocalEnergy" hdf5="no"/>
```

```
<parameter name="walkers">    256 </parameter>
<parameter name="samples">    2867200 </parameter>
<parameter name="stepsbetweensamples">    1 </parameter>
<parameter name="substeps">    5 </parameter>
<parameter name="warmupSteps">    5 </parameter>
<parameter name="blocks">    70 </parameter>
<parameter name="timestep">    1.0 </parameter>
<parameter name="usedrift">    no </parameter>
</qmc>
```

9.2 Wavefunction Optimization

9.3 Diffusion Monte Carlo

dmc method				
parameters				
name	datatype	values	default	description
targetwalkers	integer	> 0	dep.	number of walkers per node
blocks	integer	≥ 0	1	number of blocks
steps	integer	≥ 0	1	number of steps per block
warmupsteps	integer	≥ 0	0	number of steps for warming up
timestep	real	> 0	0.1	time step for each electron move
checkproperties	integer	≥ 0	100	number of steps between walker update
maxcpusecs	real	≥ 0	3.6e5	maximum allowed walltime in seconds
energyUpdateInterval	integer	≥ 0	0	trial energy update interval
refEnergy	AU	all values	dep.	reference energy
feedback	double	≥ 0	1.0	population feedback on the trial energy
useBareTau	option	yes,no	0	do not use effective time step
warmupByReconfiguration	option	yes,no	0	warm up with a fixed population
sigmaBound	double	≥ 0	10	parameter to cutoff large weights
killnode	string	yes/other	no	kill or reject walkers that cross nodes
reconfiguration	string	yes/pure/other	no	fixed population technique
branchInterval	integer	≥ 0	1	branching interval
substeps	integer	≥ 0	1	branching interval
nonlocalmoves	string	yes/other	no	run with tmoves
scaleweight	string	yes/other	yes	scale weights (CUDA only)
MaxAge	double	≥ 0	10	kill persistent walkers
MaxCopy	double	≥ 0	2	limit population growth
fastgrad	text	yes/other	yes	fast gradients
maxDisplSq	real	all values	-1	maximum particle move
storeconfigs	integer	all values	0	store configurations

Additional information:

- **targetwalkers.** A DMC run can be considered a restart run or a new run. A restart run is considered to be any method block beyond the first one, such as when a DMC method block that follows a VMC block. Alternatively, if the user reads in configurations from disk it is also considered a restart run. In the case of a restart run, the DMC driver will use the configurations from the previous run, and this variable will not be used. For a new run, if the number of walkers is less than the number of threads, then the number of walkers will be set equal to the number of threads.
- **blocks.** Number of blocks run during an DMC method block. A block consists of a number of DMC steps (steps), after which all the statistics accumulated in the block are written to

disk.

- **steps**. Number of diffusion Monte Carlo steps in a block.
- **warmupsteps**. Warm-up steps are steps at the beginning of a DMC run in which the instantaneous average energy is used to update the trial energy. During regular steps, E_{ref} is used.
- **timestep**. The timestep determines the accuracy of the imaginary time propagator. Generally, multiple time steps are used to extrapolate to the infinite time step limit. A good range of timesteps in which to perform time step extrapolation will typically have a minimum of 99% acceptance probability for each step.
- **checkproperties**. When using particle by particle driver, this variable specifies how often to reset all the variables kept in the buffer.
- **maxcpusecs**. The default is 100 hours. Once the specified time has elapsed, the program will finalize the simulation even if not all blocks are completed.
- **energyUpdateInterval**. The default is to update the trial energy at every step. Otherwise the trial energy is updated every **energyUpdateInterval** steps.

$$E_{\text{trial}} = \text{refEnergy} + \text{feedback} \cdot (\ln \text{targetWalkers} - \ln N)$$

where N is the current population.

- **refEnergy**. The default reference energy is taken from the VMC run that precedes the DMC run. This value is updated to the current mean whenever branching happens.
- **feedback**. Variable used to determine how strong to react to population fluctuations when doing population control. See the equation in **energyUpdateInterval** for more details.
- **useBareTau**. The same time step is used whether a move is rejected or not. The default is to use an effective time step when a move is rejected.
- **warmupByReconfiguration**. Warmup DMC is done with a fixed population
- **sigmaBound**. Determine the branch cutoff to limit wild weights based on the sigma and sigmaBound
- **killnode**. When running fixed-node, if a walker attempts to cross a node, the move will normally be rejected. If **killnode** = "yes", then walkers are destroyed when they cross a node.
- **reconfiguration**. If reconfiguration is "yes", then run with a fixed walker population using the reconfiguration technique.

- **branchInterval**. Number of steps between branching. The total number of DMC steps in a block will be `BranchInterval*Steps`.
- **substeps**. Same as `BranchInterval`.
- **nonlocalmoves**. DMC driver for running Hamiltonians with non-local moves. An typical usage is to simulate Hamiltonians with non-local pseudopotentials with T-Moves. Setting this equal to false will impose the locality approximation.
- **scaleweight**. Scaling weight per Umrigar/Nightengale. CUDA only.
- **MaxAge**. Set the weight of a walker to $\min(\text{currentweight}, 0.5)$ after a walker has not moved for `MaxAge` steps. Needed if persistent walkers appear during the course of a run.
- **MaxCopy**. When determining the number of copies of a walker to branch, set the number of copies equal to $\min(\text{Multiplicity}, \text{MaxCopy})$.
- **fastgrad**. Calculates gradients with either the fast version or the full-ratio version.
- **maxDisplSq** . When running a DMC calculation with particle by particle, this sets the maximum displacement allowed for a single particle move. All distance displacements larger than the max is rejected. If initialized to a negative value, it becomes equal to `Lattice(LR/rc)`.
- **sigmaBound** . Determine the branch cutoff to limit wild weights based on the sigma and `sigmaBound`
- **storeconfigs**. If `storeconfigs` is set to a non-zero value, then electron configurations during the DMC run will be saved. This option is disabled for the OpenMP version of DMC.

Listing 9.1: The following is an example of a very simple DMC section.

```
<qmc method="dmc" move="pbyp" target="e">
  <parameter name="blocks">100</parameter>
  <parameter name="steps">400</parameter>
  <parameter name="timestep">0.010</parameter>
  <parameter name="warmupsteps">100</parameter>
</qmc>
```

The time step should be adjusted for each problem individually. Please refer to the theory section on diffusion Monte Carlo.

Listing 9.2: The following is an example of running a simulation that can be restarted .

```
<qmc method="dmc" move="pbyp" checkpoint="0" dumpconfig="5">
  <parameter name="timestep"> 0.004 </parameter>
  <parameter name="blocks"> 100 </parameter>
  <parameter name="steps"> 400 </parameter>
</qmc>
```

The flags `checkpoint` and `dumpconfig` instructs `qmcpack` to output walker configurations. This also works in variational Monte Carlo. This will output an h5 file with the name `"projectid"."run-number".config.h5`. Check that this file exists before attempting a restart. To read in this file for a continuation run, specify the following:

Listing 9.3: Restart (read walkers from previous run)

```
<mcwalkerset fileroot="BH.s002" node="-1" nprocs="1" version="0 6" collected="yes"/>
```

where, BH is the project id and s002 is the calculation number to read in the walkers from the previous run.

Combining VMC and DMC in a single run (and wave function optimization can be combined in this way too) is the standard way in which QMCPACK is typical run. There is no need to run two separate jobs, as method sections can be stacked, and walkers are transfered between them.

Listing 9.4: Combined VMC and DMC run

```
<qmc method="vmc" move="pbyp" target="e">
  <parameter name="blocks">100</parameter>
  <parameter name="steps">4000</parameter>
  <parameter name="warmupsteps">100</parameter>
  <parameter name="samples">1920</parameter>
  <parameter name="walkers">1</parameter>
  <parameter name="timestep">0.5</parameter>
</qmc>
<qmc method="dmc" move="pbyp" target="e">
  <parameter name="blocks">100</parameter>
  <parameter name="steps">400</parameter>
  <parameter name="timestep">0.010</parameter>
  <parameter name="warmupsteps">100</parameter>
</qmc>
<qmc method="dmc" move="pbyp" target="e">
  <parameter name="warmupsteps">500</parameter>
  <parameter name="blocks">50</parameter>
  <parameter name="steps">100</parameter>
  <parameter name="timestep">0.005</parameter>
</qmc>
```

9.4 Reptation Monte Carlo

Like diffusion monte carlo, reptation monte carlo (RMC) is a projector based method, allowing us the ability to sample the fixed-node wavefunction. However, by exploiting the path-integral formulation of Schrödinger's equation, the RMC algorithm can offer some advantages over traditional DMC, such as sampling both the mixed and pure fixed-node distributions in polynomial time, as

well as not having population fluctuations and biases. The current implementation does not work with T-moves.

There are two adjustable parameters that affect the quality of the RMC projection: imaginary projection time β of the sampling path (commonly called a “reptile”), and the Trotter time step τ . β must be chosen to be large enough such that $e^{-\beta\hat{H}}|\Psi_T\rangle \approx |\Phi_0\rangle$ for mixed observables, and $e^{-\frac{\beta}{2}\hat{H}}|\Psi_T\rangle \approx |\Phi_0\rangle$ for pure observables. The reptile is discretized into $M = \beta/\tau$ beads at the cost of an $\mathcal{O}(\tau)$ time-step error for observables arising from the Trotter-Suzuki breakup of the short-time propagator.

The following table lists some of the more practical

vmc method				
parameters				
name	datatype	values	default	description
beta	real	> 0	dep.	reptile projection time β
timestep	real	> 0	0.1	Trotter time step τ for each electron move
beads	int	> 0	1	Number of reptile beads $M = \beta/\tau$
blocks	integer	≥ 0	1	number of blocks
steps	integer	≥ 0	1	number of steps per block
vmcpresteps	integer	≥ 0	0	propagates reptile using VMC for given number of steps
warmupsteps	integer	≥ 0	0	number of steps for warming up
MaxAge	integer	≥ 0	0	force accept for stuck reptile if age exceeds MaxAge.

Additional information:

Because of the sampling differences between DMC ensembles of walkers and RMC reptiles, the RMC block should contain the following estimator declaration to ensure correct sampling:
`<estimator name="RMC" hdf5="no">`.

- **beta or beads?** One can specify one or the other, and from the Trotter time-step, the code will construct an appropriately sized reptile. If both are given, **beta** overrides **beads**.
- **Mixed vs. Pure observables?** For all observables appearing in the `scalar.dat` file in either VMC or DMC, RMC appends the suffix `_m` or `_p` for mixed and pure estimates respectively.
- **Sampling.** For pure estimators, one should check the traces of both pure and mixed estimates. Ergodicity is a known problem in RMC. Because we use the bounce algorithm, it is possible for the reptile to bounce back and forth without changing the electron coordinates of the central beads. This might not easily show up with mixed estimators, since these are accumulated at constantly regrown ends, but pure estimates are accumulated on these central beads, and so can exhibit strong autocorrelations in pure estimate traces.

- **Propagator:** Our implementation of RMC uses Moroni’s DMC link action (symmetrized), with Umrigar’s scaled drift near nodes. In this regard, the propagator is identical to the one QMCPACK uses in DMC.
- **Sampling:** We use Ceperley’s bounce algorithm. MaxAge is used in case the reptile gets stuck, at which point the code forces move acceptance, stops accumulating statistics, and requilibrates the reptile. Very rarely will this be required. For move proposals, we use particle-by-particle VMC a total of N_e times to generate a new all-electron configuration, at which point the action is computed and the move is either accepted or rejected.

Chapter 10

Analysing QMCPACK data

10.1 Using the qmca tool

10.2 Densities and spin-densities

10.3 Energy densities

Chapter 11

Examples

WARNING: THESE EXAMPLES ARE NOT CONVERGED! YOU MUST CONVERGE PARAMETERS (SIMULATION CELL SIZE, JASTROW PARAMETER NUMBER/CUTOFF, TWIST NUMBER, DMC TIME STEP, DFT PLANE WAVE CUTOFF, DFT K-POINT MESH, ETC.) FOR REAL CALCULATIONS!

The following examples should run in serial on a modern workstation in a few hours.

11.1 Using Nexus

11.1.1 H₂O Molecule with Quantum ESPRESSO Orbitals

With BFD pseudopotentials (see Section ??) for O and H in a subdirectory named `pseudopotentials` (both UPF and FSAtom xml formats, named `O.BFD.upf`, `O.BFD.xml`, `H.BFD.upf`, and `H.BFD.xml`, respectively) and the following XYZ file (named `H2O.xyz`)

```
3
O 0.0000000000e+00 0.0000000000e+00 0.0000000000e+00
H 0.0000000000e+00 -1.4308249289e+00 1.1078707576e+00
H 0.0000000000e+00 1.4308249289e+00 1.1078707576e+00
```

in the working directory, a Python script using Nexus to generate the orbitals using Quantum ESPRESSO, then run QMCPACK to optimize the Jastrow and then do DMC for H₂O in a box is:

Listing 11.1: Nexus example for H₂O using Quantum ESPRESSO orbitals and BFD pseudopotentials

```
#!/usr/bin/env python

from nexus import settings, Job, run_project
from nexus import Structure, PhysicalSystem
```

```

from nexus import generate_pwscf
from nexus import generate_pw2qmcpack
from nexus import generate_qmcpack,vmc,loop,linear,dmc

# General Settings (Directories For I/O, Machine Type, etc.)
settings(
    pseudo_dir      = 'pseudopotentials',
    runs             = 'runs',
    results          = 'results',
    sleep            = 3,
    generate_only     = 0,
    status_only      = 0,
    machine           = 'ws1',
)

# Executables (Indicate Path If Needed)
pwscf               = 'pw.x'
pw2qmcpack          = 'pw2qmcpack.x'
qmcpack              = 'qmcapp'

# Pseudopotentials
dft_pps              = ['O.BFD.upf','H.BFD.upf']
qmc_pps              = ['O.BFD.xml','H.BFD.xml']

# Job Definitions (MPI Tasks, MP Threading, PBS Queue, Time, etc.)
scf_job              = Job(app=pwscf,serial=True)
p2q_job              = Job(app=pw2qmcpack,serial=True)
opt_job              = Job(threads=4,app=qmcpack,serial=True)
dmc_job              = Job(threads=4,app=qmcpack,serial=True)

# System To Be Simulated
structure = Structure()
structure.read_xyz('H2O.xyz')
structure.bounding_box(
    box              = 'cubic',
    scale             = 1.5
)
structure.add_kmesh(
    kgrid             = (1,1,1),
    kshift             = (0,0,0)
)
H2O_molecule = PhysicalSystem(
    structure          = structure,
    net_charge         = 0,
    net_spin           = 0,
    O                  = 6,
    H                  = 1,
)

```

```

sims = []

# DFT SCF To Generate Converged Density
scf = generate_pwscf(
    identifier      = 'scf',
    path            = '.',
    job             = scf_job,
    input_type      = 'scf',
    system          = H2O_molecule,
    pseudos         = dft_pps,
    ecut            = 50,
    ecutrho         = 400,
    conv_thr        = 1.0e-5,
    mixing_beta     = 0.7,
    mixing_mode     = 'local-TF',
    degauss         = 0.001
)
sims.append(scf)

# Convert DFT Wavefunction Into HDF5 File For QMCPACK
p2q = generate_pw2qmcpack(
    identifier      = 'p2q',
    path            = '.',
    job             = p2q_job,
    write_psiir     = False,
    dependencies    = (scf,'orbitals')
)
sims.append(p2q)

# QMC Optimization Parameters - Coarse Sampling Set
linopt1 = linear(
    energy          = 0.0,
    unreweightedvariance = 1.0,
    reweightedvariance  = 0.0,
    timestep        = 0.4,
    samples          = 8192,
    warmupsteps     = 50,
    blocks           = 64,
    substeps         = 4,
    nonlocalpp       = True,
    usebuffer        = True,
    walkers          = 1,
    minwalkers       = 0.5,
    maxweight        = 1e9,
    usedrift         = True,
    minmethod        = 'quartic',
    beta             = 0.0,
    exp0             = -16,
    bigchange        = 15.0,

```

```

        alloweddifference    = 1e-4,
        stepsize             = 0.2,
        stabilizerscale      = 1.0,
        nstabilizers         = 3
    )

# QMC Optimization Parameters - Finer Sampling Set
linopt2 = linopt1.copy()
linopt2.samples = 16384

# QMC Optimization
opt = generate_qmcpack(
    identifier    = 'opt',
    path         = '.',
    job          = opt_job,
    input_type   = 'basic',
    system       = H2O_molecule,
    twistnum     = 0,
    bconds       = 'nnn',
    pseudos      = qmc_pps,
    jastrows     = [('J1','bspline',8,4),
                    ('J2','bspline',8,4)],
    calculations = [loop(max=4,qmc=linopt1),
                    loop(max=4,qmc=linopt2)],
    dependencies = (p2q,'orbitals')
)
sims.append(opt)

# QMC VMC/DMC With Optimized Jastrow Parameters
qmc = generate_qmcpack(
    identifier    = 'dmc',
    path         = '.',
    job          = dmc_job,
    input_type   = 'basic',
    system       = H2O_molecule,
    pseudos      = qmc_pps,
    bconds       = 'nnn',
    jastrows     = [],
    calculations = [
        vmc(
            walkers           = 1,
            samplesperthread = 64,
            stepsbetweensamples = 1,
            substeps          = 5,
            warmupsteps       = 100,
            blocks             = 1,
            timestep          = 1.0,
            usedrift           = False
        ),

```

```

        dmc(
            minimumtargetwalkers = 128,
            reconfiguration       = 'no',
            warmupsteps          = 100,
            timestep              = 0.005,
            steps                 = 10,
            blocks                = 200,
            nonlocalmoves        = True
        )
    ],
    dependencies = [(p2q,'orbitals'),(opt,'jastrow')]
)
sims.append(qmc)

run_project(sims)

```

11.1.2 LiH Crystal with Quantum ESPRESSO Orbitals

With CASINO-formatted Trail-Needs pseudopotentials (see Section ??) for O and H in a subdirectory named `pseudopotentials` (both UPF and CASINO formats, named `O.TN-DF.upf`, `O.pp.data`, `H.TN-DF.upf`, and `H.pp.data`, respectively), a Python script using Nexus to generate the orbitals using Quantum ESPRESSO, then run QMCPACK to optimize the Jastrow and then do DMC for LiH with periodic boundary conditions is:

Listing 11.2: Nexus example for bulk LiH using Quantum ESPRESSO orbitals and CASINO pseudopotentials

```

#!/usr/bin/env python

from nexus import settings,Job,run_project
from nexus import generate_physical_system
from nexus import generate_pwscf
from nexus import generate_pw2qmcpack
from nexus import generate_qmcpack,vmc,loop,linear,dmc

# General Settings (Directories For I/O, Machine Type, etc.)
settings(
    pseudo_dir    = 'pseudopotentials',
    runs          = 'runs',
    results       = 'results',
    sleep         = 3,
    generate_only  = 1,
    status_only   = 0,
    machine       = 'ws1',
)

# Executables (Indicate Path If Needed)

```

```

pwscf                = 'pw.x'
pw2qmcpack           = 'pw2qmcpack.x'
qmcpack              = 'qmcapp'

# Pseudopotentials
dft_pps              = ['Li.TN-DF.upf', 'H.TN-DF.upf']
qmc_pps              = ['Li.pp.data', 'H.pp.data']

# Job Definitions (MPI Tasks, MP Threading, PBS Queue, Time, etc.)
scf_job              = Job(app=pwscf,serial=True)
nscf_job             = Job(app=pwscf,serial=True)
p2q_job             = Job(app=pw2qmcpack,serial=True)
opt_job             = Job(threads=4,app=qmcpack,serial=True)
dmc_job             = Job(threads=4,app=qmcpack,serial=True)

# System To Be Simulated
rocksalt_LiH = generate_physical_system(
    lattice          = 'cubic',
    cell             = 'primitive',
    centering        = 'F',
    atoms            = ('Li','H'),
    basis            = [[0.0,0.0,0.0],
                       [0.5,0.5,0.5]],
    basis_vectors    = 'conventional',
    constants        = 7.1,
    units            = 'B',
    kgrid            = (17,17,17),
    kshift           = (1,1,1),
    net_charge       = 0,
    net_spin         = 0,
    Li               = 1,
    H               = 1,
)

sims = []

# DFT SCF To Generate Converged Density
scf = generate_pwscf(
    identifier       = 'scf',
    path            = '.',
    job             = scf_job,
    input_type      = 'scf',
    system          = rocksalt_LiH,
    pseudos         = dft_pps,
    ecut            = 450,
    ecutrho         = 1800,
    conv_thr        = 1.0e-10,
    mixing_beta     = 0.7,
)

```



```

sims.append(scf)

# DFT NSCF To Generate Wave Function At Specified K-points
nscf = generate_pwscf(
    identifier      = 'nscf',
    path            = '.',
    job             = nscf_job,
    input_type      = 'nscf',
    system          = rocksalt_LiH,
    pseudos         = dft_pps,
    ecut            = 450,
    ecutrho         = 1800,
    conv_thr        = 1.0e-10,
    mixing_beta     = 0.7,
    kgrid           = (1,1,1),
    kshift          = (0,0,0),
    dependencies    = (scf, 'charge-density')
)
sims.append(nscf)

# Convert DFT Wavefunction Into HDF5 File For QMCPACK
p2q = generate_pw2qmcpack(
    identifier      = 'p2q',
    path            = '.',
    job             = p2q_job,
    write_psiir     = False,
    dependencies    = (nscf, 'orbitals')
)
sims.append(p2q)

# QMC Optimization Parameters - Coarse Sampling Set
linopt1 = linear(
    energy          = 0.0,
    unreweightedvariance = 1.0,
    reweightedvariance  = 0.0,
    timestep        = 0.4,
    samples          = 8192,
    warmupsteps      = 50,
    blocks           = 64,
    substeps         = 4,
    nonlocalpp       = True,
    usebuffer        = True,
    walkers          = 1,
    minwalkers       = 0.5,
    maxweight        = 1e9,
    usedrift         = True,
    minmethod        = 'quartic',
    beta             = 0.0,
    exp0             = -16,

```

```

    bigchange          = 15.0,
    alloweddifference   = 1e-4,
    stepsize           = 0.2,
    stabilizerscale     = 1.0,
    nstabilizers        = 3

# QMC Optimization Parameters - Finer Sampling Set
linopt2 = linopt1.copy()
linopt2.samples = 16384

# QMC Optimization
opt = generate_qmcpack(
    identifier      = 'opt',
    path            = '.',
    job             = opt_job,
    input_type      = 'basic',
    system          = rocksalt_LiH,
    twistnum        = 0,
    bconds          = 'ppp',
    pseudos         = qmc_pps,
    jastrows        = [('J1','bspline',8),
                      ('J2','bspline',8)],
    calculations    = [loop(max=4,qmc=linopt1),
                      loop(max=4,qmc=linopt2)],
    dependencies    = (p2q,'orbitals')
)
pp = opt.input.get('pseudos')
pp.Li.format='casino'
pp.Li['l-local']='s'
pp.Li.nrule=2
pp.Li.lmax=2
pp.Li.cutoff=2.19
pp.H.format='casino'
pp.H['l-local']='s'
pp.H.nrule=2
pp.H.lmax=2
pp.H.cutoff=0.50
sims.append(opt)

# QMC VMC/DMC With Optimized Jastrow Parameters
qmc = generate_qmcpack(
    identifier      = 'dmc',
    path            = '.',
    job             = dmc_job,
    input_type      = 'basic',
    system          = rocksalt_LiH,
    pseudos         = qmc_pps,
    bconds          = 'ppp',
    jastrows        = [],

```

```

calculations    = [
    vmc(
        walkers            = 1,
        samplesperthread   = 64,
        stepsbetweensamples = 1,
        substeps            = 5,
        warmupsteps        = 100,
        blocks              = 1,
        timestep            = 1.0,
        usedrift            = False
    ),
    dmc(
        minimumtargetwalkers = 128,
        reconfiguration      = 'no',
        warmupsteps          = 100,
        timestep             = 0.005,
        steps                = 10,
        blocks                = 200,
        nonlocalmoves        = True
    )
],
dependencies    = [(p2q,'orbitals'),(opt,'jastrow')]
)

pp = qmc.input.get('pseudos')
pp.Li.format='casino'
pp.Li['l-local']='s'
pp.Li.nrule=2
pp.Li.lmax=2
pp.Li.cutoff=2.37
pp.H.format='casino'
pp.H['l-local']='s'
pp
pp.H.lmax=2
pp.H.cutoff=0.50
sims.append(qmc)

run_project(sims)

```

Chapter 12

Lab 1: Monte Carlo Statistical Analysis

12.1 Topics covered in this Lab

This lab focuses on the basics of analyzing data from Monte Carlo (MC) calculations. In this lab, participants will use data from VMC calculations of a simple one-electron system with an analytically soluble system (the ground state of the hydrogen atom) to understand how to interpret a MC situation. Most of these analyses will also carry over to diffusion Monte Carlo (DMC) simulations. Topics covered include:

- averaging Monte Carlo variables
- the statistical error bar of mean values
- effects of autocorrelation and variance on the error bar
- the relationship between Monte Carlo timestep and autocorrelation
- the use of blocking to reduce autocorrelation
- the significance of the acceptance ratio
- the significance of the sample size
- how to determine whether a Monte Carlo run was successful
- the relationship between wavefunction quality and variance
- gauging the efficiency of Monte Carlo runs
- the cost of scaling up to larger system sizes

12.2 Lab directories and files

Lab_1_MC_Analysis/

atom	- H atom VMC calculation
H.s000.scalar.dat	- H atom VMC data
H.xml	- H atom VMC input file
autocorrelation	- varying autocorrelation
H.dat	- data for gnuplot
H.plt	- gnuplot for time step vs. E_L, tau_c
H.s000.scalar.dat	- H atom VMC data: time step = 10
H.s001.scalar.dat	- H atom VMC data: time step = 5
H.s002.scalar.dat	- H atom VMC data: time step = 2
H.s003.scalar.dat	- H atom VMC data: time step = 1
H.s004.scalar.dat	- H atom VMC data: time step = 0.5
H.s005.scalar.dat	- H atom VMC data: time step = 0.2
H.s006.scalar.dat	- H atom VMC data: time step = 0.1
H.s007.scalar.dat	- H atom VMC data: time step = 0.05
H.s008.scalar.dat	- H atom VMC data: time step = 0.02
H.s009.scalar.dat	- H atom VMC data: time step = 0.01
H.s010.scalar.dat	- H atom VMC data: time step = 0.005
H.s011.scalar.dat	- H atom VMC data: time step = 0.002
H.s012.scalar.dat	- H atom VMC data: time step = 0.001
H.s013.scalar.dat	- H atom VMC data: time step = 0.0005
H.s014.scalar.dat	- H atom VMC data: time step = 0.0002
H.s015.scalar.dat	- H atom VMC data: time step = 0.0001
H.xml	- H atom VMC input file
average	- Python scripts for average/std. dev.
average.py	- average five E_L from H atom VMC
stddev2.py	- standard deviation using (E_L)^2
stddev.py	- standard deviation around the mean
basis	- varying basis set for orbitals
H__exact.s000.scalar.dat	- H atom VMC data using STO basis
H_STO-2G.s000.scalar.dat	- H atom VMC data using STO-2G basis
H_STO-3G.s000.scalar.dat	- H atom VMC data using STO-3G basis

H_STO-6G.s000.scalar.dat	- H atom VMC data using STO-6G basis
blocking	- varying block/step ratio
H.dat	- data for gnuplot
H.plt	- gnuplot for N_block vs. E, tau_c
H.s000.scalar.dat	- H atom VMC data 50000:1 blocks:steps
H.s001.scalar.dat	- " " " " 25000:2 blocks:steps
H.s002.scalar.dat	- " " " " 12500:4 blocks:steps
H.s003.scalar.dat	- " " " " 6250: 8 blocks:steps
H.s004.scalar.dat	- " " " " 3125:16 blocks:steps
H.s005.scalar.dat	- " " " " 2500:20 blocks:steps
H.s006.scalar.dat	- " " " " 1250:40 blocks:steps
H.s007.scalar.dat	- " " " " 1000:50 blocks:steps
H.s008.scalar.dat	- " " " " 500:100 blocks:steps
H.s009.scalar.dat	- " " " " 250:200 blocks:steps
H.s010.scalar.dat	- " " " " 125:400 blocks:steps
H.s011.scalar.dat	- " " " " 100:500 blocks:steps
H.s012.scalar.dat	- " " " " 50:1000 blocks:steps
H.s013.scalar.dat	- " " " " 40:1250 blocks:steps
H.s014.scalar.dat	- " " " " 20:2500 blocks:steps
H.s015.scalar.dat	- " " " " 10:5000 blocks:steps
H.xml	- H atom VMC input file
blocks	- varying total number of blocks
H.dat	- data for gnuplot
H.plt	- gnuplot for N_block vs. E
H.s000.scalar.dat	- H atom VMC data 500 blocks
H.s001.scalar.dat	- " " " " 2000 blocks
H.s002.scalar.dat	- " " " " 8000 blocks
H.s003.scalar.dat	- " " " " 32000 blocks
H.s004.scalar.dat	- " " " " 128000 blocks
H.xml	- H atom VMC input file
dimer	- comparing no and simple Jastrow factor
H2_STO___no_jastrow.s000.scalar.dat	- H dimer VMC data without Jastrow
H2_STO_with_jastrow.s000.scalar.dat	- H dimer VMC data with Jastrow
docs	- documentation

Lab_1_MC_Analysis.pdf	- this document
Lab_1_Slides.pdf	- slides presented in the lab
nodes	- varying number of computing nodes
H.dat	- data for gnuplot
H.plt	- gnuplot for N_node vs. E
H.s000.scalar.dat	- H atom VMC data with 32 nodes
H.s001.scalar.dat	- H atom VMC data with 128 nodes
H.s002.scalar.dat	- H atom VMC data with 512 nodes
problematic	- problematic VMC run
H.s000.scalar.dat	- H atom VMC data with a problem
size	- scaling with number of particles
01_____H.s000.scalar.dat	- H atom VMC data
02_____H2.s000.scalar.dat	- H dimer " "
06_____C.s000.scalar.dat	- C atom " "
10_____CH4.s000.scalar.dat	- methane " "
12_____C2.s000.scalar.dat	- C dimer " "
16_____C2H4.s000.scalar.dat	- ethene "
18___CH4CH4.s000.scalar.dat	- methane dimer VMC data
32_C2H4C2H4.s000.scalar.dat	- ethene dimer " "
nelectron_tcpu.dat	- data for gnuplot
Nelectron_tCPU.plt	- gnuplot for N_elec vs. t_CPU

12.3 Atomic units

QMCPACK operates in Hartree atomic units to reduce the number of factors in the Schrödinger equation. Thus, the unit of length is the bohr ($5.291772 \times 10^{-11} \text{ m} = 0.529177 \text{ Å}$); the unit of energy is the hartree ($4.359744 \times 10^{-18} \text{ J} = 27.211385 \text{ eV}$). The energy of the ground state of the hydrogen atom in these units is -0.5 hartrees.

12.4 Reviewing statistics

We will practice taking the average (mean) and standard deviation of some Monte Carlo data by hand to review the basic definitions.

Enter Python's command line by typing **python** [Enter]. You will see a prompt ">>>".

The mean of a data set is given by:

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i \quad (12.1)$$

To calculate the average of five local energies from a MC calculation of the ground state of an electron in the hydrogen atom, input (truncate at the thousandths place if you cannot copy and paste; script versions are also available in the **average** directory):

```
( (-0.45298911858) + (-0.45481953564) + (-0.48066105923) + (-0.47316713469) + (-0.4620473302) )/5.
```

Then, press **[Enter]** to get:

```
>>> ((-0.45298911858) + (-0.45481953564) + (-0.48066105923) +
(-0.47316713469) + (-0.4620473302))/5.
-0.46473683566800006
```

To understand the significance of the mean, we also need the standard deviation around the mean of the data (also called the error bar), given by:

$$\sigma = \sqrt{\frac{1}{N(N-1)} \sum_{i=1}^N (x_i - \bar{x})^2} \quad (12.2)$$

To calculate the standard deviation around the mean (-0.464736835668) of these five data points, put in:

```
( (1./(5.*(5.-1.))) * ( (-0.45298911858-(-0.464736835668))**2 +
(-0.45481953564-(-0.464736835668))**2 + (-0.48066105923-(-0.464736835668))**2 + (-0.47316713469-(-0.464736835668))**2 + (-0.4620473302-(-0.464736835668))**2 ) )**0.5
```

Then, press **[Enter]** to get:

```
>>> ( (1./(5.*(5.-1.))) * ( (-0.45298911858-(-0.464736835668))**2 +
(-0.45481953564-(-0.464736835668))**2 + (-0.48066105923-(-0.464736835668))**2 +
(-0.47316713469-(-0.464736835668))**2 + (-0.4620473302-(-0.464736835668))**2 ) )**0.5
0.0053303187464332066
```


Thus, we might report this data as having a value -0.465 ± 0.005 hartrees. This calculation of the standard deviation assumes that the average for this data is fixed, but we may continually add Monte Carlo samples to the data so it is better to use an estimate of the error bar that does not rely on the overall average. Such an estimate is given by:

$$\tilde{\sigma} = \sqrt{\frac{1}{N-1} \sum_{i=1}^N [(x^2)_i - (x_i)^2]} \quad (12.3)$$

To calculate the standard deviation with this formula, input the following, which includes the square of the local energy calculated with each corresponding local energy:

```
( (1./(5.-1.)) * ( (0.60984565298-(-0.45298911858)**2) +
(0.61641291630-(-0.45481953564)**2) + (1.35860151160-(-0.48066105923)**2) +
(0.78720769003-(-0.47316713469)**2) + (0.56393677687-(-0.46204733302)**2) ) )**0.5
```

and press **[Enter]** to get:

```
>>> ((1./(5.-1.))*((0.60984565298-(-0.45298911858)**2)+
(0.61641291630-(-0.45481953564)**2)+(1.35860151160-(-0.48066105923)**2)+
(0.78720769003-(-0.47316713469)**2)+(0.56393677687-(-0.46204733302)**2))
)**0.5
0.84491636672906634
```

This much larger standard deviation, acknowledging that the mean of this small data set is not the average in the limit of infinite sampling more accurately, reports the value of the local energy as -0.5 ± 0.8 hartrees.

Type **quit()** and press **[Enter]** to exit the Python command line.

12.5 Inspecting Monte Carlo data

QMCPACK outputs data from MC calculations into files ending in `scalar.dat`. Several quantities are calculated and written for each block of Monte Carlo steps in successive columns to the right of the step index.

Change directories to `atom`, and open the file ending in `scalar.dat` with a text editor (e.g., **vi *.scalar.dat** or **emacs *.scalar.dat**). If possible, adjust the terminal so that lines do not wrap. The data will begin as follows (broken into three groups to fit on this page):

#	index	LocalEnergy	LocalEnergy_sq	LocalPotential	...
	0	-4.5298911858e-01	6.0984565298e-01	-1.1708693521e+00	
	1	-4.5481953564e-01	6.1641291630e-01	-1.1863425644e+00	
	2	-4.8066105923e-01	1.3586015116e+00	-1.1766446209e+00	
	3	-4.7316713469e-01	7.8720769003e-01	-1.1799481122e+00	
	4	-4.6204733302e-01	5.6393677687e-01	-1.1619244081e+00	
	5	-4.4313854290e-01	6.0831516179e-01	-1.2064503041e+00	
	6	-4.5064926960e-01	5.9891422196e-01	-1.1521370176e+00	
	7	-4.5687452611e-01	5.8139614676e-01	-1.1423627617e+00	
	8	-4.5018503739e-01	8.4147849706e-01	-1.1842075439e+00	
	9	-4.3862013841e-01	5.5477715836e-01	-1.2080979177e+00	

The first line begins with a #, indicating that this line does not contain MC data but rather the labels of the columns. After a blank line, the remaining lines consist of the MC data. The first column, labeled index, is an integer indicating which block of MC data is on that line. The second column contains the quantity usually of greatest interest from the simulation, the local energy. Since this simulation did not use the exact ground state wave function, it does not produce -0.5 hartrees as the local energy although the value lies within about 10%. The value of the local energy fluctuates from block to block and the closer the trial wave function is to the ground state, the smaller these fluctuations will be. The next column contains an important ingredient in estimating the error in the MC average—the square of the local energy—found by evaluating the square of the Hamiltonian.

...	Kinetic	Coulomb	BlockWeight	...
	7.1788023352e-01	-1.1708693521e+00	1.2800000000e+04	
	7.3152302871e-01	-1.1863425644e+00	1.2800000000e+04	
	6.9598356165e-01	-1.1766446209e+00	1.2800000000e+04	
	7.0678097751e-01	-1.1799481122e+00	1.2800000000e+04	
	6.9987707508e-01	-1.1619244081e+00	1.2800000000e+04	
	7.6331176120e-01	-1.2064503041e+00	1.2800000000e+04	
	7.0148774798e-01	-1.1521370176e+00	1.2800000000e+04	
	6.8548823555e-01	-1.1423627617e+00	1.2800000000e+04	
	7.3402250655e-01	-1.1842075439e+00	1.2800000000e+04	
	7.6947777925e-01	-1.2080979177e+00	1.2800000000e+04	

The fourth column from the left consists of the values of the local potential energy. In this simulation, it is identical to the Coulomb potential (contained in the sixth column) because the one electron in the simulation has only the potential energy coming from its interaction with the nucleus. In many-electron simulations, the local potential energy contains contributions from the electron-electron Coulomb interactions and the nuclear potential or pseudopotential. The fifth column contains the local kinetic energy value for each MC block, obtained from the Laplacian of the wave function. The sixth column shows the local Coulomb interaction energy. The seventh column displays the weight each line of data has in the average (the weights are identical in this simulation).

```
...      BlockCPU      AcceptRatio
        6.0178991748e-03  9.8515625000e-01
        5.8323097461e-03  9.8562500000e-01
        5.8213412744e-03  9.8531250000e-01
        5.8330412549e-03  9.8828125000e-01
        5.8108362256e-03  9.8625000000e-01
        5.8254170264e-03  9.8625000000e-01
        5.8314813086e-03  9.8679687500e-01
        5.8258469971e-03  9.8726562500e-01
        5.8158433545e-03  9.8468750000e-01
        5.7959401123e-03  9.8539062500e-01
```

The eighth column shows the CPU time (in seconds) to calculate the data in that line. The ninth column from the left contains the acceptance ratio (1 being full acceptance) for Monte Carlo steps in that line's data. Other than the block weight, all quantities vary from line to line.

Exit the text editor (**[Esc]** **:q!** **[Enter]** in vi, **[Ctrl]-x** **[Ctrl]-c** in emacs).

12.6 Averaging quantities in the MC data

QMCPACK includes the `qmca` Python tool to average quantities in the `scalar.dat` file (and also the `dmc.dat` file of DMC simulations). Without any flags, `qmca` will output the average of each column with a quantity in the `scalar.dat` file as follows.

Execute `qmca` by **`qmca *.scalar.dat`**, which for this data outputs:

```
H  series 0
```

LocalEnergy	=	-0.45446 +/-	0.00057
Variance	=	0.529 +/-	0.018
Kinetic	=	0.7366 +/-	0.0020
LocalPotential	=	-1.1910 +/-	0.0016
Coulomb	=	-1.1910 +/-	0.0016
LocalEnergy_sq	=	0.736 +/-	0.018
BlockWeight	=	12800.00000000 +/-	0.00000000
BlockCPU	=	0.00582002 +/-	0.00000067
AcceptRatio	=	0.985508 +/-	0.000048
Efficiency	=	0.00000000 +/-	0.00000000

After one blank, **qmca** prints the title of the subsequent data, gleaned from the data file name. In this case, H.s000.scalar.dat became “H series 0”. Everything before the first “.s” will be interpreted as the title, and the number between “.s” and the next “.” will be interpreted as the series number.

The first column under the title is the name of each quantity **qmca** averaged. The column to the right of the equal signs contains the average for the quantity of that line, and the column to the right of the plus-slash-minus is the statistical error bar on the quantity. All quantities calculated from MC simulations have and must be reported with a statistical error bar!

Two new quantities not present in the scalar.dat file are computed by **qmca** from the data—variance and efficiency. We will look at these later in this lab.

To view only one value, **qmca** takes the **-q (quantity)** flag. For example, the output of **qmca -q LocalEnergy *.scalar.dat** in this directory produces a single line of output:

```
H series 0 LocalEnergy = -0.454460 +/- 0.000568
```

Type **qmca -help** to see the list of all quantities and their abbreviations.

12.7 Evaluating MC simulation quality

There are several aspects of a MC simulation to consider in deciding how well it went. Besides the deviation of the average from an expected value (if there is one), the stability of the simulation in its sampling, the autocorrelation between MC steps, the value of the acceptance ratio (accepted steps over total proposed steps), and the variance in the local energy all indicate the quality of a MC simulation. We will look at these one by one.

12.7.1 Tracing MC quantities

Visualizing the evolution of MC quantities over the course of the simulation by a *trace* offers a quick picture of whether the random walk had expected behavior. `qmca` plots traces with the `-t` flag.

Type `qmca -q e -t H.s000.scalar.dat`, which produces a graph of the trace of the local energy:

The solid black line connects the values of the local energy at each MC block (labeled “samples”). The average value is marked with a horizontal, solid red line. One standard deviation above and below the average are marked with horizontal, dashed red lines.

The trace of this run is largely centered around the average with no large-scale oscillations or major shifts, indicating a good quality MC run.

Try tracing the kinetic and potential energies, seeing that their behavior is comparable to the total local energy.

Change to directory `problematic` and type `qmca -q e -t H.s000.scalar.dat` to produce this graph:

Here, the local energy samples cluster around the expected -0.5 hartrees for the first 150 samples or so and then begin to oscillate more wildly and increase erratically toward 0, indicating a poor quality MC run.

Again, trace the kinetic and potential energies in this run and see how their behavior compares to the total local energy.

12.7.2 Blocking away autocorrelation

Autocorrelation occurs when a given MC step biases subsequent MC steps, leading to samples that are not statistically independent. We must take this autocorrelation into account in order to obtain accurate statistics. `qmca` outputs autocorrelation when given the `--sac` flag.

Change to directory `autocorrelation` and type `qmca -q e --sac H.s000.scalar.dat`.

```
H series 0 LocalEnergy = -0.454982 +/- 0.000430 1.0
```

The value after the error bar on the quantity is the autocorrelation (1.0 in this case).

Proposing too small a step in configuration space, the MC *time step*, can lead to autocorrelation since the new samples will be in the neighborhood of previous samples. Type `grep timestep H.xml` to see the varying time step values in this QMCPACK input file (H.xml):

```

<parameter name="timestep">10</parameter>
<parameter name="timestep">5</parameter>
<parameter name="timestep">2</parameter>
<parameter name="timestep">1</parameter>
<parameter name="timestep">0.5</parameter>
<parameter name="timestep">0.2</parameter>
<parameter name="timestep">0.1</parameter>
<parameter name="timestep">0.05</parameter>
<parameter name="timestep">0.02</parameter>
<parameter name="timestep">0.01</parameter>
<parameter name="timestep">0.005</parameter>
<parameter name="timestep">0.002</parameter>
<parameter name="timestep">0.001</parameter>
<parameter name="timestep">0.0005</parameter>
<parameter name="timestep">0.0002</parameter>
<parameter name="timestep">0.0001</parameter>

```

Generally, as the time step decreases, the autocorrelation will increase (caveat: very large time steps will also have increasing autocorrelation). To see this, type **qmca -q e --sac *.scalar.dat** to see the energies and autocorrelation times, then plot with gnuplot by inputting **gnuplot H.plt**:

The error bar also increases with the autocorrelation.

Press **q** [**Enter**] to quit gnuplot.

To get around the bias of autocorrelation, we group the MC steps into blocks, take the average of the data in the steps of each block, and then finally average the averages in all the blocks. QMCPACK outputs the block averages as each line in the scalar.dat file. (For DMC simulations, in addition to the scalar.dat, QMCPACK outputs the quantities at each step to the dmc.dat file, which permits reblocking the data differently from the specification in the input file.)

Change directories to **blocking**. Here we look at the time step of the last data set in the **autocorrelation** directory. Verify this by typing **grep timestep H.xml** to see that all values are set to 0.001. Now to see how we will vary the blocking, type **grep -A1 blocks H.xml**. The parameter “steps” indicates the number of steps per block, and the parameter “blocks” gives the number of blocks. For this comparison, the total number of MC steps (equal to the product of “steps” and “blocks”) is fixed at 50000. Now check the effect of blocking on autocorrelation—type **qmca -q e --sac *.scalar.dat** to see the data and **gnuplot H.plt** to visualize the data:

The greatest number of steps per block produces the smallest autocorrelation time. The larger number of blocks over which to average at small step-per-block number masks the corresponding increase in error bar with increasing autocorrelation.

Press **q** [**Enter**] to quit gnuplot.

12.7.3 Balancing autocorrelation and acceptance ratio

Adjusting the time step value also affects the ratio of accepted steps to proposed steps. Stepping nearby in configuration space implies that the probability distribution is similar and thus more likely to result in an accepted move. Keeping the acceptance ratio high means the algorithm is efficiently exploring configuration space and not sticking at particular configurations. Return to the `autocorrelation` directory. Refresh your memory on the time steps in this set of simulations by `grep timestep H.xml`. Then, type `qmca -q ar *scalar.dat` to see the acceptance ratio as it varies with decreasing time step:

```
H series 0 AcceptRatio = 0.047646 +/- 0.000206
H series 1 AcceptRatio = 0.125361 +/- 0.000308
H series 2 AcceptRatio = 0.328590 +/- 0.000340
H series 3 AcceptRatio = 0.535708 +/- 0.000313
H series 4 AcceptRatio = 0.732537 +/- 0.000234
H series 5 AcceptRatio = 0.903498 +/- 0.000156
H series 6 AcceptRatio = 0.961506 +/- 0.000083
H series 7 AcceptRatio = 0.985499 +/- 0.000051
H series 8 AcceptRatio = 0.996251 +/- 0.000025
H series 9 AcceptRatio = 0.998638 +/- 0.000014
H series 10 AcceptRatio = 0.999515 +/- 0.000009
H series 11 AcceptRatio = 0.999884 +/- 0.000004
H series 12 AcceptRatio = 0.999958 +/- 0.000003
H series 13 AcceptRatio = 0.999986 +/- 0.000002
H series 14 AcceptRatio = 0.999995 +/- 0.000001
H series 15 AcceptRatio = 0.999999 +/- 0.000000
```

By series 8 (time step = 0.02), the acceptance ratio is in excess of 99%.

Considering the increase in autocorrelation and subsequent increase in error bar as time step decreases, it is important to choose a time step that trades off appropriately between acceptance ratio and autocorrelation. In this example, a time step of 0.02 occupies a spot where acceptance ratio is high (99.6%), and autocorrelation is not appreciably larger than the minimum value (1.4 vs. 1.0).

12.7.4 Considering variance

Besides autocorrelation, the dominant contributor to the error bar is the *variance* in the local energy. The variance measures the fluctuations around the average local energy, and, as the fluctuations

go to zero, the wave function reaches an exact eigenstate of the Hamiltonian. `qmca` calculates this from the local energy and local energy squared columns of the `scalar.dat`.

Type `qmca -q v H.s009.scalar.dat` to calculate the variance on the run with time step balancing autocorrelation and acceptance ratio:

```
H series 9 Variance = 0.513570 +/- 0.010589
```

Just as the total energy doesn't tell us much by itself, neither does the variance. However, comparing the ratio of the variance to the energy indicates how the magnitude of the fluctuations compares to the energy itself. Type `qmca -q ev H.s009.scalar.dat` to calculate the energy and variance on the run side by side with the ratio:

	LocalEnergy	Variance	ratio
H series 0	-0.454460 +/- 0.000568	0.529496 +/- 0.018445	1.1651

1.1651 is a very high ratio indicating the square of the fluctuations is on average larger than the value itself. In the next section, we will approach ways to improve the variance that subsequent labs will build upon.

12.8 Reducing statistical error bars

12.8.1 Increasing MC sampling

Increasing the number of MC samples in a data set reduces the error bar as the inverse of the square root of the number of samples. There are two ways to increase the number of MC samples in a simulation: running more samples in parallel and increasing the number of blocks (with fixed number of steps per block, this increases the total number of MC steps).

To see the effect of the running more samples in parallel, change to the directory `nodes`. The series here increases the number of nodes by factors of four from 32 to 128 to 512. Type `qmca -q ev *scalar.dat` and note the change in the error bar on the local energy as the number of nodes. Visualize this with `gnuplot H.plt`:

Increasing the number of blocks, unlike running in parallel, increases the total CPU time of the simulation.

Press `q` [Enter] to quit `gnuplot`.

To see the effect of increasing the block number, change to the directory **blocks**. To see how we will vary the number of blocks, type **grep -A1 blocks H.xml**. The number of steps remains fixed, thus increasing the total number of samples. Visualize the tradeoff by inputting **gnuplot H.plt**:

Press **q** [**Enter**] to quit gnuplot.

12.8.2 Improving the basis set

In all of the above examples, we are using the sum of two gaussian functions (STO-2G) to approximate what should be a simple decaying exponential (STO = Slater-type orbital) for the wave function of the ground state of the hydrogen atom. The sum of multiple copies of a function varying each copy's width and amplitude with coefficients is called a *basis set*. As we add gaussians to the basis set, the approximation improves, the variance goes toward zero and the energy goes to -0.5 hartrees. In nearly every other case, the exact function is unknown, and we add basis functions until the total energy does not change within some threshold.

Change to the directory **basis** and look at the total energy and variance as we change the wave function by typing **qmca -q ev H_***:

		LocalEnergy		Variance		ratio
H_STO-2G	series 0	-0.454460	+/- 0.000568	0.529496	+/- 0.018445	1.1651
H_STO-3G	series 0	-0.465386	+/- 0.000502	0.410491	+/- 0.010051	0.8820
H_STO-6G	series 0	-0.471332	+/- 0.000491	0.213919	+/- 0.012954	0.4539
H__exact	series 0	-0.500000	+/- 0.000000	0.000000	+/- 0.000000	-0.0000

qmca also puts out the ratio of the variance to the local energy in a column to the right of the variance error bar. A typical high quality value for this ratio is lower than 0.1 or so—none of these few-gaussian wave functions satisfy that rule of thumb.

Use qmca to plot the trace of the local energy, kinetic energy, and potential energy of H__exact—the total energy is constantly -0.5 hartree even though the kinetic and potential energies fluctuate from configuration to configuration.

12.8.3 Adding a Jastrow factor

Another route to reducing the variance is the introduction of a Jastrow factor to account for electron-electron correlation (not the statistical autocorrelation of Monte Carlo steps but the physical avoidance that electrons have of one another). To do this, we will switch to the hydrogen dimer with the exact ground state wave function of the atom (STO basis)—this will not be exact for the dimer. The ground state energy of the hydrogen dimer is -1.174 hartrees.

Change directories to `dimer` and put in `qmca -q ev *scalar.dat` to see the result of adding a simple, one-parameter Jastrow to the STO basis for the hydrogen dimer at experimental bond length:

		LocalEnergy	Variance
H2_STO___no_jastrow	series 0	-0.876548 +/- 0.005313	0.473526 +/- 0.014910
H2_STO_with_jastrow	series 0	-0.912763 +/- 0.004470	0.279651 +/- 0.016405

The energy reduces by 0.044 +/- 0.006 hartrees and the variance by 0.19 +/- 0.02. This is still 20% above the ground state energy, and subsequent labs will cover how to improve on this with improved forms of the wave function that capture more of the physics.

12.9 Scaling to larger numbers of electrons

12.9.1 Calculating the efficiency

The inverse of the product of CPU time and the variance measures the *efficiency* of an MC calculation. Use `qmca` to calculate efficiency by typing `qmca -q eff *scalar.dat` to see the efficiency of these two H₂ calculations:

H2_STO___no_jastrow	series 0	Efficiency = 16698.725453 +/- 0.000000
H2_STO_with_jastrow	series 0	Efficiency = 52912.365609 +/- 0.000000

The Jastrow factor increased the efficiency in these calculations by a factor of three, largely through the reduction in variance (check the average block CPU time to verify this claim).

12.9.2 Scaling up

To see how MC scales with increasing particle number, change directories to `size`. Here are the data from runs of increasing number of electrons for H, H₂, C, CH₄, C₂, C₂H₄, (CH₄)₂, and (C₂H₄)₂ using the STO-6G basis set for the orbitals of the Slater determinant. The file names begin with the number of electrons simulated for those data.

Use `qmca -q bc *scalar.dat` to see that the CPU time per block increases with number of electrons in the simulation, then plot the total CPU time of the simulation by `gnuplot Nelectron.tCPU.plt`:

The green pluses represent the CPU time per block at each electron number. The red line is a quadratic fit to those data. For a fixed basis set size, we expect the time to scale quadratically up to 1000s of electrons, at which point a cubic scaling term may become dominant. Knowing the scaling allows you to roughly project the calculation time for a larger number of electrons.

Press **q** [**Enter**] to quit gnuplot.

This isn't the whole story, however. The variance of the energy also increases with a fixed basis set as the number of particles increases at a faster rate than the energy decreases. To see this, type **qmca -q ev *scalar.dat**:

		LocalEnergy	Variance
01_____H	series 0	-0.471352 +/- 0.000493	0.213020 +/- 0.012950
02_____H2	series 0	-0.898875 +/- 0.000998	0.545717 +/- 0.009980
06_____C	series 0	-37.608586 +/- 0.020453	184.322000 +/- 45.481193
10_____CH4	series 0	-38.821513 +/- 0.022740	169.797871 +/- 24.765674
12_____C2	series 0	-72.302390 +/- 0.037691	491.416711 +/- 106.090103
16_____C2H4	series 0	-75.488701 +/- 0.042919	404.218115 +/- 60.196642
18___CH4CH4	series 0	-58.459857 +/- 0.039309	498.579645 +/- 92.480126
32_C2H4C2H4	series 0	-91.567283 +/- 0.048392	632.114026 +/- 69.637760

The increase in variance is not uniform, but the general trend is upward with a fixed wave function form and basis set. Subsequent labs will address how to improve the wave function in order to keep the variance manageable.

Chapter 13

Lab 2: QMC Basics

13.1 Topics covered in this Lab

This lab focuses on the basics of performing quality QMC calculations. As an example participants test an oxygen pseudopotential within DMC by calculating atomic and dimer properties, a common step prior to production runs. Topics covered include:

- converting pseudopotentials into QMCPACK's FSATOM format
- generating orbitals with Quantum Espresso
- converting orbitals into QMCPACK's ESHDF format with pw2qmcpack
- optimizing Jastrow factors with QMCPACK
- removing DMC timestep error via extrapolation
- automating QMC workflows with Nexus
- testing pseudopotentials for accuracy
- (optional) running QMCPACK for a general system of interest

13.2 Lab directories and files

Lab_2_QMC_Basics/

```
docs                                - documentation
  Lab_2_QMC_Basics.pdf              - this document
  Lab_2_Slides.pdf                  - slides presented during the lab
  Nexus.pdf                         - slides on QMCPACK automation (supplementary)

oxygen_atom                        - oxygen atom calculations
  ip_conv.py                        - tool to fit oxygen IP vs timestep
  0.q0.dmc.in.xml                   - neutral 0 DMC input file
  0.q0.dmc.qsub.in                  - " " " submission file
  0.q0.opt.in.xml                   - " " optimization input file
  0.q0.opt.qsub.in                  - " " " submission file
  0.q0.pwscf.h5                     - " " orbitals file
  0.q1.dmc.in.xml                   - charged 0 DMC input file
  0.q1.dmc.qsub.in                  - " " " submission file
  0.q1.opt.in.xml                   - " " optimization input file
  0.q1.opt.qsub.in                  - " " " submission file
  0.q1.pwscf.h5                     - " " orbitals file
  reference                         - directory w/ completed runs
  submit_0_q0_dmc                   - executable to submit neutral DMC
  submit_0_q0_opt                   - " " " " optimization
  submit_0_q1_dmc                   - " " " charged DMC
  submit_0_q1_opt                   - " " " " optimization

oxygen_dimer                       - oxygen dimer calculations
  dimer_fit.py                      - tool to fit dimer binding curve
  0_dimer.py                        - automation script for dimer calculations
  pseudopotentials                  - directory for pseudopotentials
  reference                         - directory w/ completed runs

your_system                        - calculations with your own physical system
  example.py                       - generates input files for your system
  pseudopotentials                  - directory for pseudopotentials
  reference                         - directory w/ completed runs
```

13.3 Obtaining and converting a pseudopotential for oxygen

We will use a potential from the Burkatzki-Filippi-Dolg pseudopotential database. To obtain the pseudopotential, go to <http://www.burkatzki.com/pseudos/index.2.html> and click on the “Select Pseudopotential” button. Next click on oxygen in the periodic table. Click on the empty circle next to “V5Z” (a large gaussian basis set) and click on “Next”. Select the Gamess format and click on “Retrive Potential”. Helpful information about the pseudopotential will be displayed. The desired portion is at the bottom (the last 7 lines). Copy this text into the editor of your choice and save it as `0.BFD.gamess` (be sure to include a newline at the end of the file). To transform the pseudopotential into the fsatom xml format used by QMCPACK, use the `ppconvert` tool:

```
ppconvert --gamess_pot 0.BFD.gamess --s_ref "1s(2)2p(4)" \  
--p_ref "1s(2)2p(4)" --d_ref "1s(2)2p(4)" --xml 0.BFD.xml
```

Observe the notation used to describe the reference valence configuration for this helium-core PP: `1s(2)2p(4)`. The `ppconvert` tool uses the following convention for the valence states: the first *s* state is labeled `1s` (`1s`, `2s`, `3s`, ...), the first *p* state is labeled `2p` (`2p`, `3p`, ...), the first *d* state is labeled `3d` (`3d`, `4d`, ...). Copy the resulting xml file into the `oxygen_atom` directory.

Note: the command to convert the PP into QM Espresso’s UPF format is similar:

```
ppconvert --gamess_pot 0.BFD.gamess --s_ref "1s(2)2p(4)" \  
--p_ref "1s(2)2p(4)" --d_ref "1s(2)2p(4)" --log_grid --upf 0.BFD.upf
```

For reference, the text of `0.BFD.gamess` should be:

```
0-QMC GEN 2 1  
3  
6.00000000 1 9.29793903  
55.78763416 3 8.86492204  
-38.81978498 2 8.62925665  
1  
38.41914135 2 8.71924452
```

The full QMCPACK pseudopotential is also included in `oxygen_atom/reference/0.BFD.xml`.

13.4 Optimization walkthrough with QMCPACK: neutral O atom

The aim of this section is to obtain a trial wavefunction of reasonable quality for the neutral oxygen atom. The first subsection provides background regarding the wavefunction for this system, including the specific form of the Jastrow factors used in QMCPACK. A brief discussion of wavefunction optimization is also given. The second subsection contains the actual walkthrough to follow for the lab.

Enter the `oxygen_atom` directory and copy over the oxygen pseudopotential (`0.BFD.xml`) you downloaded and converted (section ??). Alternatively, the already converted pseudopotential is located in the `oxygen_atom/reference` directory. All files prefixed with “`0.q0`” relate to the neutral oxygen atom.

Open `0.q0.opt.in.xml` with your favorite text editor. This is a QMCPACK input file configured for wavefunction optimization with the linear method. Take a minute to familiarize yourself with the general format and contents of the input file. The major sections are the simulation cell, description of particle species (electrons & ions/atoms), the trial wavefunction (orbitals, Slater determinants, and Jastrow factors), the Hamiltonian, and finally inputs describing the quantum Monte Carlo process (linear optimization in this case). Portions marked with “`<!-- ... -->`” are comments describing these sections. XML is not the easiest to read, but this can be helped by using an editor with color highlighting such as `emacs` or `vi`.

The most important parts to focus on for the purposes of this exercise are the Jastrow factors and the inputs to the linear optimization method. Input specifying the one-body electron-ion Jastrow factor corresponding to eq. ?? is

```
<jastrow type="One-Body" name="J1" function="bspline" source="ion0" print="yes">
  <correlation elementType="0" size="8" rcut="4.5" cusp="0.0">
    <coefficients id="e0" type="Array">
      0 0 0 0 0 0 0 0
    </coefficients>
  </correlation>
</jastrow>
```

The XML describes $U_1^{\uparrow/\downarrow}(r)$ as a B-spline with 8 knots, no cusp at the origin (the oxygen pseudopotential is finite at $r = 0$), and vanishing beyond 4.5 Bohr. The initial guess of zero for each of the 8 knot parameters corresponds to $U_1^{\uparrow/\downarrow}(r) = 0$. The input for the two-body electron-electron Jastrow is similar:

```

<jastrow type="Two-Body" name="J2" function="pade" print="yes">
  <correlation speciesA="u" speciesB="u">
    <var id="uu_b" name="B"> 0.6 </var>
  </correlation>
  <correlation speciesA="u" speciesB="d">
    <var id="ud_b" name="B"> 1.0 </var>
  </correlation>
</jastrow>

```

The XML describes $U_2^{\uparrow\uparrow/\downarrow\downarrow}(r)$ and $U_2^{\uparrow\downarrow}(r)$ from eq. ?? as Padé functions with initial guesses of $B^{\uparrow\uparrow} = 0.6 \text{ Bohr}^{-1}$ and $B^{\uparrow\downarrow} = 1.0 \text{ Bohr}^{-1}$ for the adjustable parameters.

The relevant portion of the input describing the linear optimization process is

```

<loop max="MAX">
  <qmc method="linear" move="pbyp" checkpoint="-1">
    <cost name="energy" > ECOST </cost>
    <cost name="unweightedvariance"> UVCOST </cost>
    <cost name="reweightedvariance" > RVCOST </cost>
    <parameter name="timestep" > TS </parameter>
    <parameter name="samples" > SAMPLES </parameter>
    <parameter name="warmupSteps" > 300 </parameter>
    <parameter name="blocks" > 800 </parameter>
    <parameter name="subSteps" > 10 </parameter>
    <parameter name="nonlocalpp" > yes </parameter>
    <parameter name="useBuffer" > yes </parameter>
    ...
  </qmc>
</loop>

```

An explanation of each input variable can be found below. The remaining variables control specialized internal details of the linear optimization algorithm. The meaning of these inputs is beyond the scope of this lab and reasonable results are often obtained keeping these values fixed.

energy Fraction of trial energy in the cost function.

unweightedvariance Fraction of unweighted trial variance in the cost function. Neglecting the weights can be more robust.

reweightedvariance Fraction of trial variance (including the full weights) in the cost function.

timestep Timestep of the VMC random walk, determines spatial distance moved by each electron during MC steps. Should be chosen such that the acceptance ratio of MC moves is around 50% (30-70% is often acceptable). Reasonable values are often between 0.2 and 0.6 Ha^{-1} .

samples Total number of MC samples collected for optimization, determines statistical error bar of cost function. Often efficient to start with a small number of samples (5-20k) and then increase (20-100k). More samples may be required if the wavefunction contains a large number of variational parameters. MUST be a multiple of the number of threads/cores (use multiples of 512 on Vesta).

warmupSteps Number of MC steps discarded as a warmup or equilibration period of the random walk. If this is too small, it will bias the optimization procedure.

blocks Number of average energy values written to output files. Should be greater than 200 for meaningful statistical analysis of output data (*e.g.* via `qmca`).

subSteps Number of MC steps in between energy evaluations. Each energy evaluation is expensive so taking a few steps to decorrelate between measurements can be more efficient. Will be less efficient with many substeps.

nonlocalpp,useBuffer If no, evaluate non-local pseudopotential derivatives approximately during optimization. This saves time and often does not affect optimization results unless the non-local contribution to the energy is large.

loop max Number of times to repeat the optimization. Using the resulting wavefunction from the previous optimization in the next one improves the results. Typical choices range between 4 and 20.

The three components of the cost function, energy, unreweighted variance, and reweighted variance should sum to one. Dedicating 100% of the cost function to unreweighted variance is often a good choice. Another common choice is to try 90/10 or 80/20 mixtures of reweighted variance and energy.

Replace `MAX`, `EVCOST`, `UVCOST`, `RVCOST`, `TS`, and `SAMPLES` in the two `loop`'s with appropriate starting values in the `0.q0.opt.in.xml` input file. Submit the optimization job to Vesta's queue by typing `./submit_0.q0_opt`. The job should only take a few minutes for reasonable values of `loop max` and `samples`.

Log file output will appear in `0.q0_opt.output`. The beginning of each linear optimization will be marked with text similar to

```
=====
Start QMCFixedSampleLinearOptimize
File Root 0_q0_opt.s011 append = no
=====
```

At the end of each optimization section the change in cost function, new values for the Jastrow parameters, and elapsed wallclock time are reported:

```
OldCost: 7.4701713964e-01 NewCost: 7.4681622535e-01 Delta Cost:-2.0091428584e-04
...
<optVariables href="O_q0_opt.s011.opt.xml">
e0_0 -9.5623201640e-01 1 1 ON 0
e0_1 -8.4728730387e-01 1 1 ON 1
e0_2 -6.8954452383e-01 1 1 ON 2
e0_3 -4.9327199567e-01 1 1 ON 3
e0_4 -3.2560096773e-01 1 1 ON 4
e0_5 -1.9567566480e-01 1 1 ON 5
e0_6 -1.2940405487e-01 1 1 ON 6
e0_7 -9.5221474839e-02 1 1 ON 7
uu_b 4.2002038228e-01 0 1 ON 8
ud_b 6.3472757070e-01 0 1 ON 9
</optVariables>
...
QMC Execution time = 7.0060820112e+00 secs
```

The cost function should decrease during each linear optimization (Delta cost < 0). Try “grep OldCost *.output”. You should see something like this:

```
OldCost: 1.3644746067e+00 NewCost: 1.1049104640e+00 Delta Cost:-2.5956414268e-01
OldCost: 1.0690085060e+00 NewCost: 8.3206148222e-01 Delta Cost:-2.3694702381e-01
OldCost: 7.8558402137e-01 NewCost: 7.2478477600e-01 Delta Cost:-6.0799245374e-02
OldCost: 7.3070322298e-01 NewCost: 7.1655770805e-01 Delta Cost:-1.4145514926e-02
OldCost: 1.2184771084e+00 NewCost: 1.1923197177e+00 Delta Cost:-2.6157390699e-02
OldCost: 6.8740347812e-01 NewCost: 6.8733036689e-01 Delta Cost:-7.3111228164e-05
OldCost: 6.9683928634e-01 NewCost: 6.9681780340e-01 Delta Cost:-2.1482934426e-05
OldCost: 6.7982953532e-01 NewCost: 6.7982948866e-01 Delta Cost:-4.6667065545e-08
OldCost: 6.8674328187e-01 NewCost: 6.8674327833e-01 Delta Cost:-3.5391565234e-09
OldCost: 7.5998537866e-01 NewCost: 7.5965629336e-01 Delta Cost:-3.2908530361e-04
OldCost: 7.0771416413e-01 NewCost: 7.0765392787e-01 Delta Cost:-6.0236255172e-05
OldCost: 7.4701713964e-01 NewCost: 7.4681622535e-01 Delta Cost:-2.0091428584e-04
```

Blocked averages of energy data, including the kinetic energy and components of the potential energy, are written to `scalar.dat` files. The first is named “`0_q0_opt.s000.scalar.dat`”, with a series number of zero (`s000`). In the end there will be `MAX1+MAX2` of them, one for each series.

When the job has finished, use the `qmca` tool to assess the effectiveness of the optimization process. To look at just the total energy and the variance, type “`qmca -q ev 0_q0_opt*scalar*`”. This will print the energy, variance, and the variance/energy ratio in Hartree units:

		LocalEnergy	Variance	ratio
0_q0_opt	series 0	-15.568764 +/- 0.003421	1.382681 +/- 0.056604	0.0888
0_q0_opt	series 1	-15.638500 +/- 0.005014	1.067662 +/- 0.019865	0.0683
0_q0_opt	series 2	-15.802163 +/- 0.002680	0.834521 +/- 0.007037	0.0528
0_q0_opt	series 3	-15.840982 +/- 0.001791	0.752242 +/- 0.009477	0.0475
0_q0_opt	series 4	-15.841584 +/- 0.003301	1.097355 +/- 0.252991	0.0693
0_q0_opt	series 5	-15.848602 +/- 0.003280	0.728377 +/- 0.019288	0.0460
0_q0_opt	series 6	-15.850839 +/- 0.001870	0.723159 +/- 0.008173	0.0456
0_q0_opt	series 7	-15.848411 +/- 0.002449	0.708589 +/- 0.007225	0.0447
...				

Plots of the data can also be obtained with the “`-p`” option (“`qmca -p -q ev 0_q0_opt*scalar*`”).

Identify which optimization series is the “best” according to your cost function. It is likely that multiple series are similar in quality. Note the `opt.xml` file corresponding to this series. This file contains the final value of the optimized Jastrow parameters to be used in the DMC calculations of the next section of the lab.

Questions and Exercises

1. What is the acceptance ratio of your optimization runs? (use “`qmca --help`” if necessary) Do you expect the Monte Carlo sampling to be efficient?
2. How do you know when the optimization process has converged?
3. Why is the mean and the error of the variance sometimes large? Consider using “`qmca -t ...`” to investigate.
4. Optimization is sometimes sensitive to initial guesses of the parameters. If you have time, try varying the initial parameters, including the cutoff radius (`rcut`) of the one-body Jastrow factor (remember to change `id` in the `<project/>` element). Do you arrive at a similar set of final Jastrow parameters? What is the lowest variance you are able to achieve?

13.5 DMC timestep extrapolation I: neutral O atom

The diffusion Monte Carlo (DMC) algorithm contains two biases in addition to the fixed node and pseudopotential approximations that are important to control: timestep and population control bias. The following subsection briefly discusses the origin of timestep and population control biases in DMC and how they can be minimized or extrapolated away. As before, the second subsection contains the lab walkthrough with QMCPACK. By the end of the section, we will have a solid DMC estimate of the ground state energy of oxygen.

In the same directory you used to perform wavefunction optimization (`oxygen.atom`) you will find a sample DMC input file for the neutral oxygen atom named `0.q0.dmc.in.xml`. Open this file in a text editor and note the differences from the optimization case. The XML describing the wavefunction is no longer present. In its place is the line

```
<include href="OPT_XML"/>
```

Replace “OPT_XML” with the `opt.xml` file corresponding to the best Jastrow parameters you found in the last section. The `include` element essentially amounts to an in-place copy and paste of the contents of the `opt.xml` file.

The QMC calculation section at the bottom is also different. The linear optimization blocks have been replaced with XML describing a VMC run followed by DMC. The input keywords are described below.

timestep Timestep of the VMC/DMC random walk. In VMC choose a timestep corresponding to an acceptance ratio of about 50%. In DMC the acceptance ratio is often above 99%.

warmupSteps Number of MC steps discarded as a warmup or equilibration period of the random walk.

steps Number of MC steps per block. Physical quantities, such as the total energy, are averaged over walkers and steps.

blocks Number of blocks. This is also the number of average energy values written to output files. Should be greater than 200 for meaningful statistical analysis of output data (*e.g.* via `qmca`). The total number of MC steps each walker takes is `blocks`×`steps`.

samples VMC only. This is the number of walkers used in subsequent DMC runs. Each DMC walker is initialized with electron positions sampled from the VMC random walk.

nonlocalmoves DMC only. If yes/no, use the locality approximation/T-moves for non-local pseudopotentials. T-moves generally improve the stability of the algorithm and restore the variational principle for small systems (T-moves version 1).

The purpose of the VMC run is to provide initial electron positions for each DMC walker. Setting `walkers = 1` in the VMC block ensures there will be only one VMC walker per execution thread. There will be a total of 512 VMC walkers in this case (see `0.q0.dmc.qsub.in`). We want the electron positions used to initialize the DMC walkers to be decorrelated from one another. A VMC walker will often decorrelate from its current position after propagating for a few Ha^{-1} in imaginary time (in general this is system dependent). This leads to a rough rule of thumb for choosing `blocks` and `steps` for the VMC run (`VWALKERS = 512` here):

$$\text{VBLOCKS} \times \text{VSTENexus} \geq \frac{\text{DWALKERS}}{\text{VWALKERS}} \frac{5 \text{ Ha}^{-1}}{\text{VTIMESTEP}} \quad (13.1)$$

Fill in the VMC XML block with appropriate values for these parameters. There should be more than one DMC walker per thread and enough walkers in total to avoid population control bias (see previous subsection).

To study timestep bias, we will perform a sequence of DMC runs over a range of timesteps (0.1 Ha^{-1} is too large and timesteps below 0.002 Ha^{-1} are probably too small). A common approach is to select a fairly large timestep to begin with and then decrease the timestep by a factor of two in each subsequent DMC run. The total amount of imaginary time the walker population propagates should be the same for each run. A simple way to accomplish this is to choose input parameters in the following way

$$\begin{aligned} \text{timestep}_n &= \text{timestep}_{n-1}/2 \\ \text{warmupSteps}_n &= \text{warmupSteps}_{n-1} \times 2 \\ \text{blocks}_n &= \text{blocks}_{n-1} \\ \text{steps}_n &= \text{steps}_{n-1} \times 2 \end{aligned} \quad (13.2)$$

Each DMC run will require about twice as much computer time as the one preceeding it. Note that the number of blocks is kept fixed for uniform statistical analysis. `blocks` \times `steps` \times `timestep` $\sim 60 \text{ Ha}^{-1}$ is sufficient for this system.

Choose an initial DMC timestep and create a sequence of N timesteps according to 13.2. Make N copies of the DMC XML block in the input file

```
<qmc method="dmc" move="pbyp">
  <parameter name="warmupSteps"      >  DWARMUP      </parameter>
  <parameter name="blocks"            >  DBLOCKS      </parameter>
  <parameter name="steps"             >  DSTENexus    </parameter>
  <parameter name="timestep"          >  DTIMESTEP    </parameter>
  <parameter name="nonlocalmoves"     >  yes          </parameter>
</qmc>
```

Fill in `DWARMUP`, `DBLOCKS`, `DSTENexus`, and `DTIMESTEP` for each DMC run according to 13.2. Submit the DMC timestep extrapolation run to the queue with `submit_0_q0_dmc`. The run should take only a few minutes to complete.

QMCPACK will create files prefixed with `0_q0_dmc`. The log file is `0_q0_dmc.output`. As before, block averaged data is written to `scalar.dat` files. In addition, DMC runs produce `dmc.dat` files which contain energy data averaged only over the walker population (one line per DMC step). The `dmc.dat` files also provide a record of the walker population at each step.

Use the `PlotTstepConv.pl` to obtain a linear fit to the timestep data (type “`PlotTstepConv.pl 0.q0.dmc.in.xml 40`”). You should see a plot similar to fig. 13.1. The tail end of the text output displays the parameters for the linear fit. The “a” parameter is the total energy extrapolated to zero timestep in Hartree units.

```
...
Final set of parameters          Asymptotic Standard Error
=====
a              = -15.8911        +/- 0.000756      (0.004757%)
b              = -0.221687       +/- 0.03757       (16.95%)
...
```

Questions and Exercises

1. What is the $\tau \rightarrow 0$ extrapolated value for the total energy?
2. What is the maximum timestep you should use if you want to calculate the total energy to an accuracy of 0.05 eV? For convenience, $1 \text{ Ha} = 27.2113846 \text{ eV}$.
3. What is the acceptance ratio for this (bias < 0.05 eV) run? Does it follow the rule of thumb for sensible DMC (acceptance ratio > 99%) ?
4. Check the fluctuations in the walker population (`qmca -t -q nw 0_q0_dmc*dmc.dat --noac`). Does the population seem to be stable?
5. (Optional) Study population control bias for the oxygen atom. Select a few population sizes (use multiples of 512 to fit cleanly on a single Vesta partition). Copy `0.q0.dmc.in.xml` to a new file and remove all but one DMC run (select a single timestep). Make one copy of




Figure 13.1: Linear fit to DMC timestep data from `PlotTstepConv.pl`.

the new file for each population, set “samples”, and choose a unique id in <project/>. Make submission files similar to `submit_0_q0_dmc` and `0.q0.dmc.qsub.in` and run one job at a time to avoid crowding the lab allocation. Use `qmca` to study the dependence of the DMC total energy on the walker population. How large is the bias compared to timestep error? What bias is incurred by following the “rule of thumb” of a couple thousand walkers? Will population control bias generally be an issue for production runs on modern parallel machines?

13.6 DMC timestep extrapolation II: O atom ionization potential

In this section, we will repeat the calculations of the prior two sections (optimization, timestep extrapolation) for the +1 charge state of the oxygen atom. Comparing the resulting 1st ionization potential (IP) with experimental data will complete our first test of the BFD oxygen pseudopotential. In actual practice, higher IP’s could also be tested prior to performing production runs.

Obtaining the timestep extrapolated DMC total energy for ionized oxygen should take much less (human) time than for the neutral case. For convenience, the necessary steps are briefly summarized below.

1. Copy the linear optimization blocks you used in `0.q0.opt.in.xml` to `0.q0.opt.in.xml`.
2. Submit the optimization job to Vesta’s queue with `submit_0_q1_opt`.
3. Identify the optimal set of parameters with `qmca`.
4. Replace `OPT_XML` in `submit_0_q1_dmc` with the `opt.xml` file containing the optimal parameters.
5. Copy the VMC and DMC blocks you used in `0.q0.dmc.in.xml` to `0.q1.dmc.in.xml`.
6. Submit the DMC timestep job to Vesta’s queue with `submit_0_q1_dmc`.
7. Obtain the DMC total energy extrapolated to zero timestep with `PlotTstepConv.pl`.

The process listed above, which excludes additional steps for orbital generation and conversion, can become tedious to perform by hand in production settings where many calculations are often required. For this reason automation tools are introduced for calculations involving the oxygen dimer in section 13.8 of the lab.

Questions and Exercises

1. What is the $\tau \rightarrow 0$ extrapolated DMC value for the 1st ionization potential of oxygen?

2. How does the extrapolated value compare to the experimental IP? Go to <http://physics.nist.gov/PhysRefData/ASD/ionEnergy.html> and enter “0 I” in the box labeled “Spectra” and click on the “Retrieve Data” button. For comparison the LDA value is 12.25 eV.
3. What can we conclude about the accuracy of the pseudopotential? What factors complicate this assessment?
4. Explore the sensitivity of the IP to the choice of timestep. Type “ip_conv.py” to view three timestep extrapolation plots: two for the $q = 0, 1$ total energies and one for the IP. Is the IP more, less, or similarly sensitive to timestep than the total energy?
5. What is the maximum timestep you should use if you want to calculate the ionization potential to an accuracy of 0.05 eV? What factor of cpu time is saved by assessing timestep convergence on the IP (a total energy difference) vs. a single total energy?
6. Are the acceptance ratio and population fluctuations reasonable for the $q = 1$ calculations?

13.7 DMC workflow automation with Nexus

Production QMC projects are often composed of many similar workflows. The simplest of these is a single DMC calculation involving four different compute jobs:

1. Orbital generation via Quantum Espresso or GAMESS.
2. Conversion of orbital data via `pw2qmcpack.x` or `convert4qmc`.
3. Optimization of Jastrow factors via QMCPACK.
4. DMC calculation via QMCPACK.

Simulation workflows quickly become more complex with increasing costs in terms of human time for the researcher. Automation tools can decrease both human time and error if used well.

The set of automation tools we will be using is known as Nexus, which is distributed with QMCPACK. Nexus is capable of generating input files, submitting and monitoring compute jobs, passing data between simulations (such as relaxed structures, orbital files, optimized Jastrow parameters, etc.), and data analysis. The user interface to Nexus is through a set of functions defined in the Python programming language. User scripts which execute simple workflows resemble input files and do not require programming experience. More complex workflows require only basic programming constructs (*e.g.* for loops and if statements). Nexus input files/scripts should be easier to navigate than QMCPACK input files and more efficient than submitting all the jobs by hand.

Nexus is driven by simple user-defined scripts that resemble keyword-driven input files. An example Nexus input file that performs a single VMC calculation is shown below. Take a moment to

read it over and especially note the comments (prefixed with “#”) explaining most of the contents. If the input syntax is unclear you may want to consult portions of appendix ??, which gives a condensed summary of Python constructs. For more information about the functionality and effective use of Nexus, consult docs/Nexus.pdf first. More information can be found in the user guide distributed with QMCPACK, although examples in this lab series and Nexus.pdf are more up to date (if qmcpack is the location of your QMCPACK distribution, the user guide can be found at qmcpack/nexus/documentation/nexus_user_guide.pdf).

```
#!/usr/bin/env python

# import Nexus functions
from nexus import settings, Job, get_machine, run_project
from nexus import generate_physical_system
from nexus import generate_qmcpack, vmc

settings(
    # Nexus settings
    pseudo_dir = './pseudopotentials', # location of PP files
    runs       = '',                  # root directory for simulations
    results    = '',                  # root directory for simulation results
    status_only = 0,                  # show simulation status, then exit
    generate_only = 0,                # generate input files, then exit
    sleep      = 3,                  # seconds between checks on sim. progress
    machine    = 'vesta',            # name of local machine
    account    = 'QMC_2014_training' # charge account for cpu time
)

vesta = get_machine('vesta') # allow max of one job at a time (lab only)
vesta.queue_size = 1

qmcjob = Job(
    # specify job parameters
    nodes = 32, # use 32 Vesta nodes
    threads = 16, # 16 OpenMP threads per node (32 MPI tasks)
    hours = 1, # wallclock limit of 1 hour
    # use QMCPACK executable
    app = '/soft/applications/qmcpack/build_XL_real/bin/qmcapp'
)

qmc_calcs = [
    # list QMC calculation methods
    vmc(
        # VMC
        walkers = 1, # 1 walker
        warmupsteps = 50, # 50 MC steps for warmup
        blocks = 200, # 200 blocks
        steps = 10, # 10 steps per block
        timestep = .4, # 0.4 1/Ha timestep
    )]

dimer = generate_physical_system(
    # make a dimer system
    type = 'dimer', # system type is dimer
)
```

```

dimer      = ('O','O'),          # dimer is two oxygen atoms
separation = 1.2074,             # separated by 1.2074 Angstrom
Lbox       = 15.0,               # simulation box is 15 Angstrom
units      = 'A',               # Angstrom is dist. unit
net_spin   = 2,                 # nup-down is 2
0          = 6                   # pseudo-oxygen has 6 valence el.
)

qmc = generate_qmcpack(          # make a qmcpack simulation
    identifier = 'example',      # prefix files with 'example'
    path       = 'scale_1.0',    # run in ./scale_1.0 directory
    system     = dimer,          # run the dimer system
    job        = qmcjob,         # set job parameters
    input_type = 'basic',        # basic qmcpack inputs given below
    pseudos    = ['O.BFD.xml'],  # list of PP's to use
    orbitals_h5 = 'O2.pwscf.h5', # file with orbitals from DFT
    bconds     = 'nnn',          # open boundary conditions
    jastrows   = [],             # no jastrow factors
    calculations = qmc_calcs     # QMC calculations to perform
)

run_project(qmc)                # write input file and submit job

```

13.8 Automated binding curve of the oxygen dimer

In this we will use Nexus to calculate the DMC total energy of the oxygen dimer over a series of bond lengths. The equilibrium bond length and binding energy of the dimer will be determined by performing a polynomial fit to the data (Morse potential fits should be preferred in production tests). Comparing these values with corresponding experimental data provides a second test of the BFD pseudopotential for oxygen.

Enter the `oxygen_dimer` directory. Copy your BFD pseudopotential from the atom runs into `oxygen_dimer/pseudopotentials`. Open `0_dimer.py` with a text editor. The overall format is similar to the example file shown in the last section. The header material, including Nexus imports, settings, and the job parameters for QMC are identical. The main difference is that optimization and DMC runs are being performed rather than a single VMC run.

Following the job parameters, inputs for the optimization method are given. The keywords should all be familiar from the QMCPACK XML input files you used previously:

```

linopt1 = linear(
    energy              = 0.0,
    unreweightedvariance = 1.0,

```

```

reweightedvariance = 0.0,
timestep           = 0.4,
samples            = 5000,
warmupsteps        = 50,
blocks             = 200,
substeps           = 1,
nonlocalpp         = True,
usebuffer          = True,
walkers            = 1,
minwalkers         = 0.5,
maxweight          = 1e9,
usedrift           = True,
minmethod          = 'quartic',
beta               = 0.025,
exp0               = -16,
bigchange          = 15.0,
allowedifference    = 1e-4,
stepsize           = 0.2,
stabilizerscale    = 1.0,
nstabilizers       = 3
)

```

Requesting multiple loop's with different numbers of samples is more compact than in XML:

```

linopt1 = ...

linopt2 = linopt1.copy()
linopt2.samples = 20000 # opt w/ 20000 samples

linopt3 = linopt1.copy()
linopt3.samples = 40000 # opt w/ 40000 samples

opt_calcs = [loop(max=8,qmc=linopt1), # loops over opt's
              loop(max=6,qmc=linopt2),
              loop(max=4,qmc=linopt3)]

```

The VMC/DMC method inputs should also look familiar:

```
qmc_calcs = [  
    vmc(  
        walkers      = 1,  
        warmupsteps  = 30,  
        blocks       = 20,  
        steps        = 10,  
        substeps     = 2,  
        timestep     = .4,  
        samples      = 2048  
    ),  
    dmc(  
        warmupsteps  = 100,  
        blocks       = 400,  
        steps        = 32,  
        timestep     = 0.01,  
        nonlocalmoves = True  
    )  
]
```

As in the example in the last section, the oxygen dimer is generated with the `generate_physical_system` function:

```
dimer = generate_physical_system(  
    type      = 'dimer',  
    dimer     = ('O', 'O'),  
    separation = 1.2074*scale,  
    Lbox      = 15.0,  
    units     = 'A',  
    net_spin  = 2,  
    0         = 6  
)
```

Similar syntax can be used to generate crystal structures or to specify systems with arbitrary atomic

configurations and simulation cells. Notice that a “**scale**” variable has been introduced to stretch or compress the dimer.

Next, objects representing QMCPACK simulations are constructed with the `generate_qmcpack` function:

```
opt = generate_qmcpack(
    identifier = 'opt',
    ...
    jastrows   = [('J1','bspline',8,4.5),
                  ('J2','pade',0.5,0.5)],
    calculations = opt_calcs
)
sims.append(opt)

qmc = generate_qmcpack(
    identifier = 'qmc',
    ...
    jastrows   = [],
    calculations = qmc_calcs,
    dependencies = (opt,'jastrow')
)
sims.append(qmc)
```

Shared details such as the run directory, job, pseudopotentials, and orbital file have been omitted (...). The “**opt**” run will optimize a 1-body B-spline Jastrow with 8 knots having a cutoff of 4.5 Bohr and a 2-body Padé Jastrow with up-up and up-down “**B**” parameters set to 0.5 1/Bohr. The Jastrow list for the DMC run is empty and a new keyword is present: **dependencies**. The usage of **dependencies** above indicates that the DMC run depends on the optimization run for the Jastrow factor. Nexus will submit the “**opt**” run first and upon completion it will scan the output, select the optimal set of parameters, pass the Jastrow information to the “**qmc**” run and then submit the DMC job. Independent job workflows are submitted in parallel when permitted (we have explicitly prevented this for this lab by setting `queue_size=1` for Vesta). No input files are written or job submissions made until the “**run_project**” function is reached.

As written, `0_dimer.py` will only perform calculations at the equilibrium separation distance of 1.2074 Angstrom. Modify the file now to perform DMC calculations across a range of separation distances with each DMC run using the Jastrow factor optimized at the equilibrium separation distance. The necessary Python for loop syntax should look something like this:

```

sims = []
for scale in [1.00,0.90,0.95,1.05,1.10]:
    ...
    dimer = ...
    if scale==1.00:
        opt = ...
        ...
    #end if
    qmc = ...
    ...
#end for
run_project(sims)

```

Note that the text inside the `for` loop and the `if` block must be indented by precisely four spaces. If you use Emacs, changes in indentation can be performed easily with `Cntrl-C >` and `Cntrl-C <` after highlighting a block of text (other editors should have similar functionality). If you see something like “`SyntaxError: invalid syntax`” print to the screen when you run `O_dimer.py` later on, consult the completed file in `oxygen_dimer/reference`.

The values of “`scale`” in the loop must be a subset of `[0.90,0.925,0.95,0.975,1.00,1.025,1.05,1.075,1.10]` since orbital files have been pre-generated with PWSCF for only these values. If other values are selected, the job will be submitted but QMCPACK will fail when it attempts to read the non-existent `02.pwscf.h5` file (in later labs we will run PWSCF to generate the orbital files directly with Nexus). Begin with the reduced set of `scale` values shown above.

Change the “`status_only`” parameter in the “`settings`” function to 1 and type “`./O_dimer.py`” at the command line. This will print the status of all simulations:

```

Project starting
  checking for file collisions
  loading cascade images
    cascade 0 checking in
  checking cascade dependencies
    all simulation dependencies satisfied
  cascade status
    setup, sent_files, submitted, finished, got_output, analyzed

```

```

000000 opt  ./scale_1.0
000000 qmc  ./scale_1.0
000000 qmc  ./scale_0.9
000000 qmc  ./scale_0.95
000000 qmc  ./scale_1.05
000000 qmc  ./scale_1.1
setup, sent_files, submitted, finished, got_output, analyzed

```

In this case, a single independent simulation “cascade” (workflow) has been identified, containing one “opt” and five dependent “qmc” runs. The six status flags (`setup`, `sent_files`, `submitted`, `finished`, `got_output`, `analyzed`) each show 0, indicating that no work has been done yet.

Now change “`status_only`” back to 0, set “`generate_only`” to 1, and run `0_dimer.py` again. This will perform a dry-run of all simulations. The dry-run should finish in about 20 seconds:

```

Project starting
  checking for file collisions
  loading cascade images
    cascade 0 checking in
  checking cascade dependencies
    all simulation dependencies satisfied

starting runs:
~~~~~

poll 0  memory 88.54 MB
  Entering ./scale_1.0 0
    writing input files 0 opt
  Entering ./scale_1.0 0
    sending required files 0 opt
    submitting job 0 opt
  Entering ./scale_1.0 1
    Would have executed: qsub --mode script --env BG_SHAREDMEMSIZE=32 opt.qsub.in

poll 1  memory 88.54 MB
  Entering ./scale_1.0 0
    copying results 0 opt
  Entering ./scale_1.0 0

```



```

    analyzing 0 opt

poll 2 memory 88.87 MB
  Entering ./scale_1.0 1
    writing input files 1 qmc
  Entering ./scale_1.0 1
    sending required files 1 qmc
    submitting job 1 qmc
  ...
  Entering ./scale_1.0 2
    Would have executed: qsub --mode script --env BG_SHAREDMEMSIZE=32 qmc.qsub.in
  ...

Project finished

```

Nexus polls the simulation status every 3 seconds and sleeps in between. The “scale_*” directories should now contain several files:

```

scale_1.0
  02.pwscf.h5
  0.BFD.xml
  opt.in.xml
  opt.qsub.in
  qmc.in.xml
  qmc.qsub.in
  sim_opt
    analyzer.p
    input.p
    sim.p
  sim_qmc
    analyzer.p
    input.p
    sim.p

```

Take a minute to inspect the generated input (`opt.in.xml`, `qmc.in.xml`) and submission (`opt.qsub.in`, `qmc.qsub.in`) files. The pseudopotential file `0.BFD.xml` has been copied into each local directory.

Two additional directories have been created: `sim_opt` and `sim_qmc`. The `sim.p` files in each directory contain the current status of each simulation. If you run `0_dimer.py` again, it should not attempt to rerun any of the simulations:

```
Project starting
  checking for file collisions
  loading cascade images
    cascade 0 checking in
    cascade 8 checking in
    cascade 2 checking in
    cascade 4 checking in
    cascade 6 checking in
  checking cascade dependencies
    all simulation dependencies satisfied

  starting runs:
  ~~~~~

  poll 0  memory 60.10 MB
Project finished
```

This way one can continue to add to the `0_dimer.py` file (*e.g.* adding more separation distances) without worrying about duplicate job submissions.

Let's actually submit the optimization and DMC jobs now. Reset the state of the simulations by removing the `sim.p` files (`rm ./scale*/sim*/sim.p`), set `generate_only` to 0, and rerun `0_dimer.py`. It should take about 15 minutes for all the jobs to complete. You may wish to open another terminal to monitor the progress of the individual jobs while the current terminal runs `0_dimer.py` in the foreground. You can begin the first exercise below once the optimization job completes.

Questions and Exercises

1. Evaluate the quality of the optimization at `scale=1.0` using the `qmca` tool. Did the optimization succeed? How does the variance compare with the neutral oxygen atom? Is the wavefunction of similar quality to the atomic case?
2. Evaluate the traces of the local energy and the DMC walker population for each separation distance with the `qmca` tool. Are there any anomalies in the runs? Is the acceptance ratio reasonable? Is the wavefunction of similar quality across all separation distances?
3. Use the `dimer_fit.py` tool located in `oxygen_dimer` to fit the oxygen dimer binding curve. To get the binding energy of the dimer, we will need the DMC energy of the atom. Before performing the fit, answer: What DMC timestep should be used for the oxygen atom results? The tool accepts three arguments ("`0_dimer.py P N E Eerr`"), `P` is the prefix of the DMC input files (should be "`qmc`" at this point), `N` is the order of the fit (use 2 to start), `E` and `Eerr` are your DMC total energy and error bar, respectively for the oxygen atom (in eV). A plot of the dimer data will be displayed and text output will show the DMC equilibrium bond length and binding energy as well as experimental values. How accurately does your fit to the DMC data reproduce the experimental values? What factors affect the accuracy of your results?
4. Refit your data with a fourth-order polynomial. How do your predictions change with a fourth-order fit? Is a fourth-order fit appropriate for the available data?
5. Add the four remaining "`scale`" values to the list in `0_dimer.py` that interpolate between the original set. Perform the DMC calculations and redo the fits. How accurately does your fit to the DMC data reproduce the experimental values? Should this pseudopotential be used in production calculations?
6. (Optional) Perform optimization runs at the extremal separation distances corresponding to `scale=[0.90,1.10]`. Are the individually optimized wavefunctions of significantly better quality than the one imported from `scale=1.00`? Why? What form of Jastrow factor might give an even better improvement?

13.9 (Optional) Running your system with QMCPACK

This section covers a fairly simple route to get started on QMC calculations of an arbitrary system of interest using Nexus (Nexus) automation system to setup input files and optionally perform the runs. The example provided in this section uses QM Espresso (PWSCF) to generate the orbitals forming the Slater determinant part of the trial wavefunction. PWSCF is a natural choice for solid state systems and it can be used for surface/slab and molecular systems as well, albeit at the price of describing additional vacuum space with plane waves.

To start out with, you will need pseudopotentials (PP's) for each element in your system in both the UPF (PWSCF) and FSATOM/XML (QMCPACK) formats. A good place to start is the

Burkatzki-Filippi-Dolg (BFD) pseudopotential database (<http://www.burkatzki.com/pseudos/index.2.html>), which we have already used in our study of the oxygen atom. The database does not contain PP's for the 4th and 5th row transition metals or any of the lanthanides or actinides. If you need a PP that is not in the BFD database, you may need to generate and test one manually (*e.g.* with OPIUM, <http://opium.sourceforge.net/>). Otherwise, use `ppconvert` as outlined in section ?? to obtain PP's in the formats used by PWSCF and QMCPACK. Enter the `your_system` lab directory and place the converted PP's in `your_system/pseudopotentials`.

Before performing production calculations (more than just the initial setup in this section) be sure to converge the plane wave energy cutoff in PWSCF as these PP's can be rather hard, sometimes requiring cutoffs in excess of 300 Ry. Depending on the system under study, the amount of memory required to represent the orbitals (QMCPACK uses 3D B-splines) becomes prohibitive and one may be forced to search for softer PP's.

Beyond pseudopotentials, all that is required to get started are the atomic positions and the dimensions/shape of the simulation cell. The Nexus file `example.py` illustrates how to setup PWSCF and QMCPACK input files by providing minimal information regarding the physical system (an 8-atom cubic cell of diamond in the example). Most of the contents should be familiar from your experience with the automated calculations of the oxygen dimer binding curve in section 13.8 (if you've skipped ahead you may want to skim that section for relevant information). The most important change is the expanded description of the physical system:

```
# details of your physical system (diamond conventional cell below)
my_project_name = 'diamond_vmc'    # directory to perform runs
my_dft_pps      = ['C.BFD.upf']    # pwscf pseudopotentials
my_qmc_pps      = ['C.BFD.xml']    # qmcpack pseudopotentials

# generate your system
#   units      : 'A'/'B' for Angstrom/Bohr
#   axes       : simulation cell axes in cartesian coordinates (a1,a2,a3)
#   elem       : list of atoms in the system
#   pos        : corresponding atomic positions in cartesian coordinates
#   kgrid      : Monkhorst-Pack grid
#   kshift     : Monkhorst-Pack shift (between 0 and 0.5)
#   net_charge : system charge in units of e
#   net_spin   : # of up spins - # of down spins
#   C = 4      : (pseudo) carbon has 4 valence electrons
my_system = generate_physical_system(
    units      = 'A',
    axes       = [[ 3.57000000e+00, 0.00000000e+00, 0.00000000e+00],
```

```

        [ 0.00000000e+00, 3.57000000e+00, 0.00000000e+00],
        [ 0.00000000e+00, 0.00000000e+00, 3.57000000e+00]],
    elem      = ['C','C','C','C','C','C','C','C'],
    pos       = [[ 0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
                  [ 8.92500000e-01, 8.92500000e-01, 8.92500000e-01],
                  [ 0.00000000e+00, 1.78500000e+00, 1.78500000e+00],
                  [ 8.92500000e-01, 2.67750000e+00, 2.67750000e+00],
                  [ 1.78500000e+00, 0.00000000e+00, 1.78500000e+00],
                  [ 2.67750000e+00, 8.92500000e-01, 2.67750000e+00],
                  [ 1.78500000e+00, 1.78500000e+00, 0.00000000e+00],
                  [ 2.67750000e+00, 2.67750000e+00, 8.92500000e-01]],
    kgrid     = (1,1,1),
    kshift    = (0,0,0),
    net_charge = 0,
    net_spin  = 0,
    C         = 4          # one line like this for each atomic species
)

my_bconds    = 'ppp'      # ppp/nnn for periodic/open BC's in QMC
                        # if nnn, center atoms about (a1+a2+a3)/2

```

If you have a system you would like to try with QMC, make a copy of `example.py` and fill in the relevant information about the pseudopotentials, simulation cell axes, and atomic species/positions. Otherwise, you can proceed with `example.py` as it is.

The other new aspects are two additional compute jobs to generate the orbitals with PWSCF and convert them into the ESHDF format with `pw2qmcpack.x`:

```

# scf run to generate orbitals
scf = generate_pwscf(
    identifier = 'scf',
    path       = my_project_name,
    job        = Job(nodes=32, hours=2, app=pwscf),
    input_type = 'scf',
    system     = my_system,
    pseudos    = my_dft_pps,
    input_dft  = 'lda',

```

```

    ecut          = 200,    # PW energy cutoff in Ry
    conv_thr      = 1e-8,
    mixing_beta   = .7,
    nosym         = True,
    wf_collect    = True
)

# conversion step to create h5 file with orbitals
p2q = generate_pw2qmcpack(
    identifier    = 'p2q',
    path          = my_project_name,
    job           = Job(cores=1, hours=2, app=pw2qmcpack),
    write_psi     = False,
    dependencies  = (scf, 'orbitals')
)

```

Set “generate_only” to 1 and type “./example.py” or similar to generate the input files. All files will be written to “./diamond_vmc” (“./[my_project_name]” if you have changed “my_project_name” in the file). The input files for PWSCF, pw2qmcpack, and QMCPACK are `scf.in`, `pw2qmcpack.in`, and `vmc.in.xml`, respectively. Take some time to inspect the generated input files. If you have questions about the file contents, or run into issues with the generation process, feel free to consult with a lab instructor.

If desired, you can submit the runs directly with `example.py`. To do this, first reset the Nexus simulation record by typing “`rm ./diamond_vmc/sim*/sim.p`” or similar and set “generate_only” back to 0. Next rerun `example.py` (you may want to redirect the text output).

Alternatively the runs can be submitted by hand:

```

qsub --mode script --env BG_SHAREDMEMSIZE=32 scf.qsub.in

(wait until JOB DONE appears in scf.output)

qsub --mode script --env BG_SHAREDMEMSIZE=32 p2q.qsub.in

```

Once the conversion process has finished the orbitals should be located in the file `diamond_vmc/pwscf_output/pwscf.pwscf.h5`. Open `diamond_vmc/vmc.in.xml` and replace “MISSING.h5” with “./pwscf_output/pwscf.pwscf.h5”. Next submit the VMC run:

```
qsub --mode script --env BG_SHAREDMEMSIZE=32 vmc.qsub.in
```

Note: If your system is large, the above process may not complete within the time frame of this lab. Working with a stripped down (but relevant) example is a good idea for exploratory runs.

Once the runs have finished, you may want to begin exploring Jastrow optimization and DMC for your system. Example calculations are provided at the end of `example.py` in the commented out text).

Chapter 14

Lab 4: Using PWSCF and QMCPACK to perform total energy calculations of condensed systems

Lab author: Luke Shulenburger¹

Creation date: July 17, 2014

The goal of this lab will be to introduce you to the somewhat specialized problems involved in performing diffusion Monte Carlo calculations on condensed matter as opposed to the atoms and molecules that were the focus of earlier labs. Calculations will be performed on two different systems. Firstly, we will perform a series of calculations on BCC beryllium focusing on the necessary methodology to limit finite size effects. Secondly, we will perform calculations on graphene as an example of a system where qmcpack's ability to handle cases with mixed periodic and open boundary conditions is useful. This example will also focus on strategies to limit memory usage for such systems. All of the calculations performed in this lab will utilize the Nexus workflow management system that vastly simplifies the process by automating the steps of generating trial wavefunctions and performing DMC calculations.

14.1 Preliminaries

For any DMC calculation, we must start with a trial wavefunction. As is typical for our calculations of condensed matter, we will produce this wavefunction using density functional theory. Specifically, we will use quantum espresso to generate a slater determinant of single particle orbitals. This is

¹Sandia National Laboratories is a multiprogram laboratory managed and operated by Sandia Corporation, a wholly owned subsidiary of Lockheed Martin Corporation, for the U.S. Department of Energy's National Nuclear Security Administration under Contract No. DE-AC04-94AL85000.

done as a three step process. First, we calculate the converged charge density by performing a DFT calculation with a fine grid of k-points to fully sample the Brillouin zone. Next, a non-self consistent calculation is performed at the specific k-points needed for the supercell and twists needed in the DMC calculation (more on this later). Finally, a wavefunction is converted from the binary representation used by quantum espresso to the portable hdf5 representation used by qmcpack.

The choice of k-points necessary to generate the wavefunctions depends on both the supercell chosen for the DMC calculation and by the supercell twist vectors needed. Recall that the wavefunction in a plane wave DFT calculation is written using Bloch's theorem as:

$$\Psi(\vec{r}) = e^{i\vec{k}\cdot\vec{r}}u(\vec{r}) \quad (14.1)$$

Where \vec{k} is confined to the first Brillouin zone of the cell chosen and $u(\vec{r})$ is periodic in this simulation cell. A plane wave DFT calculation stores the periodic part of the wavefunction as a linear combination of plane waves for each single particle orbital at all k-points selected. The symmetry of the system allows us to generate an arbitrary supercell of the primitive cell as follows: Consider the set of primitive lattice vectors, $\{\mathbf{a}_1^p, \mathbf{a}_2^p, \mathbf{a}_3^p\}$. We may write these vectors in a matrix, \mathbf{L}_p , whose rows are the primitive lattice vectors. Consider a non-singular matrix of integers, \mathbf{S} . A corresponding set of supercell lattice vectors, $\{\mathbf{a}_1^s, \mathbf{a}_2^s, \mathbf{a}_3^s\}$, can be constructed by the matrix product

$$\mathbf{a}_i^s = S_{ij}\mathbf{a}_j^p \quad (14.2)$$

If the primitive cell contains N_p atoms, the supercell will then contain $N_s = |\det(\mathbf{S})|N_p$ atoms.

Now, the wavefunction at any point in this new supercell can be related to the wavefunction in the primitive cell by finding the linear combination of primitive lattice vectors that maps this point back to the primitive cell:

$$\vec{r}' = \vec{r} + x\mathbf{a}_1^p + y\mathbf{a}_2^p + z\mathbf{a}_3^p = \vec{r} + \vec{T} \quad (14.3)$$

where x, y, z are integers. Now the wavefunction in the supercell at point \vec{r}' can be written in terms of the wavefunction in the primitive cell at \vec{r} as:

$$\Psi(\vec{r}) = \Psi(\vec{r}')e^{i\vec{T}\cdot\vec{k}} \quad (14.4)$$

where \vec{k} is confined to the first Brillouin zone of the primitive cell. We have also chosen the supercell twist vector which places a constraint on the form of the wavefunction in the supercell. The combination of these two constraints allows us to identify family of N k-points in the primitive cell that satisfy the constraints. Thus for a given supercell tiling matrix and twist angle, we can write the wavefunction everywhere in the supercell by knowing the wavefunction at N k-points in the primitive cell. This means that the memory necessary to store the wavefunction in a supercell is only linear in the size of the supercell rather than the quadratic cost if symmetry were neglected.

14.2 Total energy of BCC beryllium

As was discussed in this morning's lectures when performing calculations of periodic solids with QMC, it is essential to work with a reasonable size supercell rather than the primitive cells that are common in mean field calculations. Specifically, all of the finite size correction schemes discussed in the morning require that the exchange-correlation hole be considerably smaller than the periodic simulation cell. Additionally, finite size effects are lessened as the distance between the electrons in the cell and their periodic images increases, so it is advantageous to generate supercells that are as spherical as possible so as to maximize this distance. However, there is a competing consideration in that for calculating total energies we often want to be able to extrapolate the energy per particle to the thermodynamic limit by means of the following formula in 3 dimensions:

$$E_{\text{inf}} = C + E_N/N \quad (14.5)$$

This formula derived assuming the shape of the supercells is consistent (more specifically that the periodic distances scale uniformly with system size), meaning we will need to do a uniform tiling, ie, 2x2x2, 3x3x3 etc. As a 3x3x3 tiling is 27 times larger than the supercell and the practical limit of DMC is on the order of 200 atoms (depending on Z), sometimes it is advantageous to choose a less spherical supercell with fewer atoms rather than a more spherical one that is too expensive to tile.

In the case of a BCC crystal, it is possible to tile the one atom primitive cell to a cubic supercell by only doubling the number of electrons. This is the best possible combination of a small number of atoms that can be tiled and a regular box that maximizes the distance between periodic images. We will need to determine the tiling matrix S that generates this cubic supercell by solving the following equation for the coefficients of the S matrix:

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} s_{11} & s_{12} & s_{13} \\ s_{21} & s_{22} & s_{23} \\ s_{31} & s_{32} & s_{33} \end{bmatrix} \cdot \begin{bmatrix} 0.5 & 0.5 & -0.5 \\ -0.5 & 0.5 & 0.5 \\ 0.5 & -0.5 & 0.5 \end{bmatrix} \quad (14.6)$$

We will now use Nexus to generate the trial wavefunction for this BCC beryllium.

Fortunately, the Nexus will handle determination of the proper k-vectors given the tiling matrix. All that is needed is to place the tiling matrix in the Be-2at-setup.py file. Now the definition of the physical system is:

```
bcc_Be = generate_physical_system(  
    lattice    = 'cubic',  
    cell       = 'primitive',  
    centering  = 'I',
```

```

atoms      = 'Be',
constants  = 3.490,
units      = 'A',
net_charge = 0,
net_spin   = 0,
Be         = 2,
tiling     = [[a,b,c],[d,e,f],[g,h,i]],
kgrid      = kgrid,
kshift     = (.5,.5,.5)
)

```

Where the tiling line should be replaced with the row major tiling matrix from above. This script file will now perform a converged DFT calculation to generate the charge density in a directory called bcc-beryllium/scf and perform a non self consistend DFT calculation to generate single particle orbitals in the direcotry bcc-beryllium/nscf. Fortunately, Nexus will calculate the required k-points needed to tile the wavefunction to the supercell, so all that is necessary is the granularity of the supercell twists and whether this grid is shifted from the origin. Once this is finished, it performs the conversion from pwscf's binary format to the hdf5 format used by qmcpack. Finally, it will optimize the coefficients of one-body and two-body jastrow factors in the supercell defined by the tiling matrix.

Run these calculations by executing the script Be-2at-setup.py. You will notice that such small calcuations as are required to generate the wavefunction of Be in a one atom cell are rather inefficent to run on a high performance computer such as vesta in terms of the time spent doing calculations versus time waiting on the scheduler and booting compute nodes. One of the benefits of the portable hdf format that is used by qmcpack is that you can generate data like wavefunctions on a local workstation or other convenient resource and only use high performance clusters for the more expensive QMC calculations.

In this case, the wavefunction is generated in the directory bcc-beryllium/nscf-2at_222/pwscf_output in a file called pwscf.pwscf.h5. It can be useful for debugging purposes to be able to verify the contents of this file are what you expect. For instance, you can use the tool h5ls to check the geometry of the cell where the dft calculations were performed, or number of k-points or electrons in the calculation. This is done with the command: h5ls -d pwscf.pwscf.h5/supercell or h5ls -d pwscf.pwscf.h5/electrons.

In the course of running Be-2at-setup.py, you will get an error when attempting to perform the vmc and wavefunction optimization calculations. This is due to the fact that the wavefunction has been generated supercell twists of the form $(\pm 1/4, \pm 1/4, \pm 1/4)$. In the case that the supercell twist contains only 0 or $1/2$, it is possible to operate entirely with real arithmetic. The executabe that has been indicated in Be-2at-setup.py has been compiled for this case. Note

that where this is possible, the memory usage is a factor of two less than the general case and the calculations are somewhat faster. However, it is often necessary to perform calculations away from these special twist angles in order to reduce finite size effects. To fix this, delete the directory `bcc-beryllium/opt-2at`, change the line in near the top of `Be-2at-setup.py` from

```
qmcpack      = '/soft/applications/qmcpack/build_XL_real/bin/qmcapp'
```

to

```
qmcpack      = '/soft/applications/qmcpack/build_XL_complex/bin/qmcapp'
```

and rerun the script.

When the optimization calculation has finished, check that everything has proceeded correctly by looking at the output in the `opt-2at` directory. Firstly, you can grep the output file for Delta to see if the cost function has indeed been decreasing during the optimization. You should find something like:

```
OldCost: 4.8789147e-02 NewCost: 4.0695360e-02 Delta Cost:-8.0937871e-03
OldCost: 3.8507795e-02 NewCost: 3.8338486e-02 Delta Cost:-1.6930674e-04
OldCost: 4.1079105e-02 NewCost: 4.0898345e-02 Delta Cost:-1.8076319e-04
OldCost: 4.2681333e-02 NewCost: 4.2356598e-02 Delta Cost:-3.2473514e-04
OldCost: 3.9168577e-02 NewCost: 3.8552883e-02 Delta Cost:-6.1569350e-04
OldCost: 4.2176276e-02 NewCost: 4.2083371e-02 Delta Cost:-9.2903058e-05
OldCost: 4.3977361e-02 NewCost: 4.2865751e-02 Delta Cost:-1.1116183e-03
OldCost: 4.1420944e-02 NewCost: 4.0779569e-02 Delta Cost:-6.4137501e-04
```

Which shows that the starting wavefunction was fairly good and that most of the optimization occurred in the first step. Confirm this by using `qmca` to look at how the energy and variance changed over the course of the calculation with the command: `qmca -q ev -e 10 *.scalar.dat` executed in the `opt-2at` directory. You should get output like the following:

		LocalEnergy	Variance	ratio
opt	series 0	-2.159139 +/- 0.001897	0.047343 +/- 0.000758	0.0219
opt	series 1	-2.163752 +/- 0.001305	0.039389 +/- 0.000666	0.0182
opt	series 2	-2.160913 +/- 0.001347	0.040879 +/- 0.000682	0.0189
opt	series 3	-2.162043 +/- 0.001223	0.041183 +/- 0.001250	0.0190
opt	series 4	-2.162441 +/- 0.000865	0.039597 +/- 0.000342	0.0183
opt	series 5	-2.161287 +/- 0.000732	0.039954 +/- 0.000498	0.0185
opt	series 6	-2.163458 +/- 0.000973	0.044431 +/- 0.003583	0.0205
opt	series 7	-2.163495 +/- 0.001027	0.040783 +/- 0.000413	0.0189

Now that the optimization has completed successfully, we can perform dmc calculations. The first goal of the calculations will be to try to eliminate the one body finite size effects by twist averaging. The script `Be-2at-qmc.py` has the necessary input. Note on line 42 two twist grids are specified, (2,2,2) and (3,3,3). Change the tiling matrix in this input file as in `Be-2at-qmc.py` and start the calculations. Note that this workflow takes advantage of `qmcpack`'s ability to group jobs. If you look in the directory `dmc-2at.222` at the job submission script, (`dmc.qsub.in`) you will note that rather than operating on an xml input file, `qmcpapp` is targeting a text file called `dmc.in`. This file is a simple text file that contains the names of the 8 xml input files needed for this job, one for each twist. When operated in this mode, `qmcpack` will use `mpi` groups to run multiple copies of itself within the same `mpi` context. This is often useful both in terms of organizing calculations and also for taking advantage of the large job sizes that computer centers often encourage.

The dmc calculations in this case are designed to complete in a few minutes. When they have finished running, first look at the `scalar.dat` files corresponding to the dmc calculations at the various twists in `dmc-2at.222`. Using a command like `'qmca -q ev -e 32 *.s001.scalar.dat'` (with a suitably chosen number of blocks for the equilibration), you will see that the dmc energy in each calculation is nearly identical within the statistical uncertainty of the calculations. In the case of a large supercell, this is often indicative of a situation where the Brillouin zone is so small that the one body finite size effects are nearly converged without any twist averaging. In this case, however, this is because of the symmetry of the system. For this cubic supercell, all of the twist angles chosen in this shifted 2x2x2 grid are equivalent by symmetry. In the case where substantial resources are required to equilibrate the dmc calculations, it can be beneficial to avoid repeating such twists and instead simply weight them properly. In this case however where the equilibration is inexpensive, there is no benefit to adding such complexity as the calculations can simply be averaged together and the result is equivalent to performing a single longer calculation.

Using the command `qmc -a -q ev -e 16 *.s001.scalar.dat`, average the dmc energies in `dmc-2at.222` and `dmc-2at.333` to see whether the one body finite size effects are converged with a 3x3x3 grid of twists. As beryllium as a metal, the convergence is quite poor (0.025 Ha / Be or 0.7 eV

/ Be). If this were a production calculation it would be necessary to perform calculations on much larger grids of supercell twists to eliminate the one body finite size effects.

In this case there are several other calculations that would warrent a high priority. A script `Be-16at-qmc.py` has been provided where you can input the appropriate tiling matrix for a 16 atom cell and perform calculations to estimate the two body finite size effects which will also be quite large in the 2 atom calculations. This script will take approximately 30 minutes to run to completion, so depending on interest, you can either run it, or also work to modify the scripts to address the other technical issues that would be necessary for a production calculation such as calculating the population bias or the timestep error in the dmc calculations.

Another useful exercise would be to attempt to validate this pseudopotential by calculating the ionization potential and electron affinity of the isolated atom and comparing to the experimental values: $IP = 9.3227 \text{ eV}$, $EA = 2.4 \text{ eV}$.

14.3 Handling a 2D system: graphene

In this section we will examine a calculation of an isolated sheet of graphene. As graphene is a two dimensional system, we will take advantage of `qmcpack`'s ability to mix periodic and open boundary conditions to eliminate and spurious interaction of the sheet with its images in the z direction. Run the script `graphene-setup.py` which will generate the wavefunction and optimize one and two body jastrow factors. In the script, notice line 160: `bconds = 'ppn'` in the `generate_qmcpack` function which specifies this mix of open and periodic boundary conditions. As a consequence of this, the atoms will need to be kept away from this open boundary in the z direction as the electronic wavefunction will not be defined outside of the simulation box in this direction. For this reason, all of the atom positions in at the beginning of the file have z coordinates 7.5. At this point, run the script `graphene-setup.py`.

Aside from the change in boundary conditions, the main thing that distinguished this kind of calculation from the beryllium example above is the large amount of vacuum in the cell. While this is a very small calculation designed to run quickly in the tutorial, in general a more converged calculation would quickly become memory limited on an architecture like BG/Q. When the initial wavefunction optimization has completed to your satisfaction, run the scripts `graphene-loop-buffer.py` and `graphene-loop-mesh.py`. These examine within variational Monte Carlo two approaches to reducing the memory required to store the wavefunction. In `graphene-loop-mesh.py`, the spacing between the b-spline points is varied uniformly. The mesh spacing is a prefactor to the linear spacing between the spline points, so the memory usage goes as the cube of the meshfactor. When you run the calculations, examine the `.s000.scalar.dat` files with `qmca` to determine the lowest possible mesh spacing that preserves both the vmc energy and the variance. Similarly, the script `graphene-loop-buffer.py` uses a feature which generates two spline tables for the wavefunction. One will have half of the mesh spacing requested in the input file and will be valid everywhere. The second one will only be defined in the smallest parallelepiped that contains all of the atoms in

the simulation cell with minimum distance given by the buffer size. Again, see what the smallest possible buffer size is that preserves the vmc energy and variance.

Finally, edit the file `graphene-final.py` which will perform two DMC calculations. In the first, (qmc1) replace the following lines:

```
meshfactor    = xxx,  
precision     = '---',  
truncate      = False,  
buffer        = 0.0,
```

using the values you have determined to perform the calculation with as small as possible of wavefunction. Note that we can also use single precision arithmetic to store the wavefunction by specifying `precision='single'`. When you run the script, compare the output of the two DMC calculations in terms of energy and variance. Also see if you can calculate the fraction of memory that you were able to save by using a meshfactor other than 1, a buffer table and single precision arithmetic.

14.4 Conclusion

Upon completion of this lab, you should be able to use Nexus to perform DMC calculations on periodic solids when provided with a pseudopotential. You should also be able to reduce the size of the wavefunction in a solid state calculation in cases where memory is a limiting factor.

14.5 Acknowledgment

This tutorial was created with support from Sandia National Laboratories.

Sandia National Laboratories is a multiprogram laboratory managed and operated by Sandia Corporation, a wholly owned subsidiary of Lockheed Martin Corporation, for the U.S. Department of Energy's National Nuclear Security Administration under Contract No. DE-AC04-94AL85000.

Chapter 15

Contributing to the Manual

This section briefly describes how to contribute to the manual. It is primarily “by developers, for developers”. This section should iterate until a consistent view on style/contents is reached.

Desirable:

- Use the table templates below when describing XML input.
- Place unformatted text targeted at developers in comments. Include generously.
- Encapsulate formatted text aimed at developers (like this entire chapter), in `\dev{}`. Text encapsulated in this way will be removed from the user version of the manual by editing the definition of `\dev` in `qmcpack_manual.tex`. Existing but deprecated or partially functioning features fall in this category.

Missing sections (these are opinions, not decided priorities):

- Description of XML input in general. Discuss XML format, use of attributes and `<parameter/>`'s in general, case sensitivity (input is generally case sensitive), and behavior of QMCPACK when unrecognized XML elements are encountered (they are generally ignored without notification).
- Overview of the input file in general, broad structure, and at least one full example that works in isolation.

Information currently missing for a complete reference specification:

- Noting how many instances of each child element are allowed. Examples: `simulation-1` only, `method-1` or more, `jastrow-0` or more.

Below are template tables for describing XML elements in reference fashion. A number of examples can be found in *e.g.* Chapter 8. Preliminary style is (please weigh in with opinions): typewriter

text (`\texttt{}`) for XML element, attribute, and parameter names, normal text for literal information in datatype/values/default columns, bold (`\textbf{}`) text if an attribute/parameter must take on a particular value (values column), italics (`\textit{}`) for descriptive (non-literal) information in the values column (e.g. *anything*, *non-zero*, etc.), required/optional attributes/parameters noted by `some_attrr`/`some_attro` superscripts. Valid datatypes are text, integer, real, boolean, and arrays of each. Fixed length arrays can be noted, *e.g.* by “real array(3)”.

Template for a generic XML element:

generic element				
parent elements:	parent1 parent2			
child elements:	child1 child2 child3 ...			
attributes				
name	datatype	values	default	description
attr1 ^r	text			
attr2 ^r	integer			
attr3 ^o	real			
attr4 ^o	boolean			
attr5 ^o	text array			
attr6 ^o	integer array			
attr7 ^o	real array			
attr8 ^o	boolean array			
parameters				
name	datatype	values	default	description
param1 ^r	text			
param2 ^r	integer			
param3 ^o	real			
param4 ^o	boolean			
param5 ^o	text array			
param6 ^o	integer array			
param7 ^o	real array			
param8 ^o	boolean array			
body text				
	Long form description of body text format			

“Factory” elements are XML elements that share a tag, but whose contents change based on the value an attribute (or sometimes multiple attributes take). The attribute(s) that determine the allowed contents is referred to below as the “type selector” (*e.g.* for `<estimator/>` elements, the type selector is usually the `type` attribute). These types of elements are frequently encountered as they correspond (sometimes loosely, sometimes literally) to polymorphic classes in QMCPACK

that are built in “factories”. This name is true to the underlying code, but may be obscure to the general user (is there a better name to retain the general meaning?).

The template below should be provided each time a new “factory” type is encountered (like `<estimator/>`). The table lists all types of possible elements (see “type options” below) and any attributes that are common to all possible related elements. Specific “derived” elements are then described one at a time with the template above, noting the type selector in addition to the XML tag (*e.g.* “`estimator type=density` element”).

Template for shared information about “factory” elements.

generic factory element				
parent elements:	parent1 parent2			
child elements:	child1 child2 child3 ...			
type selector:	some attribute			
type options :	Selection1			
	Selection2			
	Selection3			
	...			
shared attributes:				
name	datatype	values	default	description
attr1	text			
attr2	integer			
...				

Chapter 16

References